

(Model Training)INM434 Natural Language Processing_Yumi Heo_code

May 8, 2024

INM434 Natural Language Processing MSc Data Science | Yumi Heo | 230003122

Google Colab Folder Link: <https://drive.google.com/drive/folders/1Yn99YR6d5iJ79NYjdZwLLUTEpmucnNqH?u>

Model Training Code Google Colab Link: <https://colab.research.google.com/drive/1nTqhQt6vPr6GIj6mX9jmtQia>

Model Test Code Google Colab Link: <https://colab.research.google.com/drive/1OgBrPqldusG9K2o68J3SBCF-pTexFih0?usp=sharing>

1 Intent Classification for Bank Customer Queries

1.1 Import the dataset 'banking77'

```
[ ]: # Mount the drive to save and load files and models
from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
[ ]: # Install datasets from huggingface
!pip install datasets
```

Collecting datasets

Downloading datasets-2.19.1-py3-none-any.whl (542 kB)

542.0/542.0

kB 8.1 MB/s eta 0:00:00

Requirement already satisfied: filelock in /usr/local/lib/python3.10/dist-packages (from datasets) (3.14.0)

Requirement already satisfied: numpy>=1.17 in /usr/local/lib/python3.10/dist-packages (from datasets) (1.25.2)

Requirement already satisfied: pyarrow>=12.0.0 in /usr/local/lib/python3.10/dist-packages (from datasets) (14.0.2)

Requirement already satisfied: pyarrow-hotfix in /usr/local/lib/python3.10/dist-packages (from datasets) (0.6)

Collecting dill<0.3.9,>=0.3.0 (from datasets)

Downloading dill-0.3.8-py3-none-any.whl (116 kB)

116.3/116.3

kB 12.5 MB/s eta 0:00:00

Requirement already satisfied: pandas in /usr/local/lib/python3.10/dist-packages (from datasets) (2.0.3)

Requirement already satisfied: requests>=2.19.0 in /usr/local/lib/python3.10/dist-packages (from datasets) (2.31.0)

Requirement already satisfied: tqdm>=4.62.1 in /usr/local/lib/python3.10/dist-packages (from datasets) (4.66.4)

Collecting xxhash (from datasets)

Downloading

xxhash-3.4.1-cp310-cp310-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (194 kB)

194.1/194.1

kB 9.8 MB/s eta 0:00:00

Collecting multiprocessing (from datasets)

Downloading multiprocessing-0.70.16-py310-none-any.whl (134 kB)

134.8/134.8

kB 14.0 MB/s eta 0:00:00

Requirement already satisfied: fsspec[http]<=2024.3.1,>=2023.1.0 in /usr/local/lib/python3.10/dist-packages (from datasets) (2023.6.0)

Requirement already satisfied: aiohttp in /usr/local/lib/python3.10/dist-packages (from datasets) (3.9.5)

Collecting huggingface-hub>=0.21.2 (from datasets)

Downloading huggingface_hub-0.23.0-py3-none-any.whl (401 kB)

401.2/401.2

kB 12.4 MB/s eta 0:00:00

Requirement already satisfied: packaging in /usr/local/lib/python3.10/dist-packages (from datasets) (24.0)

Requirement already satisfied: pyyaml>=5.1 in /usr/local/lib/python3.10/dist-packages (from datasets) (6.0.1)

Requirement already satisfied: aiosignal>=1.1.2 in /usr/local/lib/python3.10/dist-packages (from aiohttp->datasets) (1.3.1)

Requirement already satisfied: attrs>=17.3.0 in /usr/local/lib/python3.10/dist-packages (from aiohttp->datasets) (23.2.0)

Requirement already satisfied: frozenlist>=1.1.1 in /usr/local/lib/python3.10/dist-packages (from aiohttp->datasets) (1.4.1)

Requirement already satisfied: multidict<7.0,>=4.5 in /usr/local/lib/python3.10/dist-packages (from aiohttp->datasets) (6.0.5)

Requirement already satisfied: yarl<2.0,>=1.0 in /usr/local/lib/python3.10/dist-packages (from aiohttp->datasets) (1.9.4)

Requirement already satisfied: async-timeout<5.0,>=4.0 in /usr/local/lib/python3.10/dist-packages (from aiohttp->datasets) (4.0.3)

Requirement already satisfied: typing-extensions>=3.7.4.3 in /usr/local/lib/python3.10/dist-packages (from huggingface-hub>=0.21.2->datasets) (4.11.0)

Requirement already satisfied: charset-normalizer<4,>=2 in

```

/usr/local/lib/python3.10/dist-packages (from requests>=2.19.0->datasets)
(3.3.2)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-
packages (from requests>=2.19.0->datasets) (3.7)
Requirement already satisfied: urllib3<3,>=1.21.1 in
/usr/local/lib/python3.10/dist-packages (from requests>=2.19.0->datasets)
(2.0.7)
Requirement already satisfied: certifi>=2017.4.17 in
/usr/local/lib/python3.10/dist-packages (from requests>=2.19.0->datasets)
(2024.2.2)
Requirement already satisfied: python-dateutil>=2.8.2 in
/usr/local/lib/python3.10/dist-packages (from pandas->datasets) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-
packages (from pandas->datasets) (2023.4)
Requirement already satisfied: tzdata>=2022.1 in /usr/local/lib/python3.10/dist-
packages (from pandas->datasets) (2024.1)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-
packages (from python-dateutil>=2.8.2->pandas->datasets) (1.16.0)
Installing collected packages: xxhash, dill, multiprocessing, huggingface-hub,
datasets
  Attempting uninstall: huggingface-hub
    Found existing installation: huggingface-hub 0.20.3
    Uninstalling huggingface-hub-0.20.3:
      Successfully uninstalled huggingface-hub-0.20.3
Successfully installed datasets-2.19.1 dill-0.3.8 huggingface-hub-0.23.0
multiprocess-0.70.16 xxhash-3.4.1

```

```

[ ]: # Import libraries
import datasets
from datasets import load_dataset
import numpy as np
import matplotlib.pyplot as plt
import re
import string
import nltk
from nltk.corpus import stopwords
from pprint import pprint
from collections import Counter
from sklearn.metrics import precision_score, recall_score, f1_score

```

```

[ ]: # Take the dataset 'banking77'
banking77 = load_dataset('banking77')

```

```

/usr/local/lib/python3.10/dist-packages/huggingface_hub/utils/_token.py:89:
UserWarning:
The secret `HF_TOKEN` does not exist in your Colab secrets.
To authenticate with the Hugging Face Hub, create a token in your settings tab
(https://huggingface.co/settings/tokens), set it as secret in your Google Colab

```

and restart your session.

You will be able to reuse this secret in all of your notebooks.

Please note that authentication is recommended but still optional to access public models or datasets.

```
warnings.warn(
```

```
Downloading readme: 0%|          | 0.00/14.4k [00:00<?, ?B/s]
```

```
Downloading data: 0%|          | 0.00/298k [00:00<?, ?B/s]
```

```
Downloading data: 0%|          | 0.00/93.9k [00:00<?, ?B/s]
```

```
Generating train split: 0%|        | 0/10003 [00:00<?, ? examples/s]
```

```
Generating test split: 0%|         | 0/3080 [00:00<?, ? examples/s]
```

```
[ ]: # Check its structure
      banking77
```

```
[ ]: DatasetDict({
      train: Dataset({
        features: ['text', 'label'],
        num_rows: 10003
      })
      test: Dataset({
        features: ['text', 'label'],
        num_rows: 3080
      })
    })
```

```
[ ]: # Check the data format of the dataset
      banking77.cache_files

      # Arrow type
```

```
[ ]: {'train': [{'filename': '/root/.cache/huggingface/datasets/banking77/default/0.0
      .0/f54121560de48f2852f90be299010d1d6dc612ec/banking77-train.arrow'}],
      'test': [{'filename': '/root/.cache/huggingface/datasets/banking77/default/0.0.
      0/f54121560de48f2852f90be299010d1d6dc612ec/banking77-test.arrow'}]}
```

1.2 Data Preprocessing: Remove numbers, punctuations, double white spaces and apply lowercase

```
[ ]: # Define the function to remove numbers
      def remove_num(example):
          return {'text': re.sub(r'\d+', '', example['text'])}
```

```
[ ]: # Mapping the lowercase function
      banking77_wonum = banking77.map(remove_num)
```

Map: 0%| | 0/10003 [00:00<?, ? examples/s]

Map: 0%| | 0/3080 [00:00<?, ? examples/s]

```
[ ]: # Apply lowercase to all texts in the dataset
# Define the function
def lowercase(example):
    return {'text': example['text'].lower()}
```

```
[ ]: # Mapping the lowercase function
banking77_lowercase = banking77_wonum.map(lowercase)
```

Map: 0%| | 0/10003 [00:00<?, ? examples/s]

Map: 0%| | 0/3080 [00:00<?, ? examples/s]

```
[ ]: # Remove punctuations in the dataset
# Check the punctuation
string.punctuation
```

```
[ ]: '!"#$%&\'()*+,-./:;<=>?@[\\]^_`{|}~'
```

```
[ ]: # Define the removing punctuations function
def remove_punctuations(example):
    return {'text': example['text'].translate(str.maketrans('', '', string.
    ↪punctuation))}
```

```
[ ]: # Mapping the function
banking77_wopunc = banking77_lowercase.map(remove_punctuations)
```

Map: 0%| | 0/10003 [00:00<?, ? examples/s]

Map: 0%| | 0/3080 [00:00<?, ? examples/s]

```
[ ]: # Define the removing double white spaces function
def remove_doublespaces(example):
    return {'text': re.sub('\s+', ' ', example['text']).strip()}
```

```
[ ]: # Mapping the function
banking77_wods = banking77_wopunc.map(remove_doublespaces)
```

Map: 0%| | 0/10003 [00:00<?, ? examples/s]

Map: 0%| | 0/3080 [00:00<?, ? examples/s]

```
[ ]: # Take stopwords in English version
nltk.download('stopwords')
eng_stop = set(stopwords.words("english"))

# Define the removing double white spaces function
```

```
def remove_stopwords(example):
    words = example['text'].split() # Split the text into words
    filtered_words = [word for word in words if word.lower() not in eng_stop]
    ↪ # Remove stopwords
    return {'text': ' '.join(filtered_words)} # Join the filtered words back
    ↪ into a single string
```

[nltk_data] Downloading package stopwords to /root/nltk_data...

[nltk_data] Unzipping corpora/stopwords.zip.

```
[ ]: # Mapping the function
banking77_preprocessed = banking77_wods.map(remove_stopwords)
```

Map: 0%| | 0/10003 [00:00<?, ? examples/s]

Map: 0%| | 0/3080 [00:00<?, ? examples/s]

```
[ ]: # Extract training set and test set
trainset = banking77_preprocessed['train']
testset = banking77_preprocessed['test']
```

```
[ ]: # Check columns in the traininig set
trainset.column_names
```

```
[ ]: ['text', 'label']
```

```
[ ]: # Check the featrues in training set
print(trainset.features)
```

```
{'text': Value(dtype='string', id=None), 'label':
ClassLabel(names=['activate_my_card', 'age_limit', 'apple_pay_or_google_pay',
'atm_support', 'automatic_top_up', 'balance_not_updated_after_bank_transfer',
'balance_not_updated_after_cheque_or_cash_deposit', 'beneficiary_not_allowed',
'cancel_transfer', 'card_about_to_expire', 'card_acceptance', 'card_arrival',
'card_delivery_estimate', 'card_linking', 'card_not_working',
'card_payment_fee_charged', 'card_payment_not_recognised',
'card_payment_wrong_exchange_rate', 'card_swallowed', 'cash_withdrawal_charge',
'cash_withdrawal_not_recognised', 'change_pin', 'compromised_card',
'contactless_not_working', 'country_support', 'declined_card_payment',
'declined_cash_withdrawal', 'declined_transfer',
'direct_debit_payment_not_recognised', 'disposable_card_limits',
'edit_personal_details', 'exchange_charge', 'exchange_rate', 'exchange_via_app',
'extra_charge_on_statement', 'failed_transfer', 'fiat_currency_support',
'get_disposable_virtual_card', 'get_physical_card', 'getting_spare_card',
'getting_virtual_card', 'lost_or_stolen_card', 'lost_or_stolen_phone',
'order_physical_card', 'passcode_forgotten', 'pending_card_payment',
'pending_cash_withdrawal', 'pending_top_up', 'pending_transfer', 'pin_blocked',
'receiving_money', 'Refund_not_showing_up', 'request_refund',
```

```
'reverted_card_payment?', 'supported_cards_and_currencies', 'terminate_account',
'top_up_by_bank_transfer_charge', 'top_up_by_card_charge',
'top_up_by_cash_or_cheque', 'top_up_failed', 'top_up_limits', 'top_up_reverted',
'topping_up_by_card', 'transaction_charged_twice', 'transfer_fee_charged',
'transfer_into_account', 'transfer_not_received_by_recipient',
'transfer_timing', 'unable_to_verify_identity', 'verify_my_identity',
'verify_source_of_funds', 'verify_top_up', 'virtual_card_not_working',
'visa_or_mastercard', 'why_verify_identity', 'wrong_amount_of_cash_received',
'wrong_exchange_rate_for_cash_withdrawal'], id=None))
```

```
[ ]: # Check the first five elements in the training set
pprint(trainset[:5], sort_dicts=False)
pprint(testset[:5], sort_dicts=False)
```

```
{'text': ['still waiting card',
          'card still hasnt arrived weeks',
          'waiting week card still coming',
          'track card process delivery',
          'know get card lost'],
 'label': [11, 11, 11, 11, 11]}
{'text': ['locate card',
          'still received new card ordered week ago',
          'ordered card arrived help please',
          'way know card arrive',
          'card arrived yet'],
 'label': [11, 11, 11, 11, 11]}
```

All preprocessing steps have been applied.

```
[ ]: # Describe general information of the training set
pprint(trainset.info)
```

```
DatasetInfo(description='',
             citation='',
             homepage='',
             license='',
             features={'label': ClassLabel(names=['activate_my_card',
                                                  'age_limit',
                                                  'apple_pay_or_google_pay',
                                                  'atm_support',
                                                  'automatic_top_up',
                                                  'balance_not_updated_after_bank_transfer',
                                                  'balance_not_updated_after_cheque_or_cash_deposit',
                                                  'beneficiary_not_allowed',
                                                  'cancel_transfer',
                                                  'card_about_to_expire',
                                                  'card_acceptance',
                                                  'card_arrival',
                                                  'card_delivery_estimate',
```

	'card_linking',
	'card_not_working',
	'card_payment_fee_charged',
	'card_payment_not_recognised',
'card_payment_wrong_exchange_rate',	
	'card_swallowed',
	'cash_withdrawal_charge',
'cash_withdrawal_not_recognised',	
	'change_pin',
	'compromised_card',
	'contactless_not_working',
	'country_support',
	'declined_card_payment',
	'declined_cash_withdrawal',
	'declined_transfer',
'direct_debit_payment_not_recognised',	
	'disposable_card_limits',
	'edit_personal_details',
	'exchange_charge',
	'exchange_rate',
	'exchange_via_app',
	'extra_charge_on_statement',
	'failed_transfer',
	'fiat_currency_support',
	'get_disposable_virtual_card',
	'get_physical_card',
	'getting_spare_card',
	'getting_virtual_card',
	'lost_or_stolen_card',
	'lost_or_stolen_phone',
	'order_physical_card',
	'passcode_forgotten',
	'pending_card_payment',
	'pending_cash_withdrawal',
	'pending_top_up',
	'pending_transfer',
	'pin_blocked',
	'receiving_money',
	'Refund_not_showing_up',
	'request_refund',
	'reverted_card_payment?',
'supported_cards_and_currencies',	
	'terminate_account',
'top_up_by_bank_transfer_charge',	
	'top_up_by_card_charge',
	'top_up_by_cash_or_cheque',
	'top_up_failed',
	'top_up_limits',


```

        'top_up_reverted',
        'topping_up_by_card',
        'transaction_charged_twice',
        'transfer_fee_charged',
        'transfer_into_account',

        'transfer_not_received_by_recipient',

        'transfer_timing',
        'unable_to_verify_identity',
        'verify_my_identity',
        'verify_source_of_funds',
        'verify_top_up',
        'virtual_card_not_working',
        'visa_or_mastercard',
        'why_verify_identity',

        'wrong_amount_of_cash_received',
        'wrong_exchange_rate_for_cash_withdrawal'],
        id=None),
        'text': Value(dtype='string', id=None)},
    post_processed=None,
    supervised_keys=None,
    task_templates=None,
    builder_name='parquet',
    dataset_name='banking77',
    config_name='default',
    version=0.0.0,
    splits={'test': SplitInfo(name='test',
                              num_bytes=204395,
                              num_examples=3080,
                              shard_lengths=None,
                              dataset_name='banking77'),
            'train': SplitInfo(name='train',
                               num_bytes=716279,
                               num_examples=10003,
                               shard_lengths=None,
                               dataset_name='banking77')},
    download_checksums={'hf://datasets/banking77@f54121560de48f2852f90be
299010d1d6dc612ec/data/test-00000-of-00001.parquet': {'checksum': None,
                                                         'num_bytes': 93870},
                       'hf://datasets/banking77@f54121560de48f2852f90be
299010d1d6dc612ec/data/train-00000-of-00001.parquet': {'checksum': None,
                                                         'num_bytes': 298170}},
    download_size=392040,
    post_processing_size=None,
    dataset_size=920674,
    size_in_bytes=1312714)

```

```
[ ]: # Change the format as dataframe
banking77_preprocessed.set_format(type='pandas')
```

```
[ ]: # Define a dataframe for the training set
train_df = banking77_preprocessed['train'][:]
```

```
[ ]: # Check the dataframe of the training set
train_df
```

```
[ ]:
```

	text	label
0	still waiting card	11
1	card still hasnt arrived weeks	11
2	waiting week card still coming	11
3	track card process delivery	11
4	know get card lost	11
...
9998	provide support countries	24
9999	countries supporting	24
10000	countries getting support	24
10001	cards available eu	24
10002	countries represented	24

[10003 rows x 2 columns]

```
[ ]: # Check if there is null data
train_df.isnull().sum()
```

```
[ ]: text      0
label      0
dtype: int64
```

```
[ ]: train_df
```

```
[ ]:
```

	text	label
0	still waiting card	11
1	card still hasnt arrived weeks	11
2	waiting week card still coming	11
3	track card process delivery	11
4	know get card lost	11
...
9998	provide support countries	24
9999	countries supporting	24
10000	countries getting support	24
10001	cards available eu	24
10002	countries represented	24

[10003 rows x 2 columns]

```
[ ]: # Check if there is the same text after text preprocessing
# Group by text and label, then count occurrences
train_df['count'] = train_df.groupby(['text', 'label'])['text'].
    ↪transform('count')

# Filter rows where count > 1 (duplicates)
duplicates_df = train_df[train_df['count'] > 1]

# Sort it by text to group duplicates together
duplicates_df = duplicates_df.sort_values(by='text')
```

```
[ ]: duplicates_df[['text', 'label']]
```

```
[ ]:
      text  label
5028  able get cash atm    26
5011  able get cash atm    26
4178  able get visa mastercard    73
4135  able get visa mastercard    73
5796  able see refund    51
...
5460  would like cancel purchase    52
31    would like track card sent    11
98    would like track card sent    11
4949  wouldnt atm give money    26
5024  wouldnt atm give money    26
```

[1298 rows x 2 columns]

```
[ ]: # Remove duplicates based on text, keeping the first one
train_df = train_df.drop_duplicates(subset='text', keep='first')
```

```
[ ]: train_df
```


```
[ ]:
      text  label  count
0      still waiting card    11    2
1  card still hasnt arrived weeks    11    1
2  waiting week card still coming    11    1
3    track card process delivery    11    1
4      know get card lost    11    1
...
9997  moved us get card    24    1
9998  provide support countries    24    1
9999  countries supporting    24    1
10000  countries getting support    24    1
10002  countries represented    24    1
```

[9183 rows x 3 columns]

```
[ ]: # Check the text length in the training set
lengths = np.array([len(text) for text in train_df['text']])
print(f'The average is {np.mean(lengths)}. The median is {np.median(lengths)}.
↳The max length is {np.max(lengths)}.')

# Plot the histogram
plt.hist(lengths)
plt.xlabel('Text length in the training set')
plt.ylabel('Frequency')
plt.show()
```

The average is 36.17641293694871. The median is 29.0. The max length is 303.



Since the median is 29, a length of 29 will be used. The rest will be cut or filled in with padding.

```
[ ]: # Define a dataframe for the test set
test_df = banking77_preprocessed['test'][:]
```

```
[ ]: # Check the dataframe of the test set
test_df
```

```
[ ]:
      text  label
0      locate card    11
1  still received new card ordered week ago    11
2      ordered card arrived help please    11
3      way know card arrive    11
4      card arrived yet    11
...
3075      im uk still get card    24
3076      many countries support    24
3077      countries business    24
3078      countries operate    24
3079      card mailed used europe    24
```

[3080 rows x 2 columns]

```
[ ]: # Check if there is null data
test_df.isnull().sum()
```

```
[ ]: text    0
      label  0
      dtype: int64
```

```
[ ]: # Check the data type of elements in the dataframes
print(train_df['text'].dtype)
print(test_df['text'].dtype)
print(train_df['label'].dtype)
print(test_df['label'].dtype)
```

```
object
object
int64
int64
```

```
[ ]: # Check the number of labels in the training set
train_df['label'].nunique()
```


```
[ ]: 77
```

```
[ ]: # Check the balance of labels in the training set
train_df['label'].value_counts()
```

```
[ ]: label
6      181
28     178
75     178
15     176
66     168
```

```
...
41      69
18      56
10      48
72      38
23      33
Name: count, Length: 77, dtype: int64
```

```
[ ]: # Plot histogram for labels in the training set
plt.figure(figsize=(12, 8))
train_df['label'].value_counts().plot(kind='bar')
plt.xlabel('Label')
plt.ylabel('Count')
plt.title('Distribution of Labels')
plt.xticks(rotation=90)
plt.show()
```



The class in the training set is imbalanced.


```
[ ]: # Check the number of labels in the test set
test_df['label'].nunique()
```

```
[ ]: 77
```

```
[ ]: # Check the balance of labels in the test set  
test_df['label'].value_counts()
```

```
[ ]: label  
11    40  
63    40  
50    40  
64    40  
7     40  
     ..  
61    40  
10    40  
40    40  
54    40  
24    40  
Name: count, Length: 77, dtype: int64
```

```
[ ]: # Plot histogram for labels in the test set  
plt.figure(figsize=(12, 8))  
test_df['label'].value_counts().plot(kind='bar')  
plt.xlabel('Label')  
plt.ylabel('Count')  
plt.title('Distribution of Labels')  
plt.xticks(rotation=90)  
plt.show()
```



The class in the test set is balanced.

1.3 Tokenize the text in the dataset

1.3.1 For LSTM Model

```
[ ]: # Separate the training set to a training set and a validation set
# Import the library
from sklearn.model_selection import train_test_split

# Split the data to X and y
X = train_df['text']
y = train_df['label']

# Split the training set (80% for the training set, 20% for the validation set)
X_train, X_val, y_train, y_val = train_test_split(X, y, test_size=0.2,
↪stratify=train_df['label']) # Balance the class

[ ]: # Reset the index
X_train.reset_index(drop=True, inplace=True)
X_val.reset_index(drop=True, inplace=True)
y_train.reset_index(drop=True, inplace=True)
y_val.reset_index(drop=True, inplace=True)
```



```
[ ]: # Define X_test and y_test from the test set
X_test = test_df['text']
y_test = test_df['label']
```

```
[ ]: # Count the words to find the vocabulary size in the training set
# Define the function to cumulate the number of words
def count_word(X_train):
    count = Counter()
    for text in X_train.values:
        for word in text.split():
            count[word] += 1
    return count

# Count the number of words
word_counts = count_word(X_train)

print(f'The vocabulary size in the training set is {len(word_counts)}')
```

The vocabulary size in the training set is 2088

```
[ ]: # Check the inside of the vocabulary
word_counts
```

```
[ ]: Counter({'transaction': 256,
             'reverted': 49,
             'card': 1942,
             'charged': 401,
             'payment': 551,
             'made': 285,
             'cash': 518,
             'withdrawal': 212,
             'atm': 353,
             'listed': 12,
             'dont': 259,
             'remember': 28,
             'making': 40,
             'unauthorized': 11,
             'direct': 76,
             'debit': 106,
             'account': 1059,
             'please': 394,
             'explain': 61,
             'fee': 333,
             'transfer': 807,
             'verify': 95,
             'new': 256,
             'tried': 199,
```

'sending': 13,
'standard': 9,
'five': 5,
'times': 102,
'hasnt': 134,
'gone': 66,
'problem': 69,
'get': 606,
'refund': 213,
'purchase': 112,
'cost': 57,
'time': 122,
'frame': 9,
'getting': 122,
'identity': 152,
'verification': 85,
'isnt': 130,
'working': 145,
'could': 96,
'google': 41,
'pay': 133,
'top': 436,
'together': 2,
'expires': 22,
'soon': 39,
'fast': 13,
'replacement': 10,
'sent': 107,
'costs': 7,
'accidentally': 9,
'chose': 3,
'exchange': 411,
'gbp': 47,
'need': 544,
'pick': 3,
'aud': 16,
'change': 151,
'documents': 12,
'validate': 1,
'fiat': 26,
'currencies': 147,
'supported': 25,
'holding': 19,
'decline': 17,
'incoming': 1,
'next': 13,
'maybe': 8,

'let': 60,
'people': 16,
'know': 230,
'buying': 9,
'things': 21,
'another': 79,
'country': 48,
'extra': 205,
'used': 97,
'item': 72,
'also': 17,
'looks': 26,
'like': 214,
'wasnt': 49,
'aware': 11,
'think': 127,
'updated': 24,
'depositing': 7,
'cheque': 84,
'balance': 89,
'multiple': 44,
'one': 206,
'received': 181,
'incorrect': 43,
'amount': 161,
'today': 76,
'wanted': 42,
'virtual': 161,
'located': 8,
'yesterday': 52,
'topped': 24,
'didnt': 257,
'complete': 50,
'still': 273,
'pending': 304,
'processed': 14,
'cards': 214,
'shipped': 5,
'recently': 39,
'would': 274,
'cancel': 102,
'longer': 26,
'want': 213,
'us': 67,
'arrive': 24,
'assist': 19,
'completing': 4,

'im': 225,
'trying': 72,
'keep': 59,
'error': 43,
'flat': 12,
'initial': 1,
'mortgage': 7,
'find': 104,
'whats': 143,
'happening': 25,
'able': 99,
'use': 310,
'american': 35,
'express': 38,
'add': 84,
'money': 895,
'transferring': 3,
'havent': 83,
'funds': 142,
'yet': 202,
'choose': 14,
'visa': 70,
'mastercard': 77,
'order': 91,
'cant': 175,
'topup': 283,
'declined': 210,
'app': 354,
'shows': 76,
'fraudulent': 9,
'source': 30,
'long': 275,
'uk': 50,
'usually': 14,
'take': 260,
'showing': 159,
'everything': 36,
'actually': 29,
'went': 71,
'okay': 18,
'work': 169,
'shop': 7,
'physical': 59,
'delivered': 36,
'seller': 56,
'told': 11,
'paid': 41,

'reviewing': 4,
'subtract': 1,
'added': 29,
'back': 138,
'care': 3,
'promptly': 1,
'length': 3,
'credit': 75,
'transactions': 70,
'ive': 134,
'category': 1,
'days': 91,
'wont': 68,
'move': 12,
'make': 256,
'go': 198,
'checked': 47,
'receipts': 3,
'correct': 63,
'indicated': 1,
'previous': 5,
'email': 3,
'purchasing': 6,
'noticed': 55,
'rate': 257,
'reason': 73,
'needs': 32,
'verifying': 27,
'personal': 19,
'details': 69,
'got': 192,
'mugged': 8,
'took': 25,
'help': 282,
'steps': 39,
'check': 136,
'code': 43,
'topping': 59,
'pin': 243,
'number': 61,
'machine': 44,
'necessary': 6,
'gave': 29,
'pounds': 28,
'instead': 19,
'requested': 50,
'appear': 16,

'instant': 5,
'statement': 153,
'reflect': 9,
'issued': 9,
'receive': 95,
'access': 50,
'bank': 144,
'old': 30,
'ups': 16,
'done': 45,
'apple': 50,
'watch': 18,
'twice': 50,
'returned': 23,
'something': 169,
'store': 23,
'see': 220,
'appeared': 5,
'exchanging': 26,
'feature': 5,
'short': 2,
'many': 78,
'going': 155,
'wrong': 252,
'anymore': 16,
'possible': 124,
'someone': 131,
'else': 24,
'activate': 67,
'bought': 85,
'delete': 32,
'happy': 8,
'service': 29,
'youre': 6,
'providing': 2,
'smart': 3,
'phone': 73,
'lost': 79,
'stolen': 68,
'vacation': 12,
'immediately': 34,
'hold': 30,
'kids': 11,
'daughter': 28,
'passcode': 47,
'completely': 4,
'slipped': 1,

'mind': 6,
'hi': 56,
'first': 33,
'customer': 18,
'shown': 29,
'last': 54,
'half': 4,
'hour': 14,
'fix': 30,
'assistance': 14,
'really': 52,
'frustrating': 2,
'trouble': 13,
'id': 65,
'deposited': 65,
'mistake': 30,
'process': 62,
'changing': 6,
'much': 110,
'travel': 7,
'abroad': 42,
'ordering': 6,
'disposable': 149,
'methods': 5,
'accepted': 30,
'transferred': 66,
'morning': 59,
'fees': 100,
'using': 175,
'european': 24,
'hidden': 5,
'cancelled': 54,
'typically': 7,
'friend': 55,
'earlier': 57,
'even': 31,
'though': 36,
'hours': 20,
'ago': 139,
'documentation': 2,
'children': 10,
'available': 55,
'found': 40,
'put': 58,
'come': 79,
'yall': 1,
'support': 41,

'small': 5,
'needed': 21,
'applied': 44,
'recipient': 16,
'submit': 4,
'couple': 74,
'continues': 1,
'unsuccessful': 6,
'may': 29,
'status': 24,
'tell': 189,
'failed': 41,
'withdraw': 89,
'tracked': 2,
'hello': 49,
'seems': 58,
'never': 50,
'seen': 23,
'refunds': 4,
'allowed': 40,
'certain': 30,
'items': 12,
'recognize': 51,
'left': 13,
'parents': 1,
'house': 6,
'send': 59,
'spare': 3,
'meantime': 1,
'delivering': 1,
'update': 30,
'sepa': 17,
'exchanges': 30,
'moved': 17,
'acceptable': 6,
'deposit': 77,
'denied': 12,
'option': 26,
'selected': 3,
'beneficiary': 51,
'saw': 35,
'merchant': 26,
'stay': 5,
'keeps': 35,
'declining': 16,
'two': 42,
'different': 73,

'atms': 39,
'already': 73,
'alright': 7,
'notting': 9,
'hill': 9,
'must': 23,
'pound': 36,
'currency': 156,
'ways': 2,
'paying': 25,
'choice': 2,
'reader': 1,
'accept': 57,
'foreign': 55,
'weekdays': 2,
'weekends': 2,
'completed': 29,
'waiting': 59,
'way': 80,
'somewhere': 6,
'form': 7,
'payments': 73,
'prove': 16,
'wallet': 25,
'x': 6,
'approved': 9,
'separate': 3,
'worried': 12,
'post': 13,
'address': 21,
'information': 56,
'contactless': 21,
'function': 9,
'arrived': 44,
'denying': 2,
'transfers': 110,
'horrible': 2,
'request': 21,
'salary': 28,
'might': 22,
'strange': 24,
'restaurant': 19,
'wouldnt': 18,
'taking': 37,
'asked': 34,
'pretty': 12,
'slow': 2,

'increased': 5,
'addition': 2,
'happened': 50,
'downsides': 1,
'reasons': 12,
'fail': 11,
'refills': 1,
'traveling': 19,
'somebody': 9,
'seeing': 21,
'adding': 17,
'missing': 30,
'copy': 4,
'£': 24,
'countries': 39,
'often': 7,
'worked': 28,
'places': 10,
'inform': 5,
'restrictions': 17,
'right': 88,
'away': 13,
'locations': 10,
'higher': 9,
'europe': 23,
'cannot': 26,
'dispute': 20,
'verified': 24,
'per': 13,
'person': 16,
'offer': 23,
'discount': 10,
'since': 50,
'frequently': 10,
'package': 3,
'auto': 45,
'limit': 54,
'anything': 37,
'mines': 2,
'expire': 24,
'takes': 9,
'retract': 1,
'china': 42,
'urgent': 19,
'sure': 75,
'plenty': 1,
'couldnt': 21,

'week': 55,
'view': 2,
'history': 6,
'came': 38,
'accounts': 11,
'track': 22,
'wait': 44,
'rejected': 31,
'readily': 1,
'give': 64,
'dollar': 10,
'ready': 3,
'enjoy': 1,
'tutor': 1,
'session': 1,
'withdrawal': 8,
'skip': 1,
'international': 40,
'france': 10,
'forewarned': 1,
'stop': 25,
'accessing': 2,
'eu': 22,
'online': 53,
'says': 52,
'interested': 11,
'rates': 39,
'based': 1,
'live': 20,
'assessed': 2,
'forgot': 14,
'trip': 5,
'overseas': 12,
'unblock': 22,
'blocked': 22,
'giving': 7,
'message': 39,
'device': 5,
'log': 5,
'onto': 3,
'platform': 1,
'needing': 3,
'doesnt': 61,
'seem': 42,
'wondering': 16,
'several': 37,
'charges': 69,

'charge': 234,
'mine': 37,
'expiring': 8,
'avoid': 3,
'future': 4,
'checking': 19,
'authorized': 3,
'reversed': 13,
'replaced': 3,
'revert': 15,
'second': 34,
'canceled': 12,
'automatically': 26,
'convert': 7,
'company': 17,
'awful': 1,
'procedure': 10,
'happen': 12,
'quick': 5,
'expedition': 1,
'crucial': 1,
'approximately': 4,
'specify': 2,
'day': 48,
'show': 84,
'orders': 1,
'less': 25,
'activity': 5,
'withdrawals': 24,
'always': 13,
'thought': 57,
'free': 49,
'stopped': 7,
'product': 11,
'ordered': 29,
'unknown': 1,
'additional': 42,
'prior': 2,
'notification': 2,
'sorts': 1,
'required': 12,
'explained': 1,
'figure': 19,
'maximum': 11,
'via': 11,
'seconds': 3,
'hit': 5,

'follow': 4,
'directions': 3,
'calculated': 4,
'reach': 10,
'start': 11,
'list': 9,
'saying': 16,
'anticipated': 1,
'password': 20,
'reset': 16,
'unknowingly': 1,
'redeposied': 1,
'resolve': 10,
'estimated': 3,
'include': 1,
'latest': 4,
'entered': 19,
'unblocked': 7,
'theres': 33,
'somehow': 1,
'average': 2,
'shouldnt': 25,
'name': 32,
'married': 6,
'look': 35,
'owe': 1,
'beneficiary': 2,
'compromised': 11,
'swift': 15,
'open': 35,
'offspring': 1,
'ideal': 1,
'understand': 27,
'portugal': 1,
'operate': 3,
'including': 1,
'past': 20,
'month': 13,
'place': 27,
'good': 12,
'suddenly': 17,
'rewarded': 3,
'quickly': 16,
'type': 13,
'cleared': 8,
'withdrawn': 14,
'holiday': 15,

'guess': 5,
'overcharged': 11,
'withdrew': 32,
'require': 6,
'unfamiliar': 6,
'calling': 1,
'actual': 11,
'period': 3,
'swallow': 1,
'marriage': 2,
'three': 3,
'months': 11,
'cat': 2,
'fluffy': 3,
'unfortunately': 3,
'fan': 1,
'pet': 1,
'said': 15,
'nothing': 15,
'cause': 12,
'failure': 2,
'talk': 7,
'debited': 1,
'contacted': 14,
'directed': 1,
'issue': 41,
'within': 12,
'definitely': 22,
'fully': 2,
'phase': 1,
'normal': 10,
'turn': 4,
'around': 5,
'expect': 28,
'link': 31,
'friends': 20,
'double': 32,
'home': 13,
'cool': 2,
'saturday': 12,
'accident': 1,
'digits': 2,
'fixing': 2,
'purchases': 19,
'authorize': 11,
'quite': 8,
'idea': 10,

'freeze': 21,
'utilizing': 1,
'without': 29,
'authorization': 1,
'previously': 6,
'attempt': 4,
'remove': 14,
'clearly': 6,
'issuing': 1,
'frames': 1,
'recognise': 11,
'particular': 1,
'coming': 19,
'specifically': 2,
'ask': 7,
'although': 6,
'pass': 2,
'acount': 1,
'currently': 7,
'autotop': 8,
'close': 20,
'determined': 2,
'changed': 35,
'remotely': 2,
'visiting': 1,
'purchased': 41,
'low': 6,
'shopping': 16,
'tickets': 2,
'age': 24,
'child': 7,
'services': 17,
'limitations': 3,
'finishing': 1,
'directly': 9,
'matter': 7,
'bit': 12,
'concern': 1,
'notice': 11,
'weeks': 37,
'trace': 8,
'truly': 3,
'unable': 40,
'crypto': 20,
'keen': 1,
'buy': 24,
'applicationi': 1,

'thats': 13,
'expired': 8,
'kind': 13,
'looking': 23,
'price': 5,
'offered': 5,
'disappeared': 18,
'urgency': 1,
'odd': 8,
'remote': 4,
'town': 8,
'reasonable': 1,
'hundred': 2,
'realize': 5,
'single': 5,
'visit': 5,
'removed': 6,
'mistakenly': 2,
'taken': 19,
'occasions': 1,
'elaborate': 2,
'upon': 3,
'replace': 8,
'step': 3,
'businesses': 7,
'limits': 16,
'transferring': 35,
'wheres': 9,
'assess': 1,
'try': 16,
'unlimited': 2,
'reaches': 3,
'forms': 4,
'types': 10,
'disappointed': 4,
'bad': 10,
'hope': 5,
'confirm': 6,
'official': 8,
'interbank': 11,
'unknown': 4,
'submitted': 5,
'query': 1,
'thing': 18,
'receiving': 20,
'separably': 1,
'brand': 2,

'activation': 12,
'failing': 11,
'fraud': 8,
'protect': 1,
'sale': 1,
'transferis': 1,
'expiration': 7,
'near': 3,
'pull': 2,
'unusual': 5,
'believe': 26,
'didny': 1,
'retrieve': 2,
'looked': 13,
'broken': 15,
'given': 12,
'data': 2,
'exposed': 1,
'stole': 15,
'banks': 2,
'happens': 14,
'stipulations': 1,
'entail': 1,
'asap': 18,
'question': 5,
'concerning': 1,
'guys': 18,
'accurate': 5,
'unknown': 1,
'origin': 1,
'sense': 3,
'ok': 11,
'wish': 5,
'select': 2,
'preference': 2,
'recover': 2,
'swallowed': 6,
'stuck': 16,
'changes': 5,
'supplier': 1,
'cardsper': 1,
'problems': 12,
'beneficiarys': 1,
'credited': 7,
'difference': 3,
'union': 2,
'full': 14,

'banking': 3,
'mean': 21,
'little': 6,
'bill': 7,
'tired': 3,
'whatwhen': 1,
'set': 35,
'displaying': 3,
'report': 9,
'correctly': 8,
'standing': 5,
'telling': 4,
'austria': 5,
'theretotop': 1,
'exchanged': 9,
'anywhere': 22,
'l': 2,
'return': 20,
'advise': 7,
'performed': 3,
'print': 1,
'date': 18,
'unhappy': 3,
'cut': 2,
'waited': 1,
'retailers': 3,
'activated': 9,
'messed': 5,
'recent': 27,
'goes': 7,
'login': 3,
'point': 11,
'deliveries': 1,
'weird': 4,
'recognizing': 4,
'united': 6,
'states': 9,
'appreciate': 3,
'siphoned': 1,
'helping': 1,
'japan': 1,
'decorations': 1,
'local': 3,
'stores': 4,
'differed': 1,
'greatly': 1,
'showed': 5,

'hacked': 3,
'policy': 10,
'end': 10,
'gonna': 1,
'enabled': 1,
'living': 4,
'ship': 2,
'sort': 5,
'whenever': 1,
'allow': 20,
'usd': 15,
'hate': 2,
'exactly': 3,
'withdrawals': 1,
'progressseems': 1,
'wrongi': 1,
'memberships': 1,
'appears': 25,
'placed': 8,
'refunded': 15,
'ended': 2,
'info': 22,
'hiaccordingly': 1,
'accountwhat': 1,
'reverse': 11,
'apply': 8,
'reactivate': 17,
'turns': 3,
'jacket': 9,
'pocket': 3,
'withdraws': 3,
'tranfering': 1,
'dispose': 1,
'reactive': 2,
'continue': 3,
'statements': 11,
'upi': 1,
'submitting': 1,
'refilled': 2,
'attempted': 21,
'enter': 4,
'russian': 9,
'ruble': 7,
'wen': 1,
'forget': 2,
'eaten': 1,
'patentp': 1,

'proof': 6,
'mail': 15,
'deliver': 7,
'rent': 16,
'side': 8,
'travelling': 6,
'euros': 3,
'drunk': 2,
'mistyped': 1,
'restored': 1,
'seemingly': 2,
'randomly': 2,
'system': 28,
'complicated': 1,
'til': 1,
'severely': 1,
'known': 3,
'requirements': 5,
'opening': 2,
'call': 2,
'locate': 6,
'marital': 1,
'learning': 2,
'supplementary': 1,
'enough': 15,
'unsure': 5,
'topups': 27,
'gotten': 16,
'canceling': 4,
'products': 3,
'realized': 4,
'correcti': 3,
'tomorrow': 10,
'delivery': 25,
'solve': 6,
'discounts': 5,
'urgently': 16,
'duplicated': 7,
'goofed': 1,
'ate': 6,
'identify': 7,
'important': 10,
'paranoid': 1,
'factor': 1,
'configured': 2,
'hey': 12,
'theyre': 4,

'forever': 9,
'difficult': 1,
'electronic': 2,
'suffice': 1,
'confusing': 2,
'confused': 6,
'ass': 1,
'size': 2,
'supposed': 18,
'big': 3,
'regretted': 1,
'cover': 3,
'cast': 1,
'works': 13,
'bring': 1,
'identification': 12,
'swap': 1,
'cloned': 1,
'increase': 3,
'rather': 5,
'decreasing': 1,
'customers': 3,
'incentives': 1,
'youngest': 1,
'simply': 2,
'greece': 1,
'course': 1,
'investigate': 2,
'pulled': 3,
'brought': 1,
'stuff': 1,
'suspend': 1,
'expected': 10,
'recipients': 2,
'machines': 13,
'kept': 9,
'normally': 13,
'confirmed': 6,
'countrys': 3,
'pulling': 3,
'instructions': 4,
'activating': 19,
'manually': 3,
'hotel': 8,
'ran': 2,
'causing': 6,
'thanks': 6,

```

'institutions': 1,
'far': 12,
'obtain': 5,
'possibility': 3,
'faster': 7,
'warned': 4,
'explanation': 1,
'hard': 4,
'application': 2,
'cheque': 2,
'authorise': 3,
'arrival': 3,
'specific': 10,
'payee': 1,
'courier': 1,
'cone': 1,
'experiencing': 2,
'experienced': 1,
'terminated': 2,
'sorted': 2,
'inquiring': 2,
'unlocked': 2,
'instantly': 2,
'edit': 11,
'intervals': 7,
'accepting': 4,
'intended': 1,
'tracking': 8,
'eur': 15,
'anyway': 3,
'setup': 4,
'automatic': 6,
'discounted': 1,
'finish': 5,
'unsuccessfully': 2,
'deducted': 8,
...})

```

```

[ ]: # Check the most frequent words
word_counts.most_common(5)

```

```

[ ]: [('card', 1942),
      ('account', 1059),
      ('money', 895),
      ('transfer', 807),
      ('get', 606)]

```

```
[ ]: # Store the size of the vocabulary
voca_size = len(word_counts)
print(voca_size)
```

2088

```
[ ]: # Change the X data to arrays
X_train_array = X_train.to_numpy()
X_val_array = X_val.to_numpy()
X_test_array = X_test.to_numpy()
```

```
[ ]: # Change the y data to arrays
y_train_array = y_train.to_numpy()
y_val_array = y_val.to_numpy()
y_test_array = y_test.to_numpy()
```

```
[ ]: # Check the dimension of the variables
print(X_train_array.shape)
print(X_val_array.shape)
print(X_test_array.shape)
print(y_train_array.shape)
print(y_val_array.shape)
print(y_test_array.shape)
```

(7346,)

(1837,)

(3080,)

(7346,)

(1837,)

(3080,)

```
[ ]: # Check the type of the variables
print(X_train_array.__class__)
print(X_val_array.__class__)
print(X_test_array.__class__)
print(y_train_array.__class__)
print(y_val_array.__class__)
print(y_test_array.__class__)
```

<class 'numpy.ndarray'>

<class 'numpy.ndarray'>

<class 'numpy.ndarray'>

<class 'numpy.ndarray'>

<class 'numpy.ndarray'>

<class 'numpy.ndarray'>

```
[ ]: # Import tokenizer
from tensorflow.keras.preprocessing.text import Tokenizer

# Define the tokenizer with vocabulary size of the training data
tokenizer = Tokenizer(num_words=voca_size)
# Fit the tokenizer with X train
tokenizer.fit_on_texts(X_train_array)
```

```
[ ]: # Define the word index
word_index = tokenizer.word_index
word_index
```

```
[ ]: {'card': 1,
      'account': 2,
      'money': 3,
      'transfer': 4,
      'get': 5,
      'payment': 6,
      'need': 7,
      'cash': 8,
      'top': 9,
      'exchange': 10,
      'charged': 11,
      'please': 12,
      'app': 13,
      'atm': 14,
      'fee': 15,
      'use': 16,
      'pending': 17,
      'made': 18,
      'topup': 19,
      'help': 20,
      'long': 21,
      'would': 22,
      'still': 23,
      'take': 24,
      'dont': 25,
      'didnt': 26,
      'rate': 27,
      'transaction': 28,
      'new': 29,
      'make': 30,
      'wrong': 31,
      'pin': 32,
      'charge': 33,
      'know': 34,
      'im': 35,
```


'see': 36,
'like': 37,
'cards': 38,
'refund': 39,
'want': 40,
'withdrawal': 41,
'declined': 42,
'one': 43,
'extra': 44,
'yet': 45,
'tried': 46,
'go': 47,
'got': 48,
'tell': 49,
'received': 50,
'cant': 51,
'using': 52,
'work': 53,
'something': 54,
'amount': 55,
'virtual': 56,
'showing': 57,
'currency': 58,
'going': 59,
'statement': 60,
'identity': 61,
'change': 62,
'disposable': 63,
'currencies': 64,
'working': 65,
'bank': 66,
'whats': 67,
'funds': 68,
'ago': 69,
'back': 70,
'check': 71,
'hasnt': 72,
'ive': 73,
'pay': 74,
'someone': 75,
'isnt': 76,
'think': 77,
'possible': 78,
'time': 79,
'getting': 80,
'purchase': 81,
'much': 82,

'transfers': 83,
'sent': 84,
'debit': 85,
'find': 86,
'times': 87,
'cancel': 88,
'fees': 89,
'able': 90,
'used': 91,
'could': 92,
'verify': 93,
'receive': 94,
'order': 95,
'days': 96,
'balance': 97,
'withdraw': 98,
'right': 99,
'verification': 100,
'bought': 101,
'cheque': 102,
'add': 103,
'show': 104,
'havent': 105,
'way': 106,
'another': 107,
'lost': 108,
'come': 109,
'many': 110,
'mastercard': 111,
'deposit': 112,
'direct': 113,
'today': 114,
'shows': 115,
'credit': 116,
'sure': 117,
'couple': 118,
'reason': 119,
'phone': 120,
'different': 121,
'already': 122,
'payments': 123,
'item': 124,
'trying': 125,
'went': 126,
'visa': 127,
'transactions': 128,
'problem': 129,

'details': 130,
'charges': 131,
'wont': 132,
'stolen': 133,
'us': 134,
'activate': 135,
'gone': 136,
'transferred': 137,
'id': 138,
'deposited': 139,
'give': 140,
'correct': 141,
'process': 142,
'explain': 143,
'number': 144,
'doesnt': 145,
'let': 146,
'keep': 147,
'physical': 148,
'topping': 149,
'morning': 150,
'send': 151,
'waiting': 152,
'put': 153,
'seems': 154,
'cost': 155,
'earlier': 156,
'accept': 157,
'thought': 158,
'seller': 159,
'hi': 160,
'information': 161,
'noticed': 162,
'friend': 163,
'available': 164,
'foreign': 165,
'week': 166,
'last': 167,
'cancelled': 168,
'limit': 169,
'online': 170,
'yesterday': 171,
'really': 172,
'says': 173,
'recognize': 174,
'beneficiary': 175,
'complete': 176,

'uk': 177,
'requested': 178,
'access': 179,
'apple': 180,
'twice': 181,
'never': 182,
'happened': 183,
'since': 184,
'reverted': 185,
'wasnt': 186,
'hello': 187,
'free': 188,
'country': 189,
'day': 190,
'gbp': 191,
'checked': 192,
'passcode': 193,
'done': 194,
'auto': 195,
'multiple': 196,
'machine': 197,
'applied': 198,
'arrived': 199,
'wait': 200,
'incorrect': 201,
'error': 202,
'code': 203,
'wanted': 204,
'abroad': 205,
'two': 206,
'china': 207,
'seem': 208,
'additional': 209,
'google': 210,
'paid': 211,
'support': 212,
'failed': 213,
'issue': 214,
'purchased': 215,
'making': 216,
'found': 217,
'allowed': 218,
'international': 219,
'unable': 220,
'soon': 221,
'recently': 222,
'steps': 223,

'atms': 224,
'countries': 225,
'rates': 226,
'message': 227,
'express': 228,
'came': 229,
'taking': 230,
'anything': 231,
'several': 232,
'mine': 233,
'weeks': 234,
'everything': 235,
'delivered': 236,
'though': 237,
'pound': 238,
'american': 239,
'saw': 240,
'keeps': 241,
'look': 242,
'open': 243,
'changed': 244,
'transferring': 245,
'set': 246,
'immediately': 247,
'asked': 248,
'second': 249,
'first': 250,
'theres': 251,
'needs': 252,
'delete': 253,
'name': 254,
'withdrew': 255,
'double': 256,
'even': 257,
'rejected': 258,
'link': 259,
'source': 260,
'old': 261,
'hold': 262,
'fix': 263,
'mistake': 264,
'accepted': 265,
'certain': 266,
'update': 267,
'exchanges': 268,
'missing': 269,
'actually': 270,

'added': 271,
'gave': 272,
'service': 273,
'shown': 274,
'may': 275,
'completed': 276,
'ordered': 277,
'without': 278,
'remember': 279,
'pounds': 280,
'daughter': 281,
'salary': 282,
'worked': 283,
'expect': 284,
'system': 285,
'verifying': 286,
'understand': 287,
'place': 288,
'recent': 289,
'topups': 290,
'fiat': 291,
'looks': 292,
'longer': 293,
'exchanging': 294,
'option': 295,
'merchant': 296,
'cannot': 297,
'automatically': 298,
'believe': 299,
'supported': 300,
'happening': 301,
'took': 302,
'paying': 303,
'wallet': 304,
'stop': 305,
'less': 306,
'shouldnt': 307,
'appears': 308,
'delivery': 309,
'updated': 310,
'topped': 311,
'arrive': 312,
'else': 313,
'european': 314,
'status': 315,
'strange': 316,
'£': 317,

'verified': 318,
'expire': 319,
'withdrawals': 320,
'age': 321,
'buy': 322,
'returned': 323,
'store': 324,
'seen': 325,
'must': 326,
'europe': 327,
'offer': 328,
'looking': 329,
'duplicate': 330,
'expires': 331,
'might': 332,
'track': 333,
'eu': 334,
'unblock': 335,
'blocked': 336,
'definitely': 337,
'anywhere': 338,
'info': 339,
'withdrawing': 340,
'things': 341,
'needed': 342,
'address': 343,
'contactless': 344,
'request': 345,
'seeing': 346,
'couldnt': 347,
'freeze': 348,
'mean': 349,
'attempted': 350,
'hours': 351,
'dispute': 352,
'live': 353,
'password': 354,
'past': 355,
'friends': 356,
'close': 357,
'crypto': 358,
'receiving': 359,
'return': 360,
'allow': 361,
'holding': 362,
'assist': 363,
'personal': 364,

'instead': 365,
'restaurant': 366,
'traveling': 367,
'urgent': 368,
'checking': 369,
'figure': 370,
'entered': 371,
'purchases': 372,
'coming': 373,
'taken': 374,
'activating': 375,
'okay': 376,
'watch': 377,
'customer': 378,
'wouldnt': 379,
'disappeared': 380,
'thing': 381,
'asap': 382,
'guys': 383,
'date': 384,
'supposed': 385,
'€': 386,
'provide': 387,
'decline': 388,
'also': 389,
'sepa': 390,
'moved': 391,
'adding': 392,
'restrictions': 393,
'company': 394,
'suddenly': 395,
'services': 396,
'reactivate': 397,
'aud': 398,
'people': 399,
'appear': 400,
'ups': 401,
'anymore': 402,
'recipient': 403,
'declining': 404,
'prove': 405,
'person': 406,
'wondering': 407,
'saying': 408,
'reset': 409,
'quickly': 410,
'shopping': 411,

'limits': 412,
'try': 413,
'stuck': 414,
'rent': 415,
'gotten': 416,
'urgently': 417,
'revert': 418,
'swift': 419,
'holiday': 420,
'said': 421,
'nothing': 422,
'broken': 423,
'stole': 424,
'usd': 425,
'refunded': 426,
'mail': 427,
'enough': 428,
'eur': 429,
'processed': 430,
'choose': 431,
'usually': 432,
'hour': 433,
'assistance': 434,
'forgot': 435,
'withdrawn': 436,
'contacted': 437,
'remove': 438,
'happens': 439,
'full': 440,
'sending': 441,
'fast': 442,
'next': 443,
'trouble': 444,
'left': 445,
'post': 446,
'away': 447,
'per': 448,
'reversed': 449,
'always': 450,
'month': 451,
'type': 452,
'home': 453,
'thats': 454,
'kind': 455,
'looked': 456,
'works': 457,
'machines': 458,

'normally': 459,
'fine': 460,
'clear': 461,
'receiver': 462,
'listed': 463,
'documents': 464,
'flat': 465,
'move': 466,
'vacation': 467,
'items': 468,
'denied': 469,
'worried': 470,
'pretty': 471,
'reasons': 472,
'overseas': 473,
'canceled': 474,
'happen': 475,
'required': 476,
'good': 477,
'cause': 478,
'within': 479,
'saturday': 480,
'bit': 481,
'activation': 482,
'given': 483,
'problems': 484,
'hey': 485,
'identification': 486,
'far': 487,
'deal': 488,
'however': 489,
'current': 490,
'well': 491,
'large': 492,
'unauthorized': 493,
'aware': 494,
'told': 495,
'kids': 496,
'fail': 497,
'accounts': 498,
'interested': 499,
'product': 500,
'maximum': 501,
'via': 502,
'start': 503,
'compromised': 504,
'overcharged': 505,

'actual': 506,
'months': 507,
'authorize': 508,
'recognise': 509,
'notice': 510,
'interbank': 511,
'failing': 512,
'ok': 513,
'point': 514,
'reverse': 515,
'statements': 516,
'edit': 517,
'checks': 518,
'charging': 519,
'replacement': 520,
'children': 521,
'places': 522,
'locations': 523,
'discount': 524,
'frequently': 525,
'dollar': 526,
'france': 527,
'procedure': 528,
'reach': 529,
'resolve': 530,
'normal': 531,
'idea': 532,
'types': 533,
'bad': 534,
'policy': 535,
'end': 536,
'tomorrow': 537,
'important': 538,
'expected': 539,
'specific': 540,
'create': 541,
'options': 542,
'standard': 543,
'frame': 544,
'accidentally': 545,
'buying': 546,
'fraudulent': 547,
'reflect': 548,
'issued': 549,
'notting': 550,
'hill': 551,
'approved': 552,

'function': 553,
'somebody': 554,
'higher': 555,
'takes': 556,
'list': 557,
'directly': 558,
'wheres': 559,
'report': 560,
'exchanged': 561,
'activated': 562,
'states': 563,
'jacket': 564,
'russian': 565,
'forever': 566,
'kept': 567,
'issues': 568,
'terminate': 569,
'gas': 570,
'progress': 571,
'landlord': 572,
'high': 573,
'due': 574,
'contact': 575,
'maybe': 576,
'located': 577,
'mugged': 578,
'happy': 579,
'withdrawl': 580,
'expiring': 581,
'cleared': 582,
'quite': 583,
'autotop': 584,
'trace': 585,
'expired': 586,
'odd': 587,
'town': 588,
'replace': 589,
'official': 590,
'fraud': 591,
'correctly': 592,
'placed': 593,
'apply': 594,
'side': 595,
'hotel': 596,
'tracking': 597,
'deducted': 598,
'ones': 599,

'every': 600,
'arent': 601,
'costs': 602,
'depositing': 603,
'mortgage': 604,
'shop': 605,
'travel': 606,
'typically': 607,
'alright': 608,
'form': 609,
'often': 610,
'giving': 611,
'convert': 612,
'stopped': 613,
'unblocked': 614,
'talk': 615,
'ask': 616,
'currently': 617,
'child': 618,
'matter': 619,
'businesses': 620,
'expiration': 621,
'credited': 622,
'bill': 623,
'advise': 624,
'goes': 625,
'ruble': 626,
'deliver': 627,
'duplicated': 628,
'identify': 629,
'faster': 630,
'intervals': 631,
'play': 632,
'impression': 633,
'recall': 634,
'spain': 635,
'years': 636,
'concerned': 637,
'outside': 638,
'area': 639,
'say': 640,
'locked': 641,
'reflected': 642,
'posted': 643,
'comes': 644,
'purchasing': 645,
'necessary': 646,

'youre': 647,
'mind': 648,
'changing': 649,
'ordering': 650,
'unsuccessful': 651,
'house': 652,
'acceptable': 653,
'somewhere': 654,
'x': 655,
'history': 656,
'married': 657,
'require': 658,
'unfamiliar': 659,
'previously': 660,
'clearly': 661,
'although': 662,
'low': 663,
'removed': 664,
'confirm': 665,
'swallowed': 666,
'little': 667,
'united': 668,
'proof': 669,
'travelling': 670,
'locate': 671,
'solve': 672,
'ate': 673,
'confused': 674,
'confirmed': 675,
'causing': 676,
'thanks': 677,
'automatic': 678,
'entering': 679,
'block': 680,
'putting': 681,
'frequent': 682,
'car': 683,
'permission': 684,
'proving': 685,
'minutes': 686,
'lot': 687,
'feel': 688,
'letting': 689,
'fixed': 690,
'security': 691,
'results': 692,
'typo': 693,

'regarding': 694,
'incorrectly': 695,
'location': 696,
'resolved': 697,
'five': 698,
'shipped': 699,
'previous': 700,
'instant': 701,
'appeared': 702,
'feature': 703,
'methods': 704,
'hidden': 705,
'small': 706,
'stay': 707,
'increased': 708,
'inform': 709,
'trip': 710,
'device': 711,
'log': 712,
'quick': 713,
'activity': 714,
'hit': 715,
'guess': 716,
'around': 717,
'price': 718,
'offered': 719,
'realize': 720,
'single': 721,
'visit': 722,
'hope': 723,
'submitted': 724,
'unusual': 725,
'question': 726,
'accurate': 727,
'wish': 728,
'changes': 729,
'standing': 730,
'austria': 731,
'messed': 732,
'showed': 733,
'sort': 734,
'requirements': 735,
'unsure': 736,
'discounts': 737,
'rather': 738,
'obtain': 739,
'finish': 740,

'extremely': 741,
'satisfied': 742,
'passed': 743,
'modify': 744,
'appearing': 745,
'finding': 746,
'grocery': 747,
'perform': 748,
'gets': 749,
'attempts': 750,
'photos': 751,
'family': 752,
'forgotten': 753,
'associated': 754,
'lesser': 755,
'random': 756,
'thru': 757,
'tries': 758,
'terrible': 759,
'groceries': 760,
'rejecting': 761,
'frozen': 762,
'prefer': 763,
'heard': 764,
'linked': 765,
'refused': 766,
'reported': 767,
'completing': 768,
'reviewing': 769,
'completely': 770,
'half': 771,
'submit': 772,
'refunds': 773,
'copy': 774,
'future': 775,
'approximately': 776,
'follow': 777,
'calculated': 778,
'latest': 779,
'turn': 780,
'attempt': 781,
'remote': 782,
'forms': 783,
'disappointed': 784,
'unknown': 785,
'telling': 786,
'weird': 787,

'recognizing': 788,
'stores': 789,
'living': 790,
'enter': 791,
'canceling': 792,
'realized': 793,
'theyre': 794,
'instructions': 795,
'warned': 796,
'hard': 797,
'accepting': 798,
'setup': 799,
'dollars': 800,
'possibly': 801,
'deactivate': 802,
'usa': 803,
'spending': 804,
'refuse': 805,
'alternatives': 806,
'images': 807,
'readable': 808,
'economic': 809,
'switzerland': 810,
'properly': 811,
'ahead': 812,
'fund': 813,
'initiated': 814,
'closest': 815,
'following': 816,
'inside': 817,
'policies': 818,
'along': 819,
'mailed': 820,
'best': 821,
'son': 822,
'unlock': 823,
'case': 824,
'business': 825,
'resident': 826,
'rules': 827,
'attempting': 828,
'hoping': 829,
'better': 830,
'billing': 831,
'apparently': 832,
'finished': 833,
'receipt': 834,

'older': 835,
'recieve': 836,
'caused': 837,
'rubles': 838,
'stays': 839,
'expecting': 840,
'chose': 841,
'pick': 842,
'transferring': 843,
'care': 844,
'length': 845,
'receipts': 846,
'email': 847,
'smart': 848,
'spare': 849,
'selected': 850,
'separate': 851,
'package': 852,
'ready': 853,
'onto': 854,
'needing': 855,
'avoid': 856,
'authorized': 857,
'replaced': 858,
'seconds': 859,
'directions': 860,
'estimated': 861,
'operate': 862,
'rewarded': 863,
'period': 864,
'three': 865,
'fluffy': 866,
'unfortunately': 867,
'limitations': 868,
'truly': 869,
'upon': 870,
'step': 871,
'reaches': 872,
'near': 873,
'sense': 874,
'difference': 875,
'banking': 876,
'tired': 877,
'displaying': 878,
'performed': 879,
'unhappy': 880,
'retailers': 881,

'login': 882,
'appreciate': 883,
'local': 884,
'hacked': 885,
'exactly': 886,
'turns': 887,
'pocket': 888,
'withdraws': 889,
'continue': 890,
'euros': 891,
'known': 892,
'products': 893,
'correcti': 894,
'big': 895,
'cover': 896,
'increase': 897,
'customers': 898,
'pulled': 899,
'countrys': 900,
'pulling': 901,
'manually': 902,
'possibility': 903,
'authorise': 904,
'arrival': 905,
'anyway': 906,
'facing': 907,
'running': 908,
'rest': 909,
'limited': 910,
'finally': 911,
'anyone': 912,
'thinking': 913,
'suspicious': 914,
'cancellation': 915,
'provided': 916,
'gives': 917,
'forward': 918,
'sad': 919,
'decided': 920,
'eligible': 921,
'broke': 922,
'obtained': 923,
'upset': 924,
'false': 925,
'suppose': 926,
'almost': 927,
'visited': 928,

'explaining': 929,
'uses': 930,
'accessed': 931,
'numbers': 932,
'completion': 933,
'frustrated': 934,
'someones': 935,
'metro': 936,
'active': 937,
'bag': 938,
'regards': 939,
'minimum': 940,
'victim': 941,
'empty': 942,
'reside': 943,
'structure': 944,
'others': 945,
'thank': 946,
'fair': 947,
'method': 948,
'employer': 949,
'review': 950,
'according': 951,
'benow': 952,
'draw': 953,
'wondered': 954,
'separately': 955,
'settings': 956,
'spend': 957,
'realised': 958,
'run': 959,
'speak': 960,
'accepts': 961,
'guide': 962,
'nearly': 963,
'successfully': 964,
'real': 965,
'cancelling': 966,
'successful': 967,
'despite': 968,
'manage': 969,
'yo': 970,
'youve': 971,
'updating': 972,
'stating': 973,
'later': 974,
'max': 975,

```
'carry': 976,
'processing': 977,
'night': 978,
'accidentally': 979,
'typed': 980,
'simple': 981,
'havnt': 982,
'vanished': 983,
'raise': 984,
'city': 985,
'certainly': 986,
'nowhere': 987,
'functioning': 988,
'started': 989,
'doublecheck': 990,
'glare': 991,
'cheques': 992,
'reached': 993,
'member': 994,
'earth': 995,
'toppedup': 996,
'residents': 997,
'asking': 998,
'specified': 999,
'lookup': 1000,
...}
```

```
[ ]: # Define the maximum length as 29 (the median length of the training set)
max_length_train_text = 29
max_length_train_text
```

```
[ ]: 29
```

```
[ ]: # Encode the text data
X_train_sequences = tokenizer.texts_to_sequences(X_train_array)
X_val_sequences = tokenizer.texts_to_sequences(X_val_array)
X_test_sequences = tokenizer.texts_to_sequences(X_test_array)
```

```
[ ]: # Add padding to the encoded data
# Import pad_sequences
from tensorflow.keras.preprocessing.sequence import pad_sequences

# Padding and truncation will be added to post-texts
X_train_padded = pad_sequences(X_train_sequences, maxlen=max_length_train_text,
    ↳padding="post", truncating="post")
X_val_padded = pad_sequences(X_val_sequences, maxlen=max_length_train_text,
    ↳padding="post", truncating="post")
```

```
X_test_padded = pad_sequences(X_test_sequences, maxlen=max_length_train_text,
                               padding="post", truncating="post")
```

```
[ ]: # Check the dimension of the variables
print(X_train_padded.shape)
print(X_val_padded.shape)
print(X_test_padded.shape)
```

```
(7346, 29)
(1837, 29)
(3080, 29)
```

```
[ ]: # Check the first 3 elements of all X train variables
print(X_train_array[3])
print(X_train_sequences[3])
print(X_train_padded[3])
```

```
unauthorized direct debit account
[493, 113, 85, 2]
[493 113 85 2 0 0 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 0 0 0]
```

```
[ ]: # Define the folder path to save the test files
folder_path = '/content/drive/MyDrive/1. NLP CW/LSTM/test set/'

# Define the file path
Xtest_path = folder_path + 'X_test_padded.npy'
ytest_path = folder_path + 'y_test_array.npy'

# Save the test set
np.save(Xtest_path, X_test_padded)
np.save(ytest_path, y_test_array)
```

1.4 LSTM (with a word embedding layer)

1.4.1 LSTM (baseline)

```
[ ]: # Import tensorflow and fix the random seed
import tensorflow as tf
tf.random.set_seed(42)
```

```
[ ]: # Import keras from tensorflow and layers to build LSTM models
from tensorflow import keras
from tensorflow.keras import layers
```

```
[ ]: # Define the output dimension for the embedding layer and hidden units
embedding_output_dim = 100 # Random number
```

```

hidden_unit = 30 # Random number
nlabel = 77 # number of classes

# Build the baseline model
model = keras.models.Sequential()
model.add(layers.Embedding(voca_size, embedding_output_dim))
model.add(layers.LSTM(hidden_unit))
model.add(layers.Dense(nlabel, activation='softmax'))

# Compile the model
model.compile(optimizer='adam', loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

# Summary the model
model.summary()

```

Model: "sequential"

Layer (type)	Output Shape	Param #
embedding (Embedding)	(None, None, 100)	208800
lstm (LSTM)	(None, 30)	15720
dense (Dense)	(None, 77)	2387

=====
 Total params: 226907 (886.36 KB)
 Trainable params: 226907 (886.36 KB)
 Non-trainable params: 0 (0.00 Byte)
 =====

```

[ ]: # Define the folder path to save the model
folder_path = '/content/drive/MyDrive/1. NLP CW/LSTM/'

# Define the file path for the model checkpoint
model_checkpoint_path = folder_path + 'LSTM_embedding_model.keras'

# Define the model checkpoint
mc = tf.keras.callbacks.ModelCheckpoint(
    filepath=model_checkpoint_path,
    monitor='val_accuracy',
    mode='max',
    save_best_only=True)

```

```

[ ]: # Define early stopping

```

```
es = tf.keras.callbacks.EarlyStopping(monitor='val_loss', patience=3) # Random
↳ number of patience
```

```
[ ]: # Import time to measure the elapsed time
import time

# Measure time before training
start_time = time.time()

# Fit the model
history = model.fit(
    X_train_padded, y_train,
    epochs = 100, # Random number
    validation_data = (X_val_padded, y_val),
    callbacks = [mc, es],
    batch_size = 32) # Random number

# End the training time
end_time = time.time()


# Measure the training time
training_time = end_time - start_time
print("Total training time:", training_time, "seconds")
```

```
Epoch 1/100
230/230 [=====] - 11s 30ms/step - loss: 4.2625 -
accuracy: 0.0199 - val_loss: 4.0341 - val_accuracy: 0.0338
Epoch 2/100
230/230 [=====] - 5s 24ms/step - loss: 3.9935 -
accuracy: 0.0317 - val_loss: 3.9296 - val_accuracy: 0.0343
Epoch 3/100
230/230 [=====] - 7s 31ms/step - loss: 3.9381 -
accuracy: 0.0381 - val_loss: 3.8548 - val_accuracy: 0.0495
Epoch 4/100
230/230 [=====] - 5s 23ms/step - loss: 3.7424 -
accuracy: 0.0555 - val_loss: 3.6219 - val_accuracy: 0.0719
Epoch 5/100
230/230 [=====] - 6s 26ms/step - loss: 3.4658 -
accuracy: 0.0826 - val_loss: 3.3864 - val_accuracy: 0.1018
Epoch 6/100
230/230 [=====] - 7s 29ms/step - loss: 3.1828 -
accuracy: 0.1248 - val_loss: 3.1310 - val_accuracy: 0.1671
Epoch 7/100
230/230 [=====] - 6s 24ms/step - loss: 2.9293 -
accuracy: 0.1561 - val_loss: 2.9261 - val_accuracy: 0.1622
Epoch 8/100
230/230 [=====] - 8s 36ms/step - loss: 2.6887 -
```



accuracy: 0.2181 - val_loss: 2.7648 - val_accuracy: 0.2254
 Epoch 9/100
 230/230 [=====] - 5s 24ms/step - loss: 2.4905 -
 accuracy: 0.2641 - val_loss: 2.6963 - val_accuracy: 0.2471
 Epoch 10/100
 230/230 [=====] - 8s 34ms/step - loss: 2.2890 -
 accuracy: 0.3365 - val_loss: 2.4127 - val_accuracy: 0.3239
 Epoch 11/100
 230/230 [=====] - 6s 24ms/step - loss: 2.0995 -
 accuracy: 0.3897 - val_loss: 2.3355 - val_accuracy: 0.3571
 Epoch 12/100
 230/230 [=====] - 5s 22ms/step - loss: 1.9657 -
 accuracy: 0.4332 - val_loss: 2.1716 - val_accuracy: 0.4175
 Epoch 13/100
 230/230 [=====] - 7s 32ms/step - loss: 1.7997 -
 accuracy: 0.4861 - val_loss: 2.1235 - val_accuracy: 0.4175
 Epoch 14/100
 230/230 [=====] - 5s 23ms/step - loss: 1.6926 -
 accuracy: 0.5123 - val_loss: 2.0899 - val_accuracy: 0.4322
 Epoch 15/100
 230/230 [=====] - 6s 27ms/step - loss: 1.5568 -
 accuracy: 0.5573 - val_loss: 1.9327 - val_accuracy: 0.4714
 Epoch 16/100
 230/230 [=====] - 6s 25ms/step - loss: 1.4417 -
 accuracy: 0.5949 - val_loss: 1.9198 - val_accuracy: 0.4769
 Epoch 17/100
 230/230 [=====] - 5s 24ms/step - loss: 1.3537 -
 accuracy: 0.6250 - val_loss: 1.8871 - val_accuracy: 0.4921
 Epoch 18/100
 230/230 [=====] - 10s 44ms/step - loss: 1.2985 -
 accuracy: 0.6397 - val_loss: 1.7624 - val_accuracy: 0.5259
 Epoch 19/100
 230/230 [=====] - 5s 23ms/step - loss: 1.1997 -
 accuracy: 0.6710 - val_loss: 1.7091 - val_accuracy: 0.5514
 Epoch 20/100
 230/230 [=====] - 6s 27ms/step - loss: 1.1498 -
 accuracy: 0.6870 - val_loss: 1.6951 - val_accuracy: 0.5324
 Epoch 21/100
 230/230 [=====] - 6s 24ms/step - loss: 1.0622 -
 accuracy: 0.7189 - val_loss: 1.7325 - val_accuracy: 0.5596
 Epoch 22/100
 230/230 [=====] - 5s 23ms/step - loss: 0.9759 -
 accuracy: 0.7529 - val_loss: 1.5918 - val_accuracy: 0.5983
 Epoch 23/100
 230/230 [=====] - 8s 34ms/step - loss: 0.9697 -
 accuracy: 0.7509 - val_loss: 1.6296 - val_accuracy: 0.5841
 Epoch 24/100
 230/230 [=====] - 5s 23ms/step - loss: 0.8849 -

```
accuracy: 0.7753 - val_loss: 1.5478 - val_accuracy: 0.6135
Epoch 25/100
230/230 [=====] - 6s 24ms/step - loss: 0.8118 -
accuracy: 0.7920 - val_loss: 1.5075 - val_accuracy: 0.6282
Epoch 26/100
230/230 [=====] - 13s 54ms/step - loss: 0.7563 -
accuracy: 0.8128 - val_loss: 1.5453 - val_accuracy: 0.6124
Epoch 27/100
230/230 [=====] - 12s 52ms/step - loss: 0.7383 -
accuracy: 0.8161 - val_loss: 1.5192 - val_accuracy: 0.6184
Epoch 28/100
230/230 [=====] - 9s 39ms/step - loss: 0.7161 -
accuracy: 0.8176 - val_loss: 1.5145 - val_accuracy: 0.6233
Total training time: 193.95937156677246 seconds
```

```
[ ]: # Plot the loss
plt.title('Loss Plot')
plt.plot(history.history['loss'], label='train')
plt.plot(history.history['val_loss'], label='validation')
plt.legend()
plt.show()
```



```
[ ]: # Plot the accuracy
plt.title('Accuracy Plot')
plt.plot(history.history['accuracy'], label='train')
plt.plot(history.history['val_accuracy'], label='validation')
plt.legend()
plt.show()
```



```
[ ]: # Load the saved model
saved_model = tf.keras.models.load_model(model_checkpoint_path)

# Evaluate the model with the test set
loss, accuracy = saved_model.evaluate(X_test_padded, y_test_array)

print("Test Loss:", loss)
print("Test Accuracy:", round((accuracy*100), 2))
```

```
97/97 [=====] - 2s 10ms/step - loss: 1.4729 - accuracy:
0.6429
Test Loss: 1.472944736480713
Test Accuracy: 64.29
```

```
[ ]: # Import the library to check precision, recall, and F1 score
from sklearn.metrics import precision_score, recall_score, f1_score

# Check predictions with the test set
y_test_prob = saved_model.predict(X_test_padded)

# Convert probabilities to class labels
y_test_pred = np.argmax(y_test_prob, axis=1)

# Calculate precision, recall, and f1 score
precision = precision_score(y_test_array, y_test_pred, average='weighted')
recall = recall_score(y_test_array, y_test_pred, average='weighted')
f1 = f1_score(y_test_array, y_test_pred, average='weighted')

print("Precision:", round((precision*100), 2))
print("Recall:", round((recall*100), 2))
print("F1 Score:", round((f1*100), 2))
```

```
97/97 [=====] - 2s 8ms/step
Precision: 66.22
Recall: 64.29
F1 Score: 62.87
```

```
[ ]: # Error analysis
# Import the library for classification report
from sklearn.metrics import classification_report

# Print classification report
print(classification_report(y_test_array, y_test_pred))

# Check misclassified data
misclassified_data = np.where(y_test_pred != y_test_array)[0]
print(f"The number of misclassifications: {len(misclassified_data)}")

# Check the ratio of misclassifications
misclassification_ratio = (len(misclassified_data) / len(y_test_array)) * 100
# Round the number
rounded_ratio = round(misclassification_ratio, 2)
print(f"Proportion of misclassifications: {rounded_ratio}%")

# Iterate over misclassified data for error analysis
for idx in misclassified_data[:30]:
    input_text = X_test[idx]
    true_label = y_test[idx]
    predicted_label = y_test_pred[idx]

    # Print information about the misclassified data
```

```

print("Input Text:", input_text)
print("Actual Label:", true_label)
print("Predicted Label:", predicted_label)
print()

```

	precision	recall	f1-score	support
0	0.79	0.85	0.82	40
1	0.91	0.78	0.84	40
2	0.93	0.97	0.95	40
3	0.68	0.68	0.68	40
4	0.97	0.72	0.83	40
5	0.37	0.50	0.43	40
6	0.57	0.60	0.59	40
7	0.80	0.70	0.75	40
8	0.79	0.68	0.73	40
9	0.93	0.95	0.94	40
10	0.77	0.25	0.38	40
11	0.71	0.55	0.62	40
12	0.65	0.65	0.65	40
13	0.62	0.85	0.72	40
14	0.59	0.60	0.59	40
15	0.54	0.70	0.61	40
16	0.35	0.28	0.31	40
17	0.87	0.85	0.86	40
18	0.77	0.50	0.61	40
19	0.81	0.75	0.78	40
20	0.49	0.45	0.47	40
21	0.78	0.80	0.79	40
22	0.50	0.40	0.44	40
23	1.00	0.12	0.22	40
24	0.70	0.78	0.74	40
25	0.47	0.70	0.56	40
26	0.60	0.60	0.60	40
27	0.65	0.70	0.67	40
28	0.50	0.68	0.57	40
29	0.40	0.85	0.54	40
30	0.97	0.90	0.94	40
31	0.80	0.80	0.80	40
32	0.86	0.90	0.88	40
33	0.61	0.78	0.68	40
34	0.68	0.62	0.65	40
35	0.65	0.50	0.56	40
36	0.76	0.65	0.70	40
37	0.00	0.00	0.00	40
38	0.66	0.88	0.75	40
39	0.83	0.25	0.38	40

	40	0.61	0.95	0.75	40
	41	0.49	0.47	0.48	40
	42	0.79	0.85	0.82	40
	43	0.42	0.55	0.47	40
	44	0.58	0.88	0.70	40
	45	0.72	0.70	0.71	40
	46	0.82	0.90	0.86	40
	47	0.67	0.55	0.60	40
	48	0.63	0.47	0.54	40
	49	0.85	0.72	0.78	40
	50	0.67	0.55	0.60	40
	51	0.85	0.85	0.85	40
	52	0.76	0.88	0.81	40
	53	0.66	0.57	0.61	40
	54	0.35	0.72	0.47	40
	55	0.97	0.72	0.83	40
	56	0.63	0.55	0.59	40
	57	0.73	0.60	0.66	40
	58	0.62	0.62	0.62	40
	59	0.53	0.65	0.58	40
	60	0.79	0.82	0.80	40
	61	0.57	0.65	0.60	40
	62	0.53	0.68	0.59	40
	63	0.82	0.78	0.79	40
	64	0.53	0.65	0.58	40
	65	0.34	0.45	0.39	40
	66	0.57	0.53	0.55	40
	67	0.58	0.38	0.45	40
	68	0.68	0.33	0.44	40
	69	0.00	0.00	0.00	40
	70	0.72	0.78	0.75	40
	71	0.81	0.85	0.83	40
	72	1.00	0.03	0.05	40
	73	0.77	0.82	0.80	40
	74	0.35	0.90	0.50	40
	75	0.60	0.72	0.66	40
	76	0.69	0.62	0.66	40
	accuracy			0.64	3080
	macro avg	0.66	0.64	0.63	3080
	weighted avg	0.66	0.64	0.63	3080

The number of misclassifications: 1100

Proportion of misclassifications: 35.71%

Input Text: way know card arrive

Actual Label: 11

Predicted Label: 12

Input Text: card arrived yet
Actual Label: 11
Predicted Label: 13

Input Text: get card
Actual Label: 11
Predicted Label: 54

Input Text: received card
Actual Label: 11
Predicted Label: 12

Input Text: normal wait week new card
Actual Label: 11
Predicted Label: 56

Input Text: long card delivery take
Actual Label: 11
Predicted Label: 12

Input Text: still dont card weeks
Actual Label: 11
Predicted Label: 54

Input Text: still waiting card week ok
Actual Label: 11
Predicted Label: 56

Input Text: waiting longer expected bank card could provide information arrive
Actual Label: 11
Predicted Label: 62

Input Text: ive waiting longer expected card
Actual Label: 11
Predicted Label: 26

Input Text: card still hasnt arrived weeks lost
Actual Label: 11
Predicted Label: 13

Input Text: get card yet lost
Actual Label: 11
Predicted Label: 41

Input Text: ordered card weeks ago still isnt
Actual Label: 11
Predicted Label: 12

Input Text: card arrived yet
Actual Label: 11
Predicted Label: 13

Input Text: think something went wrong card delivery havent received yet
Actual Label: 11
Predicted Label: 62

Input Text: expecting new card wondering havent received yet
Actual Label: 11
Predicted Label: 5

Input Text: know card arrive
Actual Label: 11
Predicted Label: 12

Input Text: ordered card still havent received two weeks
Actual Label: 11
Predicted Label: 12

Input Text: wont card show app
Actual Label: 13
Predicted Label: 56

Input Text: add card account
Actual Label: 13
Predicted Label: 54

Input Text: put old card back system found
Actual Label: 13
Predicted Label: 31

Input Text: view card received app
Actual Label: 13
Predicted Label: 11

Input Text: website go link card
Actual Label: 13
Predicted Label: 62

Input Text: app doesnt show card received
Actual Label: 13
Predicted Label: 56

Input Text: international exchange rates
Actual Label: 32
Predicted Label: 75

Input Text: please advise exchange rate
Actual Label: 32
Predicted Label: 17

Input Text: good time exchange
Actual Label: 32
Predicted Label: 61

Input Text: much get exchange rate
Actual Label: 32
Predicted Label: 76

Input Text: made currency exchange think charged
Actual Label: 17
Predicted Label: 31

Input Text: rate low sure using right exchange rate
Actual Label: 17
Predicted Label: 76

1.4.2 LSTM (with dropout)

```
[ ]: # Define the output dimension for the embedding layer and hidden units
embedding_output_dim = 100 # Random number
hidden_unit = 30 # Random number
nlabel = 77 # number of classes

# Build the baseline model
dropout_model = keras.models.Sequential()
dropout_model.add(layers.Embedding(voca_size, embedding_output_dim))
dropout_model.add(layers.LSTM(hidden_unit, dropout=0.2)) # Random number
dropout_model.add(layers.Dense(nlabel, activation='softmax'))

# Compile the model
dropout_model.compile(optimizer='adam', loss='sparse_categorical_crossentropy',
    metrics=['accuracy'])

# Summary the model
dropout_model.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
embedding_1 (Embedding)	(None, None, 100)	208800
lstm_1 (LSTM)	(None, 30)	15720

dense_1 (Dense) (None, 77) 2387

```
=====
Total params: 226907 (886.36 KB)
Trainable params: 226907 (886.36 KB)
Non-trainable params: 0 (0.00 Byte)
-----
```

```
[ ]: # Define the folder path to save the model
folder_path = '/content/drive/MyDrive/1. NLP CW/LSTM/'

# Define the file path for the model checkpoint
model_checkpoint_path = folder_path + 'dropout_LSTM_embedding_model.keras'

# Define the model checkpoint
mc = tf.keras.callbacks.ModelCheckpoint(
    filepath=model_checkpoint_path,
    monitor='val_accuracy',
    mode='max',
    save_best_only=True)
```

```
[ ]: # Import time to measure the elapsed time
import time

# Measure time before training
start_time = time.time()

# Fit the model
history = dropout_model.fit(
    X_train_padded, y_train,
    epochs = 100, # Random number
    validation_data = (X_val_padded, y_val),
    callbacks = [mc, es],
    batch_size = 32)

# End the training time
end_time = time.time()

# Measure the training time
training_time = end_time - start_time
print("Total training time:", training_time, "seconds")
```

```
Epoch 1/100
230/230 [=====] - 10s 29ms/step - loss: 4.2711 -
accuracy: 0.0182 - val_loss: 4.0823 - val_accuracy: 0.0338
Epoch 2/100
230/230 [=====] - 7s 30ms/step - loss: 4.0866 -
```

accuracy: 0.0283 - val_loss: 4.1577 - val_accuracy: 0.0278
Epoch 3/100
230/230 [=====] - 8s 34ms/step - loss: 4.0583 -
accuracy: 0.0252 - val_loss: 4.0054 - val_accuracy: 0.0327
Epoch 4/100
230/230 [=====] - 11s 47ms/step - loss: 3.9583 -
accuracy: 0.0294 - val_loss: 3.8965 - val_accuracy: 0.0463
Epoch 5/100
230/230 [=====] - 9s 38ms/step - loss: 3.8102 -
accuracy: 0.0438 - val_loss: 3.7361 - val_accuracy: 0.0550
Epoch 6/100
230/230 [=====] - 7s 31ms/step - loss: 3.6619 -
accuracy: 0.0636 - val_loss: 3.6422 - val_accuracy: 0.0659
Epoch 7/100
230/230 [=====] - 8s 33ms/step - loss: 3.5815 -
accuracy: 0.0679 - val_loss: 3.6077 - val_accuracy: 0.0680
Epoch 8/100
230/230 [=====] - 11s 50ms/step - loss: 3.4941 -
accuracy: 0.0760 - val_loss: 3.5286 - val_accuracy: 0.0724
Epoch 9/100
230/230 [=====] - 11s 46ms/step - loss: 3.4428 -
accuracy: 0.0825 - val_loss: 3.4720 - val_accuracy: 0.0838
Epoch 10/100
230/230 [=====] - 11s 50ms/step - loss: 3.3883 -
accuracy: 0.0962 - val_loss: 3.3982 - val_accuracy: 0.0920
Epoch 11/100
230/230 [=====] - 10s 44ms/step - loss: 3.3377 -
accuracy: 0.1033 - val_loss: 3.4235 - val_accuracy: 0.0942
Epoch 12/100
230/230 [=====] - 9s 41ms/step - loss: 3.3131 -
accuracy: 0.1075 - val_loss: 3.3779 - val_accuracy: 0.1062
Epoch 13/100
230/230 [=====] - 10s 45ms/step - loss: 3.2711 -
accuracy: 0.1108 - val_loss: 3.3510 - val_accuracy: 0.1127
Epoch 14/100
230/230 [=====] - 6s 25ms/step - loss: 3.2236 -
accuracy: 0.1207 - val_loss: 3.3295 - val_accuracy: 0.1165
Epoch 15/100
230/230 [=====] - 8s 33ms/step - loss: 3.2137 -
accuracy: 0.1284 - val_loss: 3.3404 - val_accuracy: 0.1230
Epoch 16/100
230/230 [=====] - 6s 25ms/step - loss: 3.1697 -
accuracy: 0.1374 - val_loss: 3.3129 - val_accuracy: 0.1263
Epoch 17/100
230/230 [=====] - 7s 31ms/step - loss: 3.1523 -
accuracy: 0.1410 - val_loss: 3.2974 - val_accuracy: 0.1345
Epoch 18/100
230/230 [=====] - 11s 47ms/step - loss: 3.1204 -

accuracy: 0.1504 - val_loss: 3.1890 - val_accuracy: 0.1475
 Epoch 19/100
 230/230 [=====] - 11s 47ms/step - loss: 2.9976 -
 accuracy: 0.1678 - val_loss: 3.0844 - val_accuracy: 0.1579
 Epoch 20/100
 230/230 [=====] - 11s 46ms/step - loss: 2.8706 -
 accuracy: 0.1808 - val_loss: 2.9859 - val_accuracy: 0.1688
 Epoch 21/100
 230/230 [=====] - 11s 47ms/step - loss: 2.8179 -
 accuracy: 0.1899 - val_loss: 3.0967 - val_accuracy: 0.1590
 Epoch 22/100
 230/230 [=====] - 7s 30ms/step - loss: 2.7734 -
 accuracy: 0.1978 - val_loss: 2.9003 - val_accuracy: 0.1873
 Epoch 23/100
 230/230 [=====] - 8s 33ms/step - loss: 2.6465 -
 accuracy: 0.2169 - val_loss: 2.8283 - val_accuracy: 0.2150
 Epoch 24/100
 230/230 [=====] - 7s 31ms/step - loss: 2.5462 -
 accuracy: 0.2343 - val_loss: 2.7194 - val_accuracy: 0.2226
 Epoch 25/100
 230/230 [=====] - 8s 35ms/step - loss: 2.4750 -
 accuracy: 0.2473 - val_loss: 2.6554 - val_accuracy: 0.2395
 Epoch 26/100
 230/230 [=====] - 5s 23ms/step - loss: 2.3949 -
 accuracy: 0.2691 - val_loss: 2.6915 - val_accuracy: 0.2368
 Epoch 27/100
 230/230 [=====] - 7s 29ms/step - loss: 2.3553 -
 accuracy: 0.2731 - val_loss: 2.5588 - val_accuracy: 0.2695
 Epoch 28/100
 230/230 [=====] - 6s 26ms/step - loss: 2.2898 -
 accuracy: 0.2919 - val_loss: 2.5037 - val_accuracy: 0.2793
 Epoch 29/100
 230/230 [=====] - 5s 24ms/step - loss: 2.1817 -
 accuracy: 0.3185 - val_loss: 2.4204 - val_accuracy: 0.2885
 Epoch 30/100
 230/230 [=====] - 8s 35ms/step - loss: 2.0855 -
 accuracy: 0.3339 - val_loss: 2.4059 - val_accuracy: 0.2907
 Epoch 31/100
 230/230 [=====] - 6s 26ms/step - loss: 2.0326 -
 accuracy: 0.3464 - val_loss: 2.3102 - val_accuracy: 0.3103
 Epoch 32/100
 230/230 [=====] - 7s 29ms/step - loss: 1.9632 -
 accuracy: 0.3765 - val_loss: 2.2238 - val_accuracy: 0.3538
 Epoch 33/100
 230/230 [=====] - 7s 30ms/step - loss: 1.8574 -
 accuracy: 0.4063 - val_loss: 2.2024 - val_accuracy: 0.3664
 Epoch 34/100
 230/230 [=====] - 5s 22ms/step - loss: 1.7783 -

accuracy: 0.4283 - val_loss: 2.1439 - val_accuracy: 0.3685
 Epoch 35/100
 230/230 [=====] - 8s 35ms/step - loss: 1.7123 -
 accuracy: 0.4396 - val_loss: 2.0713 - val_accuracy: 0.3996
 Epoch 36/100
 230/230 [=====] - 5s 23ms/step - loss: 1.6415 -
 accuracy: 0.4639 - val_loss: 2.0214 - val_accuracy: 0.4148
 Epoch 37/100
 230/230 [=====] - 7s 29ms/step - loss: 1.5779 -
 accuracy: 0.4936 - val_loss: 1.9272 - val_accuracy: 0.4339
 Epoch 38/100
 230/230 [=====] - 6s 28ms/step - loss: 1.5008 -
 accuracy: 0.5144 - val_loss: 1.8958 - val_accuracy: 0.4502
 Epoch 39/100
 230/230 [=====] - 6s 24ms/step - loss: 1.4564 -
 accuracy: 0.5238 - val_loss: 1.8957 - val_accuracy: 0.4638
 Epoch 40/100
 230/230 [=====] - 8s 36ms/step - loss: 1.4113 -
 accuracy: 0.5411 - val_loss: 1.8400 - val_accuracy: 0.4703
 Epoch 41/100
 230/230 [=====] - 8s 34ms/step - loss: 1.3995 -
 accuracy: 0.5502 - val_loss: 1.8106 - val_accuracy: 0.4850
 Epoch 42/100
 230/230 [=====] - 11s 48ms/step - loss: 1.3116 -
 accuracy: 0.5784 - val_loss: 1.7411 - val_accuracy: 0.5019
 Epoch 43/100
 230/230 [=====] - 6s 26ms/step - loss: 1.2781 -
 accuracy: 0.5828 - val_loss: 1.7210 - val_accuracy: 0.5166
 Epoch 44/100
 230/230 [=====] - 9s 40ms/step - loss: 1.2346 -
 accuracy: 0.5992 - val_loss: 1.7140 - val_accuracy: 0.5182
 Epoch 45/100
 230/230 [=====] - 6s 28ms/step - loss: 1.1909 -
 accuracy: 0.6100 - val_loss: 1.6889 - val_accuracy: 0.5433
 Epoch 46/100
 230/230 [=====] - 6s 26ms/step - loss: 1.1664 -
 accuracy: 0.6297 - val_loss: 1.6523 - val_accuracy: 0.5422
 Epoch 47/100
 230/230 [=====] - 8s 35ms/step - loss: 1.3018 -
 accuracy: 0.5883 - val_loss: 1.7145 - val_accuracy: 0.5514
 Epoch 48/100
 230/230 [=====] - 6s 25ms/step - loss: 1.1404 -
 accuracy: 0.6510 - val_loss: 1.6129 - val_accuracy: 0.5504
 Epoch 49/100
 230/230 [=====] - 7s 32ms/step - loss: 1.1076 -
 accuracy: 0.6570 - val_loss: 1.6105 - val_accuracy: 0.5607
 Epoch 50/100
 230/230 [=====] - 6s 25ms/step - loss: 1.0322 -

```

accuracy: 0.6862 - val_loss: 1.5449 - val_accuracy: 0.5955
Epoch 51/100
230/230 [=====] - 6s 28ms/step - loss: 0.9976 -
accuracy: 0.7028 - val_loss: 1.5503 - val_accuracy: 0.6026
Epoch 52/100
230/230 [=====] - 7s 30ms/step - loss: 0.9781 -
accuracy: 0.7013 - val_loss: 1.5233 - val_accuracy: 0.6102
Epoch 53/100
230/230 [=====] - 5s 23ms/step - loss: 0.9301 -
accuracy: 0.7280 - val_loss: 1.5173 - val_accuracy: 0.6102
Epoch 54/100
230/230 [=====] - 8s 33ms/step - loss: 0.9084 -
accuracy: 0.7360 - val_loss: 1.4822 - val_accuracy: 0.6238
Epoch 55/100
230/230 [=====] - 6s 24ms/step - loss: 0.8920 -
accuracy: 0.7457 - val_loss: 1.5222 - val_accuracy: 0.6113
Epoch 56/100
230/230 [=====] - 8s 35ms/step - loss: 0.8535 -
accuracy: 0.7542 - val_loss: 1.4702 - val_accuracy: 0.6287
Epoch 57/100
230/230 [=====] - 7s 29ms/step - loss: 0.8575 -
accuracy: 0.7518 - val_loss: 1.4424 - val_accuracy: 0.6353
Epoch 58/100
230/230 [=====] - 6s 24ms/step - loss: 0.8061 -
accuracy: 0.7735 - val_loss: 1.4190 - val_accuracy: 0.6413
Epoch 59/100
230/230 [=====] - 8s 34ms/step - loss: 0.7699 -
accuracy: 0.7806 - val_loss: 1.4324 - val_accuracy: 0.6478
Epoch 60/100
230/230 [=====] - 5s 24ms/step - loss: 0.7414 -
accuracy: 0.7927 - val_loss: 1.3918 - val_accuracy: 0.6511
Epoch 61/100
230/230 [=====] - 6s 28ms/step - loss: 0.7441 -
accuracy: 0.7916 - val_loss: 1.3689 - val_accuracy: 0.6576
Epoch 62/100
230/230 [=====] - 6s 28ms/step - loss: 0.7009 -
accuracy: 0.8045 - val_loss: 1.3883 - val_accuracy: 0.6543
Epoch 63/100
230/230 [=====] - 6s 25ms/step - loss: 0.6720 -
accuracy: 0.8128 - val_loss: 1.3779 - val_accuracy: 0.6565
Epoch 64/100
230/230 [=====] - 8s 37ms/step - loss: 0.6601 -
accuracy: 0.8200 - val_loss: 1.3864 - val_accuracy: 0.6614
Total training time: 484.262811422348 seconds


```

```


[ ]: # Plot the loss
plt.title('Loss Plot')

```

```
plt.plot(history.history['loss'], label='train')
plt.plot(history.history['val_loss'], label='validation')
plt.legend()
plt.show()
```



```
[ ]: # Plot the accuracy
plt.title('Accuracy Plot')
plt.plot(history.history['accuracy'], label='train')
plt.plot(history.history['val_accuracy'], label='validation')
plt.legend()
plt.show()
```



```
[ ]: # Load the saved model
saved_model = tf.keras.models.load_model(model_checkpoint_path)

# Evaluate the model with the test set
loss, accuracy = saved_model.evaluate(X_test_padded, y_test_array)

print("Test Loss:", loss)
print("Test Accuracy:", round((accuracy*100), 2))
```

```
97/97 [=====] - 2s 8ms/step - loss: 1.3496 - accuracy:
0.6834
Test Loss: 1.349645972251892
Test Accuracy: 68.34
```

```
[ ]: # Check predictions with the test set
y_test_prob = saved_model.predict(X_test_padded)

# Convert probabilities to class labels
y_test_pred = np.argmax(y_test_prob, axis=1)

# Calculate precision, recall, and f1 score
```



```

precision = precision_score(y_test_array, y_test_pred, average='weighted')
recall = recall_score(y_test_array, y_test_pred, average='weighted')
f1 = f1_score(y_test_array, y_test_pred, average='weighted')

print("Precision:", round((precision*100), 2))
print("Recall:", round((recall*100), 2))
print("F1 Score:", round((f1*100), 2))

```

```

97/97 [=====] - 2s 9ms/step
Precision: 68.85
Recall: 68.34
F1 Score: 67.45

```

```

[ ]: # Error analysis

# Print classification report
print(classification_report(y_test_array, y_test_pred))

# Check misclassified data
misclassified_data = np.where(y_test_pred != y_test_array)[0]
print(f"The number of misclassifications: {len(misclassified_data)}")

# Check the ratio of misclassifications
misclassification_ratio = (len(misclassified_data) / len(y_test_array)) * 100
# Round the number
rounded_ratio = round(misclassification_ratio, 2)
print(f"Proportion of misclassifications: {rounded_ratio}%")

# Iterate over misclassified data for error analysis
for idx in misclassified_data[:30]:
    input_text = X_test[idx]
    true_label = y_test[idx]
    predicted_label = y_test_pred[idx]

    # Print information about the misclassified data
    print("Input Text:", input_text)
    print("Actual Label:", true_label)
    print("Predicted Label:", predicted_label)
    print()

```

	precision	recall	f1-score	support
0	0.77	0.90	0.83	40
1	0.89	0.85	0.87	40
2	0.83	0.97	0.90	40
3	0.80	0.60	0.69	40
4	0.79	0.82	0.80	40
5	0.54	0.78	0.64	40

6	0.66	0.72	0.69	40
7	0.66	0.72	0.69	40
8	0.75	0.82	0.79	40
9	0.92	0.85	0.88	40
10	0.55	0.15	0.24	40
11	0.70	0.78	0.74	40
12	0.71	0.60	0.65	40
13	0.94	0.80	0.86	40
14	0.38	0.60	0.46	40
15	0.69	0.88	0.77	40
16	0.58	0.70	0.64	40
17	0.82	0.82	0.82	40
18	0.71	0.68	0.69	40
19	0.87	0.82	0.85	40
20	0.59	0.47	0.53	40
21	0.49	0.85	0.62	40
22	0.56	0.57	0.57	40
23	1.00	0.70	0.82	40
24	0.92	0.85	0.88	40
25	0.62	0.75	0.68	40
26	0.43	0.65	0.51	40
27	0.66	0.62	0.64	40
28	0.77	0.68	0.72	40
29	0.79	0.75	0.77	40
30	0.66	0.95	0.78	40
31	0.67	0.72	0.70	40
32	0.82	0.93	0.87	40
33	0.70	0.82	0.76	40
34	0.77	0.85	0.81	40
35	0.71	0.72	0.72	40
36	0.76	0.65	0.70	40
37	0.64	0.70	0.67	40
38	0.60	0.30	0.40	40
39	0.68	0.70	0.69	40
40	0.69	0.85	0.76	40
41	0.50	0.15	0.23	40
42	0.92	0.88	0.90	40
43	0.65	0.60	0.62	40
44	0.93	0.97	0.95	40
45	0.79	0.78	0.78	40
46	0.81	0.88	0.84	40
47	0.42	0.55	0.47	40
48	0.40	0.45	0.42	40
49	0.88	0.70	0.78	40
50	0.62	0.65	0.63	40
51	0.67	0.78	0.72	40
52	0.62	0.65	0.63	40
53	0.82	0.78	0.79	40

	54	0.55	0.65	0.60	40
	55	0.94	0.78	0.85	40
	56	0.66	0.47	0.55	40
	57	0.00	0.00	0.00	40
	58	0.80	0.70	0.75	40
	59	0.32	0.45	0.37	40
	60	0.94	0.75	0.83	40
	61	0.52	0.62	0.57	40
	62	0.53	0.57	0.55	40
	63	0.79	0.85	0.82	40
	64	0.67	0.80	0.73	40
	65	0.69	0.60	0.64	40
	66	0.48	0.25	0.33	40
	67	0.73	0.68	0.70	40
	68	0.35	0.55	0.43	40
	69	0.62	0.25	0.36	40
	70	0.84	0.95	0.89	40
	71	0.84	0.95	0.89	40
	72	0.83	0.38	0.52	40
	73	0.76	0.88	0.81	40
	74	0.57	0.50	0.53	40
	75	0.62	0.60	0.61	40
	76	0.83	0.62	0.71	40
accuracy				0.68	3080
macro avg		0.69	0.68	0.67	3080
weighted avg		0.69	0.68	0.67	3080

The number of misclassifications: 975

Proportion of misclassifications: 31.66%

Input Text: locate card

Actual Label: 11

Predicted Label: 39

Input Text: way know card arrive

Actual Label: 11

Predicted Label: 12

Input Text: get card

Actual Label: 11

Predicted Label: 43

Input Text: card still hasnt arrived weeks lost

Actual Label: 11

Predicted Label: 0

Input Text: get card yet lost

Actual Label: 11

Predicted Label: 0

Input Text: think something went wrong card delivery havent received yet

Actual Label: 11

Predicted Label: 43

Input Text: expecting new card wondering havent received yet

Actual Label: 11

Predicted Label: 66

Input Text: know card arrive

Actual Label: 11

Predicted Label: 12

Input Text: ordered card still havent received two weeks

Actual Label: 11

Predicted Label: 47

Input Text: readd card app

Actual Label: 13

Predicted Label: 0

Input Text: add card account

Actual Label: 13

Predicted Label: 24

Input Text: view card received app

Actual Label: 13

Predicted Label: 11

Input Text: ive received card need know sync app

Actual Label: 13

Predicted Label: 11

Input Text: app doesnt show card received

Actual Label: 13

Predicted Label: 11

Input Text: way make old card usable app

Actual Label: 13

Predicted Label: 54

Input Text: need go app enter card info

Actual Label: 13

Predicted Label: 11

Input Text: link another card account

Actual Label: 13

Predicted Label: 0

Input Text: often exchange rates change

Actual Label: 32

Predicted Label: 33

Input Text: good time exchange

Actual Label: 32

Predicted Label: 31

Input Text: currencies exchange rate calculated

Actual Label: 32

Predicted Label: 33

Input Text: made currency exchange think charged

Actual Label: 17

Predicted Label: 31

Input Text: rate low sure using right exchange rate

Actual Label: 17

Predicted Label: 76

Input Text: rate applied foreign purchase incorrect

Actual Label: 17

Predicted Label: 2

Input Text: charged

Actual Label: 17

Predicted Label: 15

Input Text: charged wrong currency exchange purchase

Actual Label: 17

Predicted Label: 31

Input Text: conversion value card payments incorrect

Actual Label: 17

Predicted Label: 2

Input Text: paid something foreign currency noticed exchange rate incorrect

Actual Label: 17

Predicted Label: 76

Input Text: fee dont recognize statement

Actual Label: 34

Predicted Label: 15

Input Text: explain random charge

Actual Label: 34

Predicted Label: 16

Input Text: transaction credited

Actual Label: 34

Predicted Label: 64

1.4.3 Hyperparameter tuning

```
[ ]: # The code for hyperparameter tuning is derived from the Tensorflow website.  
# (https://www.tensorflow.org/tutorials/keras/keras_tuner)
```

```
# Import and install libraries for hyperparameter tuning  
import IPython  
!pip install -q -U keras-tuner  
import kerastuner as kt
```

129.1/129.1

kB 1.4 MB/s eta 0:00:00

<ipython-input-81-8b29936803b5>:7: DeprecationWarning: `import kerastuner` is deprecated, please use `import keras_tuner`.

```
import kerastuner as kt
```

```
[ ]: # The code for hyperparameter tuning is derived from the Tensorflow website.  
# (https://www.tensorflow.org/tutorials/keras/keras_tuner)
```

```
# Define the model for hyperparameter tuning  
def model_builder(hp):  
    model = keras.models.Sequential()  
    model.add(layers.Embedding(voca_size, embedding_output_dim)) # Use the same  
    ↪ dimension from the baseline model  
    hp_units = hp.Int('units', min_value = 20, max_value = 50, step = 10) # Set  
    ↪ up the hyperparameters  
    model.add(layers.LSTM(units = hp_units)) # We will check the optimal hidden  
    ↪ unit for the LSTM layer  
    model.add(layers.Dense(nlabel, activation='softmax'))  
  
    hp_learning_rate = hp.Choice('learning_rate', values = [0.01, 0.001, 0.0001])  
    ↪ # Set up the hyperparameters  
    model.compile(optimizer = keras.optimizers.Adam(learning_rate =  
    ↪ hp_learning_rate), # We will check the optimal learning rate  
                  loss = 'sparse_categorical_crossentropy',  
                  metrics = ['accuracy'])  
    return model
```

```
[ ]: # The code for hyperparameter tuning is derived from the Tensorflow website.
# (https://www.tensorflow.org/tutorials/keras/keras_tuner)

# Specify the tuner
tuner = kt.Hyperband(model_builder,
                     objective = 'val_accuracy',
                     max_epochs = 100)
```

```
[ ]: # The code for hyperparameter tuning is derived from the Tensorflow website.
# (https://www.tensorflow.org/tutorials/keras/keras_tuner)

# Set up a callback for early stopping
stop_early = tf.keras.callbacks.EarlyStopping(monitor='val_loss', patience=5)
```

```
[ ]: # The code for hyperparameter tuning is derived from the Tensorflow website.
# (https://www.tensorflow.org/tutorials/keras/keras_tuner)

# Run the tuner
tuner.search(X_train_padded, y_train, epochs = 100, validation_data = (
    X_val_padded, y_val), callbacks = [stop_early])

# Get the optimal hyperparameters
best_hps = tuner.get_best_hyperparameters(num_trials = 1)[0]

print(f"The optimal number of units: {best_hps.get('units')}. The optimal
    learning rate: {best_hps.get('learning_rate')}")
```

The optimal number of units: 40. The optimal learning rate: 0.001.

1.4.4 Tuned LSTM

```
[ ]: # Define the output dimension for the embedding layer and hidden units
embedding_output_dim = 100 # Random number
hidden_unit = 40 # Number from hyperparameter tuning
nlabel = 77 # number of classes

# Build the baseline model
tuned_model = keras.models.Sequential()
tuned_model.add(layers.Embedding(voca_size, embedding_output_dim))
tuned_model.add(layers.LSTM(hidden_unit))
tuned_model.add(layers.Dense(nlabel, activation='softmax'))

# Compile the model
tuned_model.compile(optimizer='adam', loss='sparse_categorical_crossentropy',
    metrics=['accuracy']) # 0.001 is the default of Adam

# Summary the model
```

```
tuned_model.summary()
```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
embedding_2 (Embedding)	(None, None, 100)	208800
lstm_2 (LSTM)	(None, 40)	22560
dense_2 (Dense)	(None, 77)	3157

=====
Total params: 234517 (916.08 KB)
Trainable params: 234517 (916.08 KB)
Non-trainable params: 0 (0.00 Byte)
=====

```
[ ]: # Define the folder path to save the model
folder_path = '/content/drive/MyDrive/1. NLP CW/LSTM/'

# Define the file path for the model checkpoint
model_checkpoint_path = folder_path + 'tuned_LSTM_embedding_model.keras'

# Define the model checkpoint
mc = tf.keras.callbacks.ModelCheckpoint(
    filepath=model_checkpoint_path,
    monitor='val_accuracy',
    mode='max',
    save_best_only=True)
```

```
[ ]: # Import time to measure the elapsed time
import time

# Measure time before training
start_time = time.time()

# Fit the model
history = tuned_model.fit(
    X_train_padded, y_train,
    epochs = 100, # Random number
    validation_data = (X_val_padded, y_val),
    callbacks = [mc, es],
    batch_size = 32)

# End the training time
end_time = time.time()
```



```
# Measure the training time
training_time = end_time - start_time
print("Total training time:", training_time, "seconds")
```

```
Epoch 1/100
230/230 [=====] - 9s 28ms/step - loss: 4.2946 -
accuracy: 0.0165 - val_loss: 4.0789 - val_accuracy: 0.0272
Epoch 2/100
230/230 [=====] - 8s 34ms/step - loss: 4.0052 -
accuracy: 0.0297 - val_loss: 3.9721 - val_accuracy: 0.0321
Epoch 3/100
230/230 [=====] - 5s 23ms/step - loss: 3.9831 -
accuracy: 0.0282 - val_loss: 3.9574 - val_accuracy: 0.0299
Epoch 4/100
230/230 [=====] - 6s 27ms/step - loss: 3.9313 -
accuracy: 0.0282 - val_loss: 3.9680 - val_accuracy: 0.0338
Epoch 5/100
230/230 [=====] - 8s 36ms/step - loss: 3.9267 -
accuracy: 0.0274 - val_loss: 3.9583 - val_accuracy: 0.0310
Epoch 6/100
230/230 [=====] - 6s 25ms/step - loss: 3.9227 -
accuracy: 0.0320 - val_loss: 3.9566 - val_accuracy: 0.0359
Epoch 7/100
230/230 [=====] - 7s 33ms/step - loss: 3.8577 -
accuracy: 0.0359 - val_loss: 3.8474 - val_accuracy: 0.0463
Epoch 8/100
230/230 [=====] - 5s 23ms/step - loss: 3.8067 -
accuracy: 0.0388 - val_loss: 3.8269 - val_accuracy: 0.0397
Epoch 9/100
230/230 [=====] - 7s 31ms/step - loss: 3.7778 -
accuracy: 0.0406 - val_loss: 3.7698 - val_accuracy: 0.0474
Epoch 10/100
230/230 [=====] - 6s 24ms/step - loss: 3.6627 -
accuracy: 0.0501 - val_loss: 3.9440 - val_accuracy: 0.0419
Epoch 11/100
230/230 [=====] - 5s 24ms/step - loss: 3.5898 -
accuracy: 0.0542 - val_loss: 3.5175 - val_accuracy: 0.0648
Epoch 12/100
230/230 [=====] - 8s 36ms/step - loss: 3.4428 -
accuracy: 0.0672 - val_loss: 3.4485 - val_accuracy: 0.0784
Epoch 13/100
230/230 [=====] - 13s 59ms/step - loss: 3.3126 -
accuracy: 0.0894 - val_loss: 3.3826 - val_accuracy: 0.1056
Epoch 14/100
230/230 [=====] - 14s 61ms/step - loss: 3.1820 -
accuracy: 0.1115 - val_loss: 3.2661 - val_accuracy: 0.1225
```

Epoch 15/100
230/230 [=====] - 14s 62ms/step - loss: 3.0260 - accuracy: 0.1417 - val_loss: 3.0991 - val_accuracy: 0.1410

Epoch 16/100
230/230 [=====] - 14s 62ms/step - loss: 2.8354 - accuracy: 0.1756 - val_loss: 2.9109 - val_accuracy: 0.1824

Epoch 17/100
230/230 [=====] - 9s 40ms/step - loss: 2.6705 - accuracy: 0.2081 - val_loss: 2.8185 - val_accuracy: 0.2009

Epoch 18/100
230/230 [=====] - 6s 25ms/step - loss: 2.4823 - accuracy: 0.2554 - val_loss: 2.6742 - val_accuracy: 0.2640

Epoch 19/100
230/230 [=====] - 7s 28ms/step - loss: 2.3056 - accuracy: 0.3056 - val_loss: 2.5004 - val_accuracy: 0.2847

Epoch 20/100
230/230 [=====] - 7s 29ms/step - loss: 2.1421 - accuracy: 0.3471 - val_loss: 2.3917 - val_accuracy: 0.3130

Epoch 21/100
230/230 [=====] - 6s 25ms/step - loss: 2.0168 - accuracy: 0.3855 - val_loss: 2.3748 - val_accuracy: 0.3435

Epoch 22/100
230/230 [=====] - 8s 35ms/step - loss: 1.8786 - accuracy: 0.4239 - val_loss: 2.1859 - val_accuracy: 0.4028

Epoch 23/100
230/230 [=====] - 6s 25ms/step - loss: 1.7576 - accuracy: 0.4608 - val_loss: 2.1545 - val_accuracy: 0.3990

Epoch 24/100
230/230 [=====] - 7s 29ms/step - loss: 1.7189 - accuracy: 0.4763 - val_loss: 2.0777 - val_accuracy: 0.4366

Epoch 25/100
230/230 [=====] - 6s 27ms/step - loss: 1.5775 - accuracy: 0.5347 - val_loss: 1.9711 - val_accuracy: 0.4736

Epoch 26/100
230/230 [=====] - 5s 23ms/step - loss: 1.4813 - accuracy: 0.5585 - val_loss: 1.9505 - val_accuracy: 0.4769

Epoch 27/100
230/230 [=====] - 9s 38ms/step - loss: 1.3838 - accuracy: 0.5996 - val_loss: 1.8017 - val_accuracy: 0.5253

Epoch 28/100
230/230 [=====] - 6s 25ms/step - loss: 1.2723 - accuracy: 0.6327 - val_loss: 1.7867 - val_accuracy: 0.5329

Epoch 29/100
230/230 [=====] - 7s 31ms/step - loss: 1.2030 - accuracy: 0.6576 - val_loss: 1.7268 - val_accuracy: 0.5297

Epoch 30/100
230/230 [=====] - 6s 25ms/step - loss: 1.1218 - accuracy: 0.6779 - val_loss: 1.6611 - val_accuracy: 0.5525

Epoch 31/100
230/230 [=====] - 7s 29ms/step - loss: 1.0285 - accuracy: 0.7069 - val_loss: 1.5686 - val_accuracy: 0.5819

Epoch 32/100
230/230 [=====] - 9s 38ms/step - loss: 1.0353 - accuracy: 0.7022 - val_loss: 1.6081 - val_accuracy: 0.5634

Epoch 33/100
230/230 [=====] - 5s 24ms/step - loss: 0.9411 - accuracy: 0.7291 - val_loss: 1.4938 - val_accuracy: 0.6064

Epoch 34/100
230/230 [=====] - 7s 30ms/step - loss: 0.8535 - accuracy: 0.7539 - val_loss: 1.5312 - val_accuracy: 0.5765

Epoch 35/100
230/230 [=====] - 6s 26ms/step - loss: 0.8156 - accuracy: 0.7723 - val_loss: 1.4590 - val_accuracy: 0.6184

Epoch 36/100
230/230 [=====] - 6s 24ms/step - loss: 0.7677 - accuracy: 0.7902 - val_loss: 1.3816 - val_accuracy: 0.6554

Epoch 37/100
230/230 [=====] - 8s 33ms/step - loss: 0.7014 - accuracy: 0.8121 - val_loss: 1.3796 - val_accuracy: 0.6456

Epoch 38/100
230/230 [=====] - 5s 24ms/step - loss: 0.6795 - accuracy: 0.8124 - val_loss: 1.3965 - val_accuracy: 0.6483

Epoch 39/100
230/230 [=====] - 7s 31ms/step - loss: 0.6476 - accuracy: 0.8288 - val_loss: 1.3664 - val_accuracy: 0.6522

Epoch 40/100
230/230 [=====] - 6s 25ms/step - loss: 0.5844 - accuracy: 0.8460 - val_loss: 1.2906 - val_accuracy: 0.6728

Epoch 41/100
230/230 [=====] - 6s 25ms/step - loss: 0.5609 - accuracy: 0.8526 - val_loss: 1.2927 - val_accuracy: 0.6832

Epoch 42/100
230/230 [=====] - 8s 33ms/step - loss: 0.5291 - accuracy: 0.8633 - val_loss: 1.3004 - val_accuracy: 0.6783

Epoch 43/100
230/230 [=====] - 6s 24ms/step - loss: 0.5322 - accuracy: 0.8587 - val_loss: 1.2343 - val_accuracy: 0.6903


Epoch 44/100
230/230 [=====] - 8s 37ms/step - loss: 0.4936 - accuracy: 0.8641 - val_loss: 1.2784 - val_accuracy: 0.6794

Epoch 45/100
230/230 [=====] - 5s 24ms/step - loss: 0.5428 - accuracy: 0.8542 - val_loss: 1.2315 - val_accuracy: 0.6908


Epoch 46/100
230/230 [=====] - 6s 27ms/step - loss: 0.4640 - accuracy: 0.8767 - val_loss: 1.3079 - val_accuracy: 0.6685

Epoch 47/100
230/230 [=====] - 7s 30ms/step - loss: 0.4211 - accuracy: 0.8859 - val_loss: 1.2118 - val_accuracy: 0.7137
Epoch 48/100
230/230 [=====] - 6s 24ms/step - loss: 0.3998 - accuracy: 0.8940 - val_loss: 1.2706 - val_accuracy: 0.6897
Epoch 49/100
230/230 [=====] - 8s 34ms/step - loss: 0.4356 - accuracy: 0.8825 - val_loss: 1.2313 - val_accuracy: 0.7088
Epoch 50/100
230/230 [=====] - 6s 25ms/step - loss: 0.3856 - accuracy: 0.8995 - val_loss: 1.2514 - val_accuracy: 0.7109
Total training time: 361.4873790740967 seconds

```
[ ]: # Plot the loss
plt.title('Loss Plot')
plt.plot(history.history['loss'], label='train')
plt.plot(history.history['val_loss'], label='validation')
plt.legend()
plt.show()
```



```
[ ]: # Plot the accuracy
plt.title('Accuracy Plot')
plt.plot(history.history['accuracy'], label='train')
plt.plot(history.history['val_accuracy'], label='validation')
plt.legend()
plt.show()
```



```
[ ]: # Load the saved model
saved_model = tf.keras.models.load_model(model_checkpoint_path)

# Evaluate the model with the test set
loss, accuracy = saved_model.evaluate(X_test_padded, y_test_array)

print("Test Loss:", loss)
print("Test Accuracy:", round((accuracy*100), 2))
```

```
97/97 [=====] - 2s 7ms/step - loss: 1.1920 - accuracy:
0.7286
Test Loss: 1.1920427083969116
Test Accuracy: 72.86
```

```
[ ]: # Check predictions with the test set
y_test_prob = saved_model.predict(X_test_padded)

# Convert probabilities to class labels
y_test_pred = np.argmax(y_test_prob, axis=1)

# Calculate precision, recall, and f1 score
precision = precision_score(y_test_array, y_test_pred, average='weighted')
recall = recall_score(y_test_array, y_test_pred, average='weighted')
f1 = f1_score(y_test_array, y_test_pred, average='weighted')

print("Precision:", round((precision*100), 2))
print("Recall:", round((recall*100), 2))
print("F1 Score:", round((f1*100), 2))
```

```
97/97 [=====] - 1s 6ms/step
Precision: 74.3
Recall: 72.86
F1 Score: 72.14
```

```
[ ]: # Error analysis

# Print classification report
print(classification_report(y_test_array, y_test_pred))

# Check misclassified data
misclassified_data = np.where(y_test_pred != y_test_array)[0]
print(f"The number of misclassifications: {len(misclassified_data)}")

# Check the ratio of misclassifications
misclassification_ratio = (len(misclassified_data) / len(y_test_array)) * 100
# Round the number
rounded_ratio = round(misclassification_ratio, 2)
print(f"Proportion of misclassifications: {rounded_ratio}%")

# Iterate over misclassified data for error analysis
for idx in misclassified_data[:30]:
    input_text = X_test[idx]
    true_label = y_test[idx]
    predicted_label = y_test_pred[idx]

    # Print information about the misclassified data
    print("Input Text:", input_text)
    print("Actual Label:", true_label)
    print("Predicted Label:", predicted_label)
    print()
```

```
precision    recall  f1-score   support
```

0	0.93	0.93	0.93	40
1	0.68	0.75	0.71	40
2	1.00	0.95	0.97	40
3	0.80	0.80	0.80	40
4	0.97	0.85	0.91	40
5	0.58	0.78	0.67	40
6	0.76	0.78	0.77	40
7	0.89	0.78	0.83	40
8	0.65	0.82	0.73	40
9	0.90	0.90	0.90	40
10	0.50	0.40	0.44	40
11	0.65	0.65	0.65	40
12	0.48	0.72	0.58	40
13	0.71	0.85	0.77	40
14	0.63	0.60	0.62	40
15	0.68	0.80	0.74	40
16	0.64	0.62	0.63	40
17	0.74	0.88	0.80	40
18	0.86	0.60	0.71	40
19	0.87	0.82	0.85	40
20	0.59	0.68	0.63	40
21	0.74	0.78	0.76	40
22	0.68	0.65	0.67	40
23	0.81	0.62	0.70	40
24	0.92	0.82	0.87	40
25	0.70	0.65	0.68	40
26	0.58	0.78	0.67	40
27	0.74	0.80	0.77	40
28	0.79	0.78	0.78	40
29	0.70	0.75	0.72	40
30	0.77	0.85	0.81	40
31	0.85	0.85	0.85	40
32	0.94	0.80	0.86	40
33	0.76	0.72	0.74	40
34	0.59	0.72	0.65	40
35	0.69	0.68	0.68	40
36	0.84	0.68	0.75	40
37	0.67	0.55	0.60	40
38	0.75	0.95	0.84	40
39	0.77	0.75	0.76	40
40	0.57	0.97	0.72	40
41	0.67	0.65	0.66	40
42	0.95	0.90	0.92	40
43	0.64	0.72	0.68	40
44	0.97	0.85	0.91	40
45	0.76	0.65	0.70	40
46	0.75	0.75	0.75	40

47	0.65	0.65	0.65	40
48	0.56	0.55	0.56	40
49	0.88	0.70	0.78	40
50	0.70	0.80	0.74	40
51	0.88	0.75	0.81	40
52	0.77	0.68	0.72	40
53	0.66	0.78	0.71	40
54	0.60	0.78	0.67	40
55	0.97	0.85	0.91	40
56	0.88	0.72	0.79	40
57	0.75	0.82	0.79	40
58	0.62	0.70	0.66	40
59	0.75	0.75	0.75	40
60	0.79	0.75	0.77	40
61	0.80	0.70	0.75	40
62	0.68	0.65	0.67	40
63	0.88	0.90	0.89	40
64	0.83	0.75	0.79	40
65	0.69	0.62	0.66	40
66	0.60	0.60	0.60	40
67	0.58	0.62	0.60	40
68	0.00	0.00	0.00	40
69	1.00	0.03	0.05	40
70	0.87	0.82	0.85	40
71	0.89	1.00	0.94	40
72	0.82	0.23	0.35	40
73	0.92	0.82	0.87	40
74	0.37	1.00	0.54	40
75	0.89	0.80	0.84	40
76	0.81	0.65	0.72	40
accuracy				0.73 3080
macro avg				0.74 0.73 0.72 3080
weighted avg				0.74 0.73 0.72 3080

The number of misclassifications: 836

Proportion of misclassifications: 27.14%

Input Text: locate card

Actual Label: 11

Predicted Label: 12

Input Text: way know card arrive

Actual Label: 11

Predicted Label: 12

Input Text: card arrived yet

Actual Label: 11

Predicted Label: 12

Input Text: get card

Actual Label: 11

Predicted Label: 12

Input Text: long card delivery take

Actual Label: 11

Predicted Label: 12

Input Text: still dont card weeks

Actual Label: 11

Predicted Label: 14

Input Text: ive waiting longer expected card

Actual Label: 11

Predicted Label: 14

Input Text: hasnt card delivered

Actual Label: 11

Predicted Label: 12

Input Text: get card yet lost

Actual Label: 11

Predicted Label: 13

Input Text: status card ordered

Actual Label: 11

Predicted Label: 9

Input Text: card arrived yet

Actual Label: 11

Predicted Label: 12

Input Text: know card arrive

Actual Label: 11

Predicted Label: 12

Input Text: im starting think card lost still hasnt arrived help

Actual Label: 11

Predicted Label: 41

Input Text: tracking info available

Actual Label: 11

Predicted Label: 12

Input Text: wont card show app

Actual Label: 13

Predicted Label: 41

Input Text: add card account

Actual Label: 13

Predicted Label: 38

Input Text: put old card back system found

Actual Label: 13

Predicted Label: 30

Input Text: app doesnt show card received

Actual Label: 13

Predicted Label: 41

Input Text: way make old card usable app

Actual Label: 13

Predicted Label: 54

Input Text: could help reactivate card previously lost found morning jacket

Actual Label: 13

Predicted Label: 42

Input Text: often exchange rates change

Actual Label: 32

Predicted Label: 31

Input Text: good time exchange

Actual Label: 32

Predicted Label: 50

Input Text: exchange rate like app

Actual Label: 32

Predicted Label: 17

Input Text: currencies exchange rate calculated

Actual Label: 32

Predicted Label: 31

Input Text: much get exchange rate

Actual Label: 32

Predicted Label: 17

Input Text: exchange rate would

Actual Label: 32

Predicted Label: 17

Input Text: exchange rate like

Actual Label: 32

Predicted Label: 17

Input Text: rate get determined
Actual Label: 32
Predicted Label: 76

Input Text: made currency exchange think charged
Actual Label: 17
Predicted Label: 31

Input Text: charged
Actual Label: 17
Predicted Label: 34

1.4.5 Tuned LSTM (with dropout)

```
[ ]: # Define the output dimension for the embedding layer and hidden units
embedding_output_dim = 100 # Random number
hidden_unit = 40 # Number from hyperparameter tuning
nlabel = 77 # number of classes

# Build the baseline model
dropout_tuned_model = keras.models.Sequential()
dropout_tuned_model.add(layers.Embedding(voca_size, embedding_output_dim))
dropout_tuned_model.add(layers.LSTM(hidden_unit, dropout=0.2)) # Random number
dropout_tuned_model.add(layers.Dense(nlabel, activation='softmax'))

# Compile the model
dropout_tuned_model.compile(optimizer='adam',
    ↪ loss='sparse_categorical_crossentropy', metrics=['accuracy']) # 0.001 is the
    ↪ default of Adam

# Summary the model
dropout_tuned_model.summary()
```

Model: "sequential_5"

Layer (type)	Output Shape	Param #
embedding_5 (Embedding)	(None, None, 100)	208800
lstm_5 (LSTM)	(None, 40)	22560
dense_5 (Dense)	(None, 77)	3157

Total params: 234517 (916.08 KB)
Trainable params: 234517 (916.08 KB)

Non-trainable params: 0 (0.00 Byte)

```
[ ]: # Define the folder path to save the model
folder_path = '/content/drive/MyDrive/1. NLP CW/LSTM/'

# Define the file path for the model checkpoint
model_checkpoint_path = folder_path + 'dropout_tuned_LSTM_embedding_model.keras'

# Define the model checkpoint
mc = tf.keras.callbacks.ModelCheckpoint(
    filepath=model_checkpoint_path,
    monitor='val_accuracy',
    mode='max',
    save_best_only=True)
```

```
[ ]: # Import time to measure the elapsed time
import time

# Measure time before training
start_time = time.time()

# Fit the model
history = dropout_tuned_model.fit(
    X_train_padded, y_train,
    epochs = 100, # Random number
    validation_data = (X_val_padded, y_val),
    callbacks = [mc, es],
    batch_size = 32)

# End the training time
end_time = time.time()

# Measure the training time
training_time = end_time - start_time
print("Total training time:", training_time, "seconds")
```

Epoch 1/100

230/230 [=====] - 12s 39ms/step - loss: 4.2630 - accuracy: 0.0189 - val_loss: 4.1256 - val_accuracy: 0.0212

Epoch 2/100

230/230 [=====] - 7s 31ms/step - loss: 3.9792 - accuracy: 0.0313 - val_loss: 3.9345 - val_accuracy: 0.0299

Epoch 3/100

230/230 [=====] - 7s 32ms/step - loss: 3.9507 - accuracy: 0.0313 - val_loss: 3.9015 - val_accuracy: 0.0305

Epoch 4/100

230/230 [=====] - 6s 26ms/step - loss: 3.9079 -


accuracy: 0.0289 - val_loss: 3.9015 - val_accuracy: 0.0327
 Epoch 5/100
 230/230 [=====] - 8s 34ms/step - loss: 3.8954 -
 accuracy: 0.0294 - val_loss: 3.8779 - val_accuracy: 0.0327
 Epoch 6/100
 230/230 [=====] - 11s 46ms/step - loss: 4.0020 -
 accuracy: 0.0316 - val_loss: 3.9115 - val_accuracy: 0.0343
 Epoch 7/100
 230/230 [=====] - 8s 34ms/step - loss: 3.9081 -
 accuracy: 0.0285 - val_loss: 3.8978 - val_accuracy: 0.0327
 Epoch 8/100
 230/230 [=====] - 5s 24ms/step - loss: 3.8826 -
 accuracy: 0.0313 - val_loss: 3.8469 - val_accuracy: 0.0419
 Epoch 9/100
 230/230 [=====] - 9s 37ms/step - loss: 3.8721 -
 accuracy: 0.0350 - val_loss: 3.8135 - val_accuracy: 0.0414
 Epoch 10/100
 230/230 [=====] - 13s 55ms/step - loss: 3.7973 -
 accuracy: 0.0422 - val_loss: 3.7652 - val_accuracy: 0.0457
 Epoch 11/100
 230/230 [=====] - 10s 44ms/step - loss: 3.6923 -
 accuracy: 0.0505 - val_loss: 3.6624 - val_accuracy: 0.0474
 Epoch 12/100
 230/230 [=====] - 12s 52ms/step - loss: 3.6041 -
 accuracy: 0.0557 - val_loss: 3.5686 - val_accuracy: 0.0729
 Epoch 13/100
 230/230 [=====] - 6s 26ms/step - loss: 3.4120 -
 accuracy: 0.0785 - val_loss: 3.3685 - val_accuracy: 0.0942
 Epoch 14/100
 230/230 [=====] - 8s 35ms/step - loss: 3.1942 -
 accuracy: 0.0939 - val_loss: 3.2385 - val_accuracy: 0.0974
 Epoch 15/100
 230/230 [=====] - 6s 26ms/step - loss: 3.0150 -
 accuracy: 0.1156 - val_loss: 3.0965 - val_accuracy: 0.1241
 Epoch 16/100
 230/230 [=====] - 8s 34ms/step - loss: 2.8881 -
 accuracy: 0.1334 - val_loss: 3.0342 - val_accuracy: 0.1361
 Epoch 17/100
 230/230 [=====] - 6s 26ms/step - loss: 2.7694 -
 accuracy: 0.1538 - val_loss: 2.9277 - val_accuracy: 0.1541
 Epoch 18/100
 230/230 [=====] - 6s 26ms/step - loss: 2.6670 -
 accuracy: 0.1740 - val_loss: 2.8175 - val_accuracy: 0.1807
 Epoch 19/100
 230/230 [=====] - 7s 30ms/step - loss: 2.5498 -
 accuracy: 0.1992 - val_loss: 2.7681 - val_accuracy: 0.1818
 Epoch 20/100
 230/230 [=====] - 6s 27ms/step - loss: 2.4741 -

accuracy: 0.2164 - val_loss: 2.6662 - val_accuracy: 0.2020
 Epoch 21/100
 230/230 [=====] - 8s 36ms/step - loss: 2.3711 -
 accuracy: 0.2389 - val_loss: 2.5793 - val_accuracy: 0.2439
 Epoch 22/100
 230/230 [=====] - 6s 24ms/step - loss: 2.2750 -
 accuracy: 0.2721 - val_loss: 2.5388 - val_accuracy: 0.2330
 Epoch 23/100
 230/230 [=====] - 8s 34ms/step - loss: 2.1622 -
 accuracy: 0.2974 - val_loss: 2.3749 - val_accuracy: 0.2880
 Epoch 24/100
 230/230 [=====] - 6s 26ms/step - loss: 2.0376 -
 accuracy: 0.3366 - val_loss: 2.2583 - val_accuracy: 0.3217
 Epoch 25/100
 230/230 [=====] - 7s 29ms/step - loss: 1.9376 -
 accuracy: 0.3711 - val_loss: 2.1459 - val_accuracy: 0.3636
 Epoch 26/100
 230/230 [=====] - 7s 32ms/step - loss: 1.8185 -
 accuracy: 0.4249 - val_loss: 2.0330 - val_accuracy: 0.4143
 Epoch 27/100
 230/230 [=====] - 6s 27ms/step - loss: 1.6616 -
 accuracy: 0.4781 - val_loss: 1.9790 - val_accuracy: 0.4371
 Epoch 28/100
 230/230 [=====] - 9s 38ms/step - loss: 1.5570 -
 accuracy: 0.5174 - val_loss: 1.8680 - val_accuracy: 0.4823
 Epoch 29/100
 230/230 [=====] - 7s 31ms/step - loss: 1.4374 -
 accuracy: 0.5553 - val_loss: 1.7090 - val_accuracy: 0.5139
 Epoch 30/100
 230/230 [=====] - 8s 36ms/step - loss: 1.2950 -
 accuracy: 0.6055 - val_loss: 1.7371 - val_accuracy: 0.5210
 Epoch 31/100
 230/230 [=====] - 6s 27ms/step - loss: 1.2320 -
 accuracy: 0.6276 - val_loss: 1.6056 - val_accuracy: 0.5700
 Epoch 32/100
 230/230 [=====] - 8s 33ms/step - loss: 1.1562 -
 accuracy: 0.6654 - val_loss: 1.6430 - val_accuracy: 0.5487
 Epoch 33/100
 230/230 [=====] - 6s 27ms/step - loss: 1.0923 -
 accuracy: 0.6835 - val_loss: 1.4850 - val_accuracy: 0.6075
 Epoch 34/100
 230/230 [=====] - 7s 29ms/step - loss: 1.1673 -
 accuracy: 0.6454 - val_loss: 1.4920 - val_accuracy: 0.6026
 Epoch 35/100
 230/230 [=====] - 7s 31ms/step - loss: 0.9790 -
 accuracy: 0.7133 - val_loss: 1.4176 - val_accuracy: 0.6217
 Epoch 36/100
 230/230 [=====] - 6s 25ms/step - loss: 0.9349 -

accuracy: 0.7284 - val_loss: 1.4201 - val_accuracy: 0.6309
 Epoch 37/100
 230/230 [=====] - 8s 37ms/step - loss: 0.8739 -
 accuracy: 0.7483 - val_loss: 1.3790 - val_accuracy: 0.6364
 Epoch 38/100
 230/230 [=====] - 6s 26ms/step - loss: 0.8562 -
 accuracy: 0.7592 - val_loss: 1.3934 - val_accuracy: 0.6434
 Epoch 39/100
 230/230 [=====] - 8s 34ms/step - loss: 0.7994 -
 accuracy: 0.7736 - val_loss: 1.2857 - val_accuracy: 0.6734
 Epoch 40/100
 230/230 [=====] - 6s 26ms/step - loss: 0.7100 -
 accuracy: 0.8037 - val_loss: 1.2523 - val_accuracy: 0.6946
 Epoch 41/100
 230/230 [=====] - 6s 28ms/step - loss: 0.6713 -
 accuracy: 0.8121 - val_loss: 1.2592 - val_accuracy: 0.6903
 Epoch 42/100
 230/230 [=====] - 8s 33ms/step - loss: 0.6246 -
 accuracy: 0.8307 - val_loss: 1.2244 - val_accuracy: 0.7050
 Epoch 43/100
 230/230 [=====] - 6s 28ms/step - loss: 0.5993 -
 accuracy: 0.8337 - val_loss: 1.2247 - val_accuracy: 0.7137
 Epoch 44/100
 230/230 [=====] - 9s 37ms/step - loss: 0.5454 -
 accuracy: 0.8527 - val_loss: 1.1916 - val_accuracy: 0.7137
 Epoch 45/100
 230/230 [=====] - 6s 27ms/step - loss: 0.5321 -
 accuracy: 0.8545 - val_loss: 1.1619 - val_accuracy: 0.7213
 Epoch 46/100
 230/230 [=====] - 9s 37ms/step - loss: 0.5138 -
 accuracy: 0.8616 - val_loss: 1.1556 - val_accuracy: 0.7235
 Epoch 47/100
 230/230 [=====] - 6s 26ms/step - loss: 0.5126 -
 accuracy: 0.8579 - val_loss: 1.1608 - val_accuracy: 0.7289
 Epoch 48/100
 230/230 [=====] - 8s 36ms/step - loss: 0.4501 -
 accuracy: 0.8816 - val_loss: 1.1435 - val_accuracy: 0.7295
 Epoch 49/100
 230/230 [=====] - 6s 28ms/step - loss: 0.4402 -
 accuracy: 0.8805 - val_loss: 1.1190 - val_accuracy: 0.7305
 Epoch 50/100
 230/230 [=====] - 7s 31ms/step - loss: 0.4019 -
 accuracy: 0.8900 - val_loss: 1.1093 - val_accuracy: 0.7382
 Epoch 51/100
 230/230 [=====] - 7s 29ms/step - loss: 0.3875 -
 accuracy: 0.8908 - val_loss: 1.0944 - val_accuracy: 0.7441
 Epoch 52/100
 230/230 [=====] - 6s 27ms/step - loss: 0.3842 -

```
accuracy: 0.8950 - val_loss: 1.0782 - val_accuracy: 0.7441
Epoch 53/100
230/230 [=====] - 8s 35ms/step - loss: 0.3607 -
accuracy: 0.9035 - val_loss: 1.1300 - val_accuracy: 0.7322
Epoch 54/100
230/230 [=====] - 6s 27ms/step - loss: 0.3953 -
accuracy: 0.8937 - val_loss: 1.0888 - val_accuracy: 0.7469
Epoch 55/100
230/230 [=====] - 9s 37ms/step - loss: 0.3558 -
accuracy: 0.9055 - val_loss: 1.0863 - val_accuracy: 0.7392
Total training time: 408.2757978439331 seconds
```


```
[ ]: # Plot the loss
plt.title('Loss Plot')
plt.plot(history.history['loss'], label='train')
plt.plot(history.history['val_loss'], label='validation')
plt.legend()
plt.show()
```



```
[ ]: # Plot the accuracy
plt.title('Accuracy Plot')
```



```
plt.plot(history.history['accuracy'], label='train')
plt.plot(history.history['val_accuracy'], label='validation')
plt.legend()
plt.show()
```



```
[ ]: # Load the saved model
saved_model = tf.keras.models.load_model(model_checkpoint_path)

# Evaluate the model with the test set
loss, accuracy = saved_model.evaluate(X_test_padded, y_test_array)

print("Test Loss:", loss)
print("Test Accuracy:", round((accuracy*100), 2))
```

```
97/97 [=====] - 2s 7ms/step - loss: 1.0432 - accuracy:
0.7659
Test Loss: 1.0432207584381104
Test Accuracy: 76.59
```

```
[ ]: # Check predictions with the test set
y_test_prob = saved_model.predict(X_test_padded)
```

```

# Convert probabilities to class labels
y_test_pred = np.argmax(y_test_prob, axis=1)

# Calculate precision, recall, and f1 score
precision = precision_score(y_test_array, y_test_pred, average='weighted')
recall = recall_score(y_test_array, y_test_pred, average='weighted')
f1 = f1_score(y_test_array, y_test_pred, average='weighted')

print("Precision:", round((precision*100), 2))
print("Recall:", round((recall*100), 2))
print("F1 Score:", round((f1*100), 2))

```

```

97/97 [=====] - 1s 7ms/step
Precision: 77.65
Recall: 76.59
F1 Score: 76.47

```

```

[ ]: # Error analysis

# Print classification report
print(classification_report(y_test_array, y_test_pred))

# Check misclassified data
misclassified_data = np.where(y_test_pred != y_test_array)[0]
print(f"The number of misclassifications: {len(misclassified_data)}")

# Check the ratio of misclassifications
misclassification_ratio = (len(misclassified_data) / len(y_test_array)) * 100
# Round the number
rounded_ratio = round(misclassification_ratio, 2)
print(f"Proportion of misclassifications: {rounded_ratio}%")

# Iterate over misclassified data for error analysis
for idx in misclassified_data[:30]:
    input_text = X_test[idx]
    true_label = y_test[idx]
    predicted_label = y_test_pred[idx]

    # Print information about the misclassified data
    print("Input Text:", input_text)
    print("Actual Label:", true_label)
    print("Predicted Label:", predicted_label)
    print()

```

	precision	recall	f1-score	support
0	0.95	0.95	0.95	40

1	1.00	0.95	0.97	40
2	0.97	0.97	0.97	40
3	0.92	0.60	0.73	40
4	0.91	0.80	0.85	40
5	0.55	0.78	0.65	40
6	0.85	0.85	0.85	40
7	0.82	0.78	0.79	40
8	0.87	0.82	0.85	40
9	0.97	0.88	0.92	40
10	0.58	0.45	0.51	40
11	0.85	0.82	0.84	40
12	0.56	0.62	0.59	40
13	0.90	0.93	0.91	40
14	0.77	0.85	0.81	40
15	0.72	0.78	0.75	40
16	0.59	0.55	0.57	40
17	0.85	0.88	0.86	40
18	0.75	0.82	0.79	40
19	0.80	0.82	0.81	40
20	0.67	0.65	0.66	40
21	0.97	0.78	0.86	40
22	0.73	0.55	0.63	40
23	0.97	0.85	0.91	40
24	0.97	0.80	0.88	40
25	0.57	0.78	0.66	40
26	0.64	0.80	0.71	40
27	0.76	0.78	0.77	40
28	0.88	0.75	0.81	40
29	0.74	0.78	0.76	40
30	0.95	0.93	0.94	40
31	0.77	0.75	0.76	40
32	0.88	0.88	0.88	40
33	0.74	0.72	0.73	40
34	0.79	0.68	0.73	40
35	0.71	0.72	0.72	40
36	0.81	0.75	0.78	40
37	0.81	0.62	0.70	40
38	0.82	1.00	0.90	40
39	0.64	0.85	0.73	40
40	0.81	0.95	0.87	40
41	0.69	0.68	0.68	40
42	0.90	0.88	0.89	40
43	0.55	0.70	0.62	40
44	0.98	1.00	0.99	40
45	0.76	0.78	0.77	40
46	0.73	0.82	0.78	40
47	0.63	0.68	0.65	40
48	0.79	0.57	0.67	40

	49	0.83	0.75	0.79	40
	50	0.89	0.85	0.87	40
	51	0.70	0.80	0.74	40
	52	0.74	0.72	0.73	40
	53	0.64	0.85	0.73	40
	54	0.62	0.78	0.69	40
	55	0.89	0.82	0.86	40
	56	0.78	0.70	0.74	40
	57	0.95	0.93	0.94	40
	58	0.85	0.70	0.77	40
	59	0.78	0.88	0.82	40
	60	0.76	0.88	0.81	40
	61	0.82	0.68	0.74	40
	62	0.74	0.72	0.73	40
	63	0.62	0.82	0.71	40
	64	0.76	0.78	0.77	40
	65	0.62	0.57	0.60	40
	66	0.69	0.55	0.61	40
	67	0.66	0.68	0.67	40
	68	0.71	0.38	0.49	40
	69	0.50	0.20	0.29	40
	70	0.92	0.90	0.91	40
	71	0.93	0.97	0.95	40
	72	0.81	0.65	0.72	40
	73	0.78	0.88	0.82	40
	74	0.43	0.85	0.57	40
	75	0.82	0.70	0.76	40
	76	0.68	0.65	0.67	40
	accuracy			0.77	3080
	macro avg	0.78	0.77	0.76	3080
	weighted avg	0.78	0.77	0.76	3080

The number of misclassifications: 721
Proportion of misclassifications: 23.41%
Input Text: get card
Actual Label: 11
Predicted Label: 43

Input Text: long card delivery take
Actual Label: 11
Predicted Label: 12

Input Text: get card yet lost
Actual Label: 11
Predicted Label: 41

Input Text: status card ordered

Actual Label: 11

Predicted Label: 53

Input Text: long new card take arrive

Actual Label: 11

Predicted Label: 12

Input Text: think something went wrong card delivery havent received yet

Actual Label: 11

Predicted Label: 12

Input Text: im starting think card lost still hasnt arrived help

Actual Label: 11

Predicted Label: 18

Input Text: add card account

Actual Label: 13

Predicted Label: 43

Input Text: link credit card

Actual Label: 13

Predicted Label: 47

Input Text: way make old card usable app

Actual Label: 13

Predicted Label: 61

Input Text: good time exchange

Actual Label: 32

Predicted Label: 50

Input Text: exchange rate like app

Actual Label: 32

Predicted Label: 76

Input Text: exchange rate use

Actual Label: 32

Predicted Label: 76

Input Text: much get exchange rate

Actual Label: 32

Predicted Label: 76

Input Text: kind foreign exchange rate get exchange money

Actual Label: 32

Predicted Label: 76

Input Text: made currency exchange think charged

Actual Label: 17

Predicted Label: 31

Input Text: rate low sure using right exchange rate

Actual Label: 17

Predicted Label: 76

Input Text: charged

Actual Label: 17

Predicted Label: 15

Input Text: hi dont think exchange rate right need check official interbank
exchange please

Actual Label: 17

Predicted Label: 76

Input Text: conversion value card payments incorrect

Actual Label: 17

Predicted Label: 15

Input Text: im okay fee statement

Actual Label: 34

Predicted Label: 63

Input Text: would like refund extra pound charged

Actual Label: 34

Predicted Label: 63

Input Text: statement extra charges

Actual Label: 34

Predicted Label: 63

Input Text: transaction credited

Actual Label: 34

Predicted Label: 8

Input Text: fee come

Actual Label: 34

Predicted Label: 15

Input Text: many fees statement

Actual Label: 34

Predicted Label: 15

Input Text: euro fee come

Actual Label: 34

Predicted Label: 2

Input Text: euro fee statement
Actual Label: 34
Predicted Label: 15

Input Text: extra charge app told aware
Actual Label: 34
Predicted Label: 16

Input Text: new customer happened look app charge familiar could tell extra charge
Actual Label: 34
Predicted Label: 15

The LSTM with baseline architecture and dropout has the best performance.

1.5 LSTM (with Word2Vec)

1.5.1 Set up the Word2Vec model

```
[ ]: # Import and install the library and file for Word2Vec
import gensim

!pip install gdown # Install google download
!gdown https://drive.google.com/uc?id=1Av37IVBQAAntSe1X3M0A15gvowQzd2_j #
↳ Download the Word2Vec (GoogleNews-vectors-negative300.bin.gz)

# Define the Word2Vec model
word2vec_model = gensim.models.KeyedVectors.
↳ load_word2vec_format('GoogleNews-vectors-negative300.bin.gz', binary=True)
```

Requirement already satisfied: gdown in /usr/local/lib/python3.10/dist-packages (5.1.0)

Requirement already satisfied: beautifulsoup4 in /usr/local/lib/python3.10/dist-packages (from gdown) (4.12.3)

Requirement already satisfied: filelock in /usr/local/lib/python3.10/dist-packages (from gdown) (3.14.0)

Requirement already satisfied: requests[socks] in /usr/local/lib/python3.10/dist-packages (from gdown) (2.31.0)

Requirement already satisfied: tqdm in /usr/local/lib/python3.10/dist-packages (from gdown) (4.66.4)

Requirement already satisfied: soupsieve>1.2 in /usr/local/lib/python3.10/dist-packages (from beautifulsoup4->gdown) (2.5)

Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.10/dist-packages (from requests[socks]->gdown) (3.3.2)

Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (from requests[socks]->gdown) (3.7)

Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.10/dist-packages (from requests[socks]->gdown) (2.0.7)

```
Requirement already satisfied: certifi>=2017.4.17 in
/usr/local/lib/python3.10/dist-packages (from requests[socks]->gdown) (2024.2.2)
Requirement already satisfied: PySocks!=1.5.7,>=1.5.6 in
/usr/local/lib/python3.10/dist-packages (from requests[socks]->gdown) (1.7.1)
Downloading...
From (original):
https://drive.google.com/uc?id=1Av37IVBQAAntSe1X3MOA15gvowQzd2_j
From (redirected): https://drive.google.com/uc?id=1Av37IVBQAAntSe1X3MOA15gvowQzd
2_j&confirm=t&uuid=4335f280-bffc-48d7-a335-0876702f458b
To: /content/GoogleNews-vectors-negative300.bin.gz
100% 1.65G/1.65G [00:18<00:00, 87.9MB/s]
```

```
[ ]: # Check the model's dimension
print(f'Word2Vec: {word2vec_model.vectors.shape}')
```

Word2Vec: (3000000, 300)

```
[ ]: # Define the embedding matrix
embedding_matrix = np.zeros((voca_size+1, 300)) # For future dimension matching
↳ with the word_index, add 1 to the vocabulary size, and match 300 from
↳ Word2Vec
print(f'The shape of embedding matrix: {np.shape(embedding_matrix)}')
```

The shape of embedding matrix: (2089, 300)

```
[ ]: # Match words in the word_index to those in the word2vec_model for creating
↳ embedding matrix indices
def extract_vector(word):
    if word in word2vec_model:
        return word2vec_model[word]
    else:
        return None
```

```
[ ]: # Match the index and word to create an embedding matrix
for word, index in word_index.items():
    vector_value = extract_vector(word)
    if vector_value is not None:
        embedding_matrix[index] = vector_value
```

```
[ ]: # Check the word 'card' vectors in the Word2Vec model
print(word2vec_model['card'])
```

```
[-1.63085938e-01  1.43554688e-01  1.97265625e-01  1.57226562e-01
 4.12597656e-02  2.43164062e-01 -2.21679688e-01 -5.02929688e-02
-4.88281250e-02 -3.45703125e-01  4.18090820e-03 -8.10546875e-02
-2.00195312e-01 -2.27539062e-01  4.41894531e-02 -1.96533203e-02
 2.83203125e-01  9.52148438e-02  8.74023438e-02 -2.69531250e-01
 2.27539062e-01  5.34667969e-02 -1.63574219e-02  3.97949219e-02]
```


-2.97851562e-02	-1.83105469e-02	-8.64257812e-02	9.96093750e-02
3.33984375e-01	1.70898438e-02	-4.88281250e-02	-9.57031250e-02
2.73437500e-01	1.33789062e-01	-5.90820312e-02	6.93359375e-02
4.51660156e-02	1.31835938e-01	-1.74804688e-01	-2.12402344e-02
-1.53198242e-02	8.78906250e-03	3.78906250e-01	-1.36718750e-01
2.13867188e-01	-2.43164062e-01	2.70996094e-02	4.88281250e-02
-3.56445312e-02	-3.19824219e-02	-3.90625000e-01	2.40234375e-01
9.96093750e-02	-1.36718750e-01	-1.52343750e-01	3.14941406e-02
-1.62109375e-01	-1.44531250e-01	2.33398438e-01	-1.55273438e-01
1.92382812e-01	1.95312500e-01	-1.94335938e-01	1.54296875e-01
-1.16210938e-01	1.23535156e-01	-2.98828125e-01	-1.29882812e-01
-1.09375000e-01	1.28906250e-01	2.08007812e-01	1.16699219e-01
4.12109375e-01	-2.24609375e-01	1.53320312e-01	-6.73828125e-02
1.68945312e-01	1.11328125e-01	2.39257812e-01	2.81250000e-01
-1.22680664e-02	-1.40625000e-01	-9.17968750e-02	-3.32031250e-02
1.07421875e-01	-2.16064453e-02	-4.80957031e-02	-1.40625000e-01
-1.13281250e-01	1.34277344e-02	-1.46484375e-02	1.48925781e-02
-2.71484375e-01	-3.75000000e-01	-2.79296875e-01	-4.10156250e-02
-9.96093750e-02	3.28125000e-01	-1.78222656e-02	-2.13867188e-01
8.44726562e-02	-1.63085938e-01	2.31445312e-01	1.75781250e-01
-1.55273438e-01	-1.49414062e-01	3.55468750e-01	-4.62890625e-01
1.04980469e-01	-1.07910156e-01	7.32421875e-02	-1.77734375e-01
-5.98144531e-02	-6.98852539e-03	-1.87500000e-01	1.84570312e-01
1.28906250e-01	1.04003906e-01	5.03906250e-01	1.36718750e-01
-3.17382812e-02	8.54492188e-03	1.67236328e-02	2.28515625e-01
1.17797852e-02	-2.94921875e-01	2.04467773e-03	2.65625000e-01
-1.51367188e-01	-9.81445312e-02	1.82617188e-01	-2.08984375e-01
-2.83203125e-01	-3.71093750e-01	-1.21307373e-03	8.69140625e-02
-5.71289062e-02	1.93359375e-01	-2.79296875e-01	-3.71093750e-02
-3.98437500e-01	-2.40234375e-01	4.41894531e-02	3.49121094e-02
1.04980469e-01	-7.71484375e-02	1.12792969e-01	-6.29882812e-02
-1.68945312e-01	-7.56835938e-02	8.44726562e-02	2.63671875e-01
-4.17480469e-02	2.24609375e-02	1.41601562e-01	-7.81250000e-02
1.23046875e-01	2.51464844e-02	4.51660156e-03	5.12695312e-02
9.88769531e-03	-2.63671875e-01	3.45703125e-01	-8.34960938e-02
-1.04980469e-01	1.44531250e-01	1.71875000e-01	-3.96484375e-01
-3.35937500e-01	-1.33056641e-02	-3.68652344e-02	-3.28125000e-01
-1.54418945e-02	2.70996094e-02	-2.61718750e-01	1.87988281e-02
6.54296875e-02	-9.03320312e-02	2.41210938e-01	1.03515625e-01
-1.40625000e-01	-1.57226562e-01	-3.00781250e-01	3.22265625e-02
1.05957031e-01	-2.61718750e-01	-1.52343750e-01	4.73632812e-02
-4.37500000e-01	4.41406250e-01	-1.20605469e-01	-1.25000000e-01
-3.39843750e-01	-3.30078125e-01	2.89062500e-01	1.18164062e-01
-1.06933594e-01	1.22680664e-02	-2.65625000e-01	-2.02148438e-01
-1.29882812e-01	1.98242188e-01	-7.66601562e-02	1.31835938e-01
1.89453125e-01	-8.20312500e-02	-1.66992188e-01	1.67968750e-01
-3.88183594e-02	8.05664062e-02	-2.53906250e-02	2.73437500e-01
-9.86328125e-02	1.52343750e-01	1.07910156e-01	3.65234375e-01

```

-3.75366211e-03  1.23023987e-04 -2.49023438e-01  5.88378906e-02
 7.27539062e-02  7.03125000e-02  8.49609375e-02 -1.63085938e-01
-2.23632812e-01  8.88671875e-02 -1.00097656e-01  3.54003906e-02
-1.42822266e-02 -2.35351562e-01 -3.06640625e-01  3.23486328e-03
 2.45117188e-01 -8.74023438e-02 -2.86865234e-02 -4.12597656e-02
 9.27734375e-02 -1.20849609e-02  1.43554688e-01 -6.39648438e-02
 2.17773438e-01 -5.98144531e-02 -6.17675781e-02 -3.37890625e-01
-2.99072266e-03 -4.27246094e-02  2.67333984e-02  3.56445312e-02
-6.31713867e-03  7.12890625e-02  1.03515625e-01 -6.20117188e-02
 1.52587891e-03 -4.21875000e-01 -8.48388672e-03  1.31835938e-01
 2.64892578e-02  1.58203125e-01  2.11181641e-02 -5.05371094e-02
 1.15234375e-01  4.46777344e-02 -1.75781250e-01 -3.06640625e-01
-1.51367188e-01 -1.09375000e-01 -1.50390625e-01  7.91015625e-02
-1.36718750e-01  5.00488281e-02 -2.23632812e-01 -8.98437500e-02
-2.81250000e-01  2.13867188e-01  5.20019531e-02  3.32031250e-02
 1.87500000e-01 -2.50000000e-01 -1.50390625e-01  3.76953125e-01
 1.29882812e-01  1.48010254e-03 -9.91210938e-02  1.59179688e-01
-1.65039062e-01 -1.15722656e-01  8.20312500e-02  8.93554688e-02
 1.38671875e-01  1.38549805e-02  1.08032227e-02  1.62109375e-01
-9.86328125e-02 -5.02929688e-02  2.18505859e-02 -1.29882812e-01
-4.68750000e-02 -1.04492188e-01 -1.25000000e-01  1.13281250e-01]

```

```

[ ]: # Check the index of the word 'card' in word_index
print(f"The index of card in word_index: {word_index['card']}")

```

The index of card in word_index: 1

```

[ ]: # Check the word 'card' vectors in the embedding matrix
print(embedding_matrix[1])

```

```

[-1.63085938e-01  1.43554688e-01  1.97265625e-01  1.57226562e-01
 4.12597656e-02  2.43164062e-01 -2.21679688e-01 -5.02929688e-02
-4.88281250e-02 -3.45703125e-01  4.18090820e-03 -8.10546875e-02
-2.00195312e-01 -2.27539062e-01  4.41894531e-02 -1.96533203e-02
 2.83203125e-01  9.52148438e-02  8.74023438e-02 -2.69531250e-01
 2.27539062e-01  5.34667969e-02 -1.63574219e-02  3.97949219e-02
-2.97851562e-02 -1.83105469e-02 -8.64257812e-02  9.96093750e-02
 3.33984375e-01  1.70898438e-02 -4.88281250e-02 -9.57031250e-02
 2.73437500e-01  1.33789062e-01 -5.90820312e-02  6.93359375e-02
 4.51660156e-02  1.31835938e-01 -1.74804688e-01 -2.12402344e-02
-1.53198242e-02  8.78906250e-03  3.78906250e-01 -1.36718750e-01
 2.13867188e-01 -2.43164062e-01  2.70996094e-02  4.88281250e-02
-3.56445312e-02 -3.19824219e-02 -3.90625000e-01  2.40234375e-01
 9.96093750e-02 -1.36718750e-01 -1.52343750e-01  3.14941406e-02
-1.62109375e-01 -1.44531250e-01  2.33398438e-01 -1.55273438e-01
 1.92382812e-01  1.95312500e-01 -1.94335938e-01  1.54296875e-01
-1.16210938e-01  1.23535156e-01 -2.98828125e-01 -1.29882812e-01
-1.09375000e-01  1.28906250e-01  2.08007812e-01  1.16699219e-01]

```

4.12109375e-01 -2.24609375e-01 1.53320312e-01 -6.73828125e-02
 1.68945312e-01 1.11328125e-01 2.39257812e-01 2.81250000e-01
 -1.22680664e-02 -1.40625000e-01 -9.17968750e-02 -3.32031250e-02
 1.07421875e-01 -2.16064453e-02 -4.80957031e-02 -1.40625000e-01
 -1.13281250e-01 1.34277344e-02 -1.46484375e-02 1.48925781e-02
 -2.71484375e-01 -3.75000000e-01 -2.79296875e-01 -4.10156250e-02
 -9.96093750e-02 3.28125000e-01 -1.78222656e-02 -2.13867188e-01
 8.44726562e-02 -1.63085938e-01 2.31445312e-01 1.75781250e-01
 -1.55273438e-01 -1.49414062e-01 3.55468750e-01 -4.62890625e-01
 1.04980469e-01 -1.07910156e-01 7.32421875e-02 -1.77734375e-01
 -5.98144531e-02 -6.98852539e-03 -1.87500000e-01 1.84570312e-01
 1.28906250e-01 1.04003906e-01 5.03906250e-01 1.36718750e-01
 -3.17382812e-02 8.54492188e-03 1.67236328e-02 2.28515625e-01
 1.17797852e-02 -2.94921875e-01 2.04467773e-03 2.65625000e-01
 -1.51367188e-01 -9.81445312e-02 1.82617188e-01 -2.08984375e-01
 -2.83203125e-01 -3.71093750e-01 -1.21307373e-03 8.69140625e-02
 -5.71289062e-02 1.93359375e-01 -2.79296875e-01 -3.71093750e-02
 -3.98437500e-01 -2.40234375e-01 4.41894531e-02 3.49121094e-02
 1.04980469e-01 -7.71484375e-02 1.12792969e-01 -6.29882812e-02
 -1.68945312e-01 -7.56835938e-02 8.44726562e-02 2.63671875e-01
 -4.17480469e-02 2.24609375e-02 1.41601562e-01 -7.81250000e-02
 1.23046875e-01 2.51464844e-02 4.51660156e-03 5.12695312e-02
 9.88769531e-03 -2.63671875e-01 3.45703125e-01 -8.34960938e-02
 -1.04980469e-01 1.44531250e-01 1.71875000e-01 -3.96484375e-01
 -3.35937500e-01 -1.33056641e-02 -3.68652344e-02 -3.28125000e-01
 -1.54418945e-02 2.70996094e-02 -2.61718750e-01 1.87988281e-02
 6.54296875e-02 -9.03320312e-02 2.41210938e-01 1.03515625e-01
 -1.40625000e-01 -1.57226562e-01 -3.00781250e-01 3.22265625e-02
 1.05957031e-01 -2.61718750e-01 -1.52343750e-01 4.73632812e-02
 -4.37500000e-01 4.41406250e-01 -1.20605469e-01 -1.25000000e-01
 -3.39843750e-01 -3.30078125e-01 2.89062500e-01 1.18164062e-01
 -1.06933594e-01 1.22680664e-02 -2.65625000e-01 -2.02148438e-01
 -1.29882812e-01 1.98242188e-01 -7.66601562e-02 1.31835938e-01
 1.89453125e-01 -8.20312500e-02 -1.66992188e-01 1.67968750e-01
 -3.88183594e-02 8.05664062e-02 -2.53906250e-02 2.73437500e-01
 -9.86328125e-02 1.52343750e-01 1.07910156e-01 3.65234375e-01
 -3.75366211e-03 1.23023987e-04 -2.49023438e-01 5.88378906e-02
 7.27539062e-02 7.03125000e-02 8.49609375e-02 -1.63085938e-01
 -2.23632812e-01 8.88671875e-02 -1.00097656e-01 3.54003906e-02
 -1.42822266e-02 -2.35351562e-01 -3.06640625e-01 3.23486328e-03
 2.45117188e-01 -8.74023438e-02 -2.86865234e-02 -4.12597656e-02
 9.27734375e-02 -1.20849609e-02 1.43554688e-01 -6.39648438e-02
 2.17773438e-01 -5.98144531e-02 -6.17675781e-02 -3.37890625e-01
 -2.99072266e-03 -4.27246094e-02 2.67333984e-02 3.56445312e-02
 -6.31713867e-03 7.12890625e-02 1.03515625e-01 -6.20117188e-02
 1.52587891e-03 -4.21875000e-01 -8.48388672e-03 1.31835938e-01
 2.64892578e-02 1.58203125e-01 2.11181641e-02 -5.05371094e-02
 1.15234375e-01 4.46777344e-02 -1.75781250e-01 -3.06640625e-01

```

-1.51367188e-01 -1.09375000e-01 -1.50390625e-01 7.91015625e-02
-1.36718750e-01 5.00488281e-02 -2.23632812e-01 -8.98437500e-02
-2.81250000e-01 2.13867188e-01 5.20019531e-02 3.32031250e-02
1.87500000e-01 -2.50000000e-01 -1.50390625e-01 3.76953125e-01
1.29882812e-01 1.48010254e-03 -9.91210938e-02 1.59179688e-01
-1.65039062e-01 -1.15722656e-01 8.20312500e-02 8.93554688e-02
1.38671875e-01 1.38549805e-02 1.08032227e-02 1.62109375e-01
-9.86328125e-02 -5.02929688e-02 2.18505859e-02 -1.29882812e-01
-4.68750000e-02 -1.04492188e-01 -1.25000000e-01 1.13281250e-01]

```

```

[ ]: # Check the word 'topup' vectors in the Word2Vec model
print(word2vec_model['topup'])

```

```

[ 5.17578125e-02 -1.61132812e-01 -1.33789062e-01 1.93359375e-01
-1.39770508e-02 1.69921875e-01 1.06933594e-01 -2.08007812e-01
7.81250000e-02 1.11816406e-01 3.82995605e-03 2.55126953e-02
-2.75390625e-01 -9.71679688e-02 -2.84423828e-02 1.31835938e-01
2.27539062e-01 2.29492188e-02 1.18652344e-01 6.03027344e-02
-1.92871094e-02 -6.15234375e-02 1.73828125e-01 1.72851562e-01
2.50244141e-02 -1.54296875e-01 -1.62109375e-01 1.02539062e-01
1.50146484e-02 -7.42187500e-02 -2.60009766e-02 4.98046875e-02
-3.32031250e-02 -1.01074219e-01 -5.34667969e-02 -2.45117188e-01
-1.04003906e-01 -1.25976562e-01 -7.86132812e-02 6.44531250e-02
-2.86865234e-02 -3.44238281e-02 3.75976562e-02 -1.87500000e-01
-4.78515625e-02 -3.53515625e-01 -1.01318359e-02 3.97949219e-02
-1.40625000e-01 -1.69921875e-01 3.11279297e-02 -5.27343750e-02
1.09863281e-01 -7.35473633e-03 -4.27246094e-02 3.00292969e-02
-3.08593750e-01 1.87988281e-02 2.20947266e-02 -1.17675781e-01
-1.70898438e-01 -2.04101562e-01 -7.51953125e-02 1.94091797e-02
-1.42578125e-01 -1.42578125e-01 -2.14843750e-01 7.71484375e-02
-3.09753418e-03 1.10839844e-01 -1.66015625e-02 -2.01416016e-02
1.69921875e-01 -9.91210938e-02 -4.76074219e-02 -2.30468750e-01
1.49414062e-01 1.42578125e-01 -2.25830078e-02 2.06298828e-02
-7.95898438e-02 1.29882812e-01 3.51562500e-02 2.03125000e-01
-3.39355469e-02 7.12890625e-02 -4.41894531e-02 3.80859375e-02
-3.80859375e-02 2.55859375e-01 -8.74023438e-02 -2.44140625e-02
-1.21093750e-01 -1.15722656e-01 -1.11694336e-02 1.20117188e-01
-7.03125000e-02 -7.12890625e-02 3.30078125e-01 4.63867188e-02
-5.00488281e-02 -7.03125000e-02 -7.91015625e-02 1.96289062e-01
-1.24511719e-01 -7.03125000e-02 1.53320312e-01 5.02929688e-02
2.69531250e-01 1.39160156e-02 -7.22656250e-02 -1.26953125e-01
-2.19726562e-01 3.78417969e-02 -3.80859375e-02 1.30859375e-01
-1.10351562e-01 -1.52343750e-01 1.64062500e-01 -7.91015625e-02
2.05078125e-01 -1.33789062e-01 -3.05175781e-02 1.25976562e-01
9.52148438e-02 9.57031250e-02 -2.29492188e-01 1.22558594e-01
1.14257812e-01 -4.95605469e-02 -7.42187500e-02 1.01074219e-01
-1.31835938e-01 -3.12500000e-01 -1.01562500e-01 -5.81054688e-02
1.30859375e-01 2.13623047e-02 -1.19140625e-01 2.65625000e-01]

```

```

5.34667969e-02 -1.57226562e-01 -1.21593475e-04 1.58203125e-01
9.27734375e-02 -4.36401367e-03 -9.91210938e-02 -6.34765625e-02
-1.09375000e-01 -2.30712891e-02 1.40625000e-01 -1.58203125e-01
-1.31835938e-01 -2.61718750e-01 1.78710938e-01 -1.61132812e-01
-1.74560547e-02 9.66796875e-02 -2.59765625e-01 5.73730469e-03
-1.40625000e-01 -5.61523438e-02 -1.52343750e-01 -2.17285156e-02
-1.05957031e-01 -3.75976562e-02 1.26953125e-01 -2.38037109e-03
-2.61718750e-01 -1.48437500e-01 -1.57226562e-01 -9.13085938e-02
7.47070312e-02 -1.49414062e-01 -9.57031250e-02 7.51953125e-02
2.38281250e-01 -1.79687500e-01 -1.28906250e-01 5.63964844e-02
-1.40625000e-01 -1.51367188e-01 2.74658203e-03 9.22851562e-02
8.64257812e-02 -1.16210938e-01 -5.54199219e-02 1.99218750e-01
-3.10546875e-01 7.85827637e-04 -2.13867188e-01 -2.08007812e-01
-1.66015625e-01 -1.18164062e-01 -1.77001953e-02 5.59082031e-02
1.17187500e-02 -3.61328125e-02 -8.34960938e-02 -9.37500000e-02
1.06445312e-01 1.48437500e-01 -1.23046875e-01 1.11816406e-01
3.58886719e-02 9.57031250e-02 -3.11279297e-02 -1.01562500e-01
-6.73828125e-02 4.30297852e-03 1.97265625e-01 1.59179688e-01
3.61328125e-02 1.09375000e-01 -1.45507812e-01 1.11328125e-01
2.24609375e-02 -1.13769531e-01 -1.79687500e-01 2.50244141e-02
-5.88378906e-02 2.44140625e-02 2.57568359e-02 -9.33837891e-03
-6.78710938e-02 1.73828125e-01 -1.55273438e-01 2.70996094e-02
-2.50244141e-02 -9.37500000e-02 -3.36914062e-02 -7.27539062e-02
7.95898438e-02 -4.15039062e-02 5.71289062e-02 6.93359375e-02
-2.06298828e-02 1.57226562e-01 -3.39355469e-02 1.02539062e-01
7.86132812e-02 2.85644531e-02 5.59082031e-02 3.44238281e-02
6.98242188e-02 -1.04370117e-02 2.89306641e-02 1.53320312e-01
1.12304688e-01 -9.22851562e-02 -3.85742188e-02 1.22558594e-01
-1.28906250e-01 1.54296875e-01 -3.58581543e-03 1.07421875e-01
3.24707031e-02 -3.32031250e-02 -1.87988281e-02 3.41796875e-02
2.07031250e-01 -1.81640625e-01 -1.67968750e-01 2.14843750e-02
-1.84570312e-01 -9.71679688e-02 -1.19140625e-01 -5.59082031e-02
-3.27148438e-02 -4.76074219e-03 -6.88476562e-02 -4.71191406e-02
-1.16699219e-01 1.67968750e-01 1.00097656e-01 -1.84326172e-02
-1.53198242e-02 -7.42187500e-02 -2.50244141e-02 6.83593750e-02
-1.05957031e-01 2.25585938e-01 9.91821289e-04 -4.07714844e-02
1.09863281e-01 1.57165527e-03 2.80761719e-02 2.04101562e-01
1.78710938e-01 2.55859375e-01 4.95605469e-02 2.03857422e-02
6.44531250e-02 -7.61718750e-02 -9.76562500e-02 2.46582031e-02
-8.88671875e-02 -1.98974609e-02 -1.02539062e-01 1.50299072e-03]

```

```

[ ]: # Check the index of the word 'topup' in word_index
print(f"The index of topup in word_index: {word_index['topup']}")

```

The index of topup in word_index: 19

```
[ ]: # Check the word 'topup' vectors in the embedding matrix
print(embedding_matrix[20])
```

```
[ 0.04980469  0.06640625  0.03833008  0.02355957 -0.02148438  0.20898438
 0.06396484 -0.02282715 -0.04101562 -0.26757812  0.1015625  0.10693359
-0.06591797  0.16992188 -0.0078125  -0.07861328 -0.06591797  0.12060547
-0.00390625 -0.02770996  0.19238281  0.13183594  0.16113281 -0.07324219
-0.22167969 -0.05102539 -0.12255859  0.06298828  0.01080322 -0.12695312
 0.04614258 -0.01794434 -0.03222656 -0.21484375 -0.01696777  0.0098877
-0.00976562 -0.05175781  0.12011719  0.04980469  0.01867676  0.05712891
-0.04492188 -0.16113281 -0.08105469 -0.09960938 -0.19824219 -0.00109863
 0.01239014  0.23144531 -0.06738281  0.08105469 -0.0177002  -0.12402344
-0.14746094 -0.10253906 -0.23046875 -0.03149414  0.03125  -0.09033203
 0.08251953 -0.09326172 -0.21679688  0.06103516  0.046875  -0.03466797
-0.05908203  0.12695312 -0.0025177  0.08251953  0.0703125  -0.01037598
 0.08447266  0.006073  -0.12988281 -0.06689453  0.18359375  0.21191406
 0.01495361 -0.04907227 -0.01525879  0.07080078 -0.04418945 -0.01153564
-0.05273438  0.06982422 -0.14453125 -0.06103516 -0.12402344  0.10058594
 0.17675781  0.12353516 -0.03710938 -0.38085938  0.15136719 -0.22753906
 0.03149414  0.09228516  0.01525879  0.07666016 -0.078125  -0.07080078
 0.05688477  0.10058594 -0.11425781 -0.02893066 -0.19921875 -0.08642578
 0.08349609  0.00552368 -0.19433594  0.05249023 -0.08251953 -0.12597656
 0.12792969  0.09814453 -0.12255859  0.10644531  0.0703125  0.13378906
 0.01550293  0.16894531 -0.10644531 -0.078125  -0.09716797 -0.02685547
-0.07470703 -0.19628906  0.23632812  0.07470703  0.04711914 -0.00427246
 0.0133667  -0.04882812 -0.02416992 -0.03588867  0.14550781 -0.02282715
 0.10546875  0.08349609 -0.00830078 -0.04931641 -0.09521484  0.11914062
 0.07470703 -0.04223633 -0.10742188 -0.22167969  0.05541992 -0.08398438
 0.00098419  0.09667969 -0.13867188  0.13476562  0.17871094  0.03149414
-0.06494141 -0.03613281 -0.27148438 -0.08691406  0.20800781  0.10546875
-0.06079102 -0.09863281  0.06884766 -0.19433594  0.03857422 -0.04589844
 0.1484375  -0.140625  0.00564575  0.03564453 -0.06542969 -0.13476562
 0.12060547 -0.01446533  0.2109375  0.21679688 -0.20507812  0.15332031
-0.28515625 -0.13769531  0.19140625 -0.05126953 -0.23242188 -0.08691406
-0.14257812 -0.06982422  0.15820312 -0.23535156 -0.06982422  0.01220703
 0.28515625 -0.09667969 -0.08837891  0.0625  0.15917969 -0.11767578
-0.18164062 -0.35546875 -0.05102539  0.12255859 -0.125  -0.01550293
-0.09521484  0.00643921 -0.13769531 -0.11865234 -0.02819824  0.02758789
-0.1171875  -0.0324707  -0.0098877  0.12255859  0.125  -0.01220703
 0.21972656  0.14160156 -0.09033203 -0.05566406 -0.03759766  0.04785156
-0.04980469 -0.15625  -0.09912109 -0.08691406 -0.10693359 -0.01293945
 0.00540161  0.04882812  0.01586914  0.10058594  0.03588867 -0.10351562
 0.03540039 -0.14453125 -0.13671875 -0.15820312  0.00994873  0.00726318
-0.10791016  0.04394531 -0.00314331  0.04077148 -0.03466797 -0.03271484
 0.140625  0.04492188  0.12988281 -0.09326172  0.13183594  0.04248047
 0.06396484  0.30078125  0.0559082  0.0039978  0.1328125  0.0390625
-0.1796875  -0.05053711 -0.16992188  0.04858398 -0.14941406  0.17675781]
```

```

0.09130859 0.08496094 0.06738281 0.03393555 -0.08300781 -0.06982422
0.06225586 0.12353516 0.05004883 0.06738281 0.05029297 -0.0324707
0.04003906 -0.20703125 -0.07861328 -0.00389099 -0.08398438 -0.11621094
0.05078125 0.24023438 -0.13183594 0.02600098 -0.23339844 -0.11474609
0.07568359 -0.12695312 -0.06787109 0.08496094 -0.00476074 -0.10595703
0.00970459 0.14648438 -0.04711914 0.20703125 0.01196289 0.03564453]

```

1.5.2 LSTM (baseline)

```

[ ]: # Define the output dimension for the embedding layer and hidden units
hidden_unit = 30
nlabel = 77

model = keras.models.Sequential()
e = layers.Embedding(voca_size+1, 300, weights=[embedding_matrix],
    ↪input_length=max_length_train_text, trainable=False)
model.add(e)
model.add(layers.LSTM(hidden_unit))
model.add(layers.Dense(nlabel, activation='softmax'))

# Compile the model
model.compile(optimizer='adam', loss='sparse_categorical_crossentropy',
    ↪metrics=['accuracy'])

# Summary the model
model.summary()

```

Model: "sequential_6"

Layer (type)	Output Shape	Param #
embedding_6 (Embedding)	(None, 29, 300)	626700
lstm_6 (LSTM)	(None, 30)	39720
dense_6 (Dense)	(None, 77)	2387

=====
 Total params: 668807 (2.55 MB)
 Trainable params: 42107 (164.48 KB)
 Non-trainable params: 626700 (2.39 MB)
 =====

```

[ ]: # Define the folder path to save the model
folder_path = '/content/drive/MyDrive/1. NLP CW/LSTM/'

# Define the file path for the model checkpoint

```

```
model_checkpoint_path = folder_path + 'LSTM_word2vec_model.keras'
```

```
# Define the model checkpoint
```

```
mc = tf.keras.callbacks.ModelCheckpoint(  
    filepath=model_checkpoint_path,  
    monitor='val_accuracy',  
    mode='max',  
    save_best_only=True)
```

```
[ ]: # Define early stopping
```

```
es = tf.keras.callbacks.EarlyStopping(monitor='val_loss', patience=3) # Random_  
↳ number of patience
```

```
[ ]: # Import time to measure the elapsed time
```

```
import time
```

```
# Measure time before training
```

```
start_time = time.time()
```

```
# Fit the model
```

```
history = model.fit(  
    X_train_padded, y_train,  
    epochs = 100,  
    validation_data = (X_val_padded, y_val),  
    callbacks = [mc, es],  
    batch_size = 32)
```

```
# End the training time
```

```
end_time = time.time()
```

```
# Measure the training time
```

```
training_time = end_time - start_time  
print("Training time:", training_time, "seconds")
```

Epoch 1/100

230/230 [=====] - 17s 60ms/step - loss: 4.1156 -
accuracy: 0.0242 - val_loss: 3.8803 - val_accuracy: 0.0397

Epoch 2/100

230/230 [=====] - 10s 43ms/step - loss: 3.8250 -
accuracy: 0.0403 - val_loss: 3.7300 - val_accuracy: 0.0332

Epoch 3/100

230/230 [=====] - 12s 54ms/step - loss: 3.6964 -
accuracy: 0.0448 - val_loss: 3.6242 - val_accuracy: 0.0577

Epoch 4/100

230/230 [=====] - 11s 47ms/step - loss: 3.6885 -
accuracy: 0.0483 - val_loss: 3.7484 - val_accuracy: 0.0457

Epoch 5/100

230/230 [=====] - 8s 36ms/step - loss: 3.5730 - accuracy: 0.0559 - val_loss: 3.4992 - val_accuracy: 0.0626
Epoch 6/100
230/230 [=====] - 10s 43ms/step - loss: 3.4550 - accuracy: 0.0595 - val_loss: 3.4535 - val_accuracy: 0.0735
Epoch 7/100
230/230 [=====] - 8s 37ms/step - loss: 3.3899 - accuracy: 0.0670 - val_loss: 3.3481 - val_accuracy: 0.0751
Epoch 8/100
230/230 [=====] - 8s 36ms/step - loss: 3.3292 - accuracy: 0.0788 - val_loss: 3.2886 - val_accuracy: 0.0947
Epoch 9/100
230/230 [=====] - 7s 30ms/step - loss: 3.2287 - accuracy: 0.0896 - val_loss: 3.1921 - val_accuracy: 0.0974
Epoch 10/100
230/230 [=====] - 7s 30ms/step - loss: 3.1885 - accuracy: 0.0961 - val_loss: 3.1957 - val_accuracy: 0.1051
Epoch 11/100
230/230 [=====] - 5s 20ms/step - loss: 3.1033 - accuracy: 0.1089 - val_loss: 3.0856 - val_accuracy: 0.1083
Epoch 12/100
230/230 [=====] - 12s 51ms/step - loss: 3.0548 - accuracy: 0.1075 - val_loss: 3.0339 - val_accuracy: 0.1165
Epoch 13/100
230/230 [=====] - 5s 21ms/step - loss: 3.0292 - accuracy: 0.1201 - val_loss: 3.0214 - val_accuracy: 0.1394
Epoch 14/100
230/230 [=====] - 5s 22ms/step - loss: 2.8912 - accuracy: 0.1465 - val_loss: 2.9558 - val_accuracy: 0.1486
Epoch 15/100
230/230 [=====] - 7s 31ms/step - loss: 2.8010 - accuracy: 0.1663 - val_loss: 2.8664 - val_accuracy: 0.1584
Epoch 16/100
230/230 [=====] - 5s 22ms/step - loss: 2.7962 - accuracy: 0.1687 - val_loss: 2.7211 - val_accuracy: 0.1987
Epoch 17/100
230/230 [=====] - 8s 33ms/step - loss: 2.6303 - accuracy: 0.1983 - val_loss: 2.6730 - val_accuracy: 0.2003
Epoch 18/100
230/230 [=====] - 5s 23ms/step - loss: 2.5544 - accuracy: 0.2111 - val_loss: 2.5998 - val_accuracy: 0.2123
Epoch 19/100
230/230 [=====] - 5s 23ms/step - loss: 2.4715 - accuracy: 0.2275 - val_loss: 2.5657 - val_accuracy: 0.2308
Epoch 20/100
230/230 [=====] - 7s 30ms/step - loss: 2.4096 - accuracy: 0.2365 - val_loss: 2.4372 - val_accuracy: 0.2580
Epoch 21/100

230/230 [=====] - 5s 23ms/step - loss: 2.3002 - accuracy: 0.2755 - val_loss: 2.3618 - val_accuracy: 0.2885
Epoch 22/100
230/230 [=====] - 6s 27ms/step - loss: 2.2456 - accuracy: 0.2908 - val_loss: 2.2972 - val_accuracy: 0.2961
Epoch 23/100
230/230 [=====] - 6s 26ms/step - loss: 2.1545 - accuracy: 0.3168 - val_loss: 2.2933 - val_accuracy: 0.3048
Epoch 24/100
230/230 [=====] - 5s 22ms/step - loss: 2.1261 - accuracy: 0.3207 - val_loss: 2.2585 - val_accuracy: 0.3190
Epoch 25/100
230/230 [=====] - 7s 31ms/step - loss: 2.0376 - accuracy: 0.3435 - val_loss: 2.2165 - val_accuracy: 0.3337
Epoch 26/100
230/230 [=====] - 5s 23ms/step - loss: 1.9811 - accuracy: 0.3601 - val_loss: 2.0442 - val_accuracy: 0.3620
Epoch 27/100
230/230 [=====] - 5s 23ms/step - loss: 1.8978 - accuracy: 0.3906 - val_loss: 2.1619 - val_accuracy: 0.3473
Epoch 28/100
230/230 [=====] - 7s 29ms/step - loss: 1.8436 - accuracy: 0.3972 - val_loss: 1.9467 - val_accuracy: 0.3930
Epoch 29/100
230/230 [=====] - 5s 22ms/step - loss: 1.8666 - accuracy: 0.3982 - val_loss: 1.9213 - val_accuracy: 0.4039
Epoch 30/100
230/230 [=====] - 6s 25ms/step - loss: 1.7474 - accuracy: 0.4336 - val_loss: 1.8573 - val_accuracy: 0.4295
Epoch 31/100
230/230 [=====] - 6s 26ms/step - loss: 1.6830 - accuracy: 0.4481 - val_loss: 1.8397 - val_accuracy: 0.4235
Epoch 32/100
230/230 [=====] - 5s 21ms/step - loss: 1.6706 - accuracy: 0.4536 - val_loss: 1.9097 - val_accuracy: 0.4208
Epoch 33/100
230/230 [=====] - 7s 30ms/step - loss: 1.6372 - accuracy: 0.4692 - val_loss: 1.7732 - val_accuracy: 0.4540
Epoch 34/100
230/230 [=====] - 5s 22ms/step - loss: 1.5439 - accuracy: 0.4954 - val_loss: 1.7658 - val_accuracy: 0.4311
Epoch 35/100
230/230 [=====] - 7s 30ms/step - loss: 1.5710 - accuracy: 0.4833 - val_loss: 1.8079 - val_accuracy: 0.4398
Epoch 36/100
230/230 [=====] - 7s 32ms/step - loss: 1.5137 - accuracy: 0.4986 - val_loss: 1.6672 - val_accuracy: 0.4752
Epoch 37/100

```

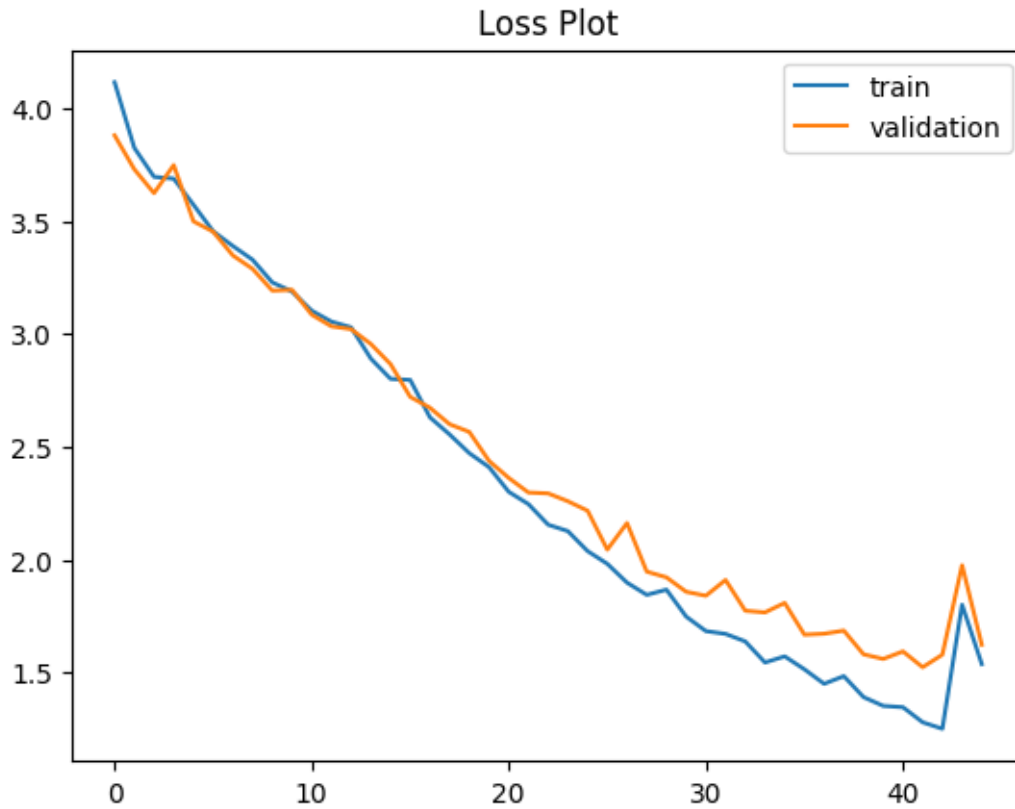
230/230 [=====] - 5s 21ms/step - loss: 1.4493 -
accuracy: 0.5159 - val_loss: 1.6713 - val_accuracy: 0.4736
Epoch 38/100
230/230 [=====] - 7s 29ms/step - loss: 1.4833 -
accuracy: 0.5117 - val_loss: 1.6851 - val_accuracy: 0.4665
Epoch 39/100
230/230 [=====] - 5s 23ms/step - loss: 1.3904 -
accuracy: 0.5399 - val_loss: 1.5800 - val_accuracy: 0.5112
Epoch 40/100
230/230 [=====] - 5s 21ms/step - loss: 1.3515 -
accuracy: 0.5535 - val_loss: 1.5596 - val_accuracy: 0.5133
Epoch 41/100
230/230 [=====] - 9s 38ms/step - loss: 1.3457 -
accuracy: 0.5576 - val_loss: 1.5929 - val_accuracy: 0.5112
Epoch 42/100
230/230 [=====] - 5s 21ms/step - loss: 1.2789 -
accuracy: 0.5862 - val_loss: 1.5229 - val_accuracy: 0.5215
Epoch 43/100
230/230 [=====] - 7s 31ms/step - loss: 1.2500 -
accuracy: 0.5828 - val_loss: 1.5785 - val_accuracy: 0.5106
Epoch 44/100
230/230 [=====] - 8s 33ms/step - loss: 1.7998 -
accuracy: 0.4834 - val_loss: 1.9750 - val_accuracy: 0.4491
Epoch 45/100
230/230 [=====] - 9s 41ms/step - loss: 1.5368 -
accuracy: 0.5402 - val_loss: 1.6225 - val_accuracy: 0.5253
Training time: 317.1865930557251 seconds

```


```

[ ]: # Plot the loss
plt.title('Loss Plot')
plt.plot(history.history['loss'], label='train')
plt.plot(history.history['val_loss'], label='validation')
plt.legend()
plt.show()

```



```
[ ]: # Plot the accuracy
plt.title('Accuracy Plot')
plt.plot(history.history['accuracy'], label='train')
plt.plot(history.history['val_accuracy'], label='validation')
plt.legend()
plt.show()
```



```
[ ]: # Load the saved model
saved_model = tf.keras.models.load_model(model_checkpoint_path)

# Evaluate the model with the test set
loss, accuracy = saved_model.evaluate(X_test_padded, y_test_array)

print("Test Loss:", loss)
print("Test Accuracy:", round((accuracy*100), 2))
```

```
97/97 [=====] - 4s 16ms/step - loss: 1.6655 - accuracy:
0.5049
Test Loss: 1.6655056476593018
Test Accuracy: 50.49
```

```
[ ]: # Check predictions with the test set
y_test_prob = saved_model.predict(X_test_padded)

# Convert probabilities to class labels
y_test_pred = np.argmax(y_test_prob, axis=1)

# Calculate precision, recall, and f1 score
```

```
precision = precision_score(y_test_array, y_test_pred, average='weighted')
recall = recall_score(y_test_array, y_test_pred, average='weighted')
f1 = f1_score(y_test_array, y_test_pred, average='weighted')

print("Precision:", round((precision*100), 2))
print("Recall:", round((recall*100), 2))
print("F1 Score:", round((f1*100), 2))
```

97/97 [=====] - 3s 16ms/step

Precision: 49.75

Recall: 50.49

F1 Score: 46.57

/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344:
 UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels
 with no predicted samples. Use `zero_division` parameter to control this
 behavior.

```
_warn_prf(average, modifier, msg_start, len(result))
```

[]: *# Error analysis*

```
# Print classification report
print(classification_report(y_test_array, y_test_pred))

# Check misclassified data
misclassified_data = np.where(y_test_pred != y_test_array)[0]
print(f"The number of misclassifications: {len(misclassified_data)}")

# Check the ratio of misclassifications
misclassification_ratio = (len(misclassified_data) / len(y_test_array)) * 100
# Round the number
rounded_ratio = round(misclassification_ratio, 2)
print(f"Proportion of misclassifications: {rounded_ratio}%")

# Iterate over misclassified data for error analysis
for idx in misclassified_data[:30]:
    input_text = X_test[idx]
    true_label = y_test[idx]
    predicted_label = y_test_pred[idx]

    # Print information about the misclassified data
    print("Input Text:", input_text)
    print("Actual Label:", true_label)
    print("Predicted Label:", predicted_label)
    print()
```

```
precision    recall  f1-score   support
```

0	0.60	0.15	0.24	40
1	0.29	0.57	0.39	40
2	0.97	0.95	0.96	40
3	0.00	0.00	0.00	40
4	0.83	0.85	0.84	40
5	0.37	0.62	0.46	40
6	0.65	0.78	0.70	40
7	0.50	0.72	0.59	40
8	0.55	0.82	0.66	40
9	0.86	0.62	0.72	40
10	0.44	0.30	0.36	40
11	0.43	0.82	0.56	40
12	0.65	0.28	0.39	40
13	0.36	0.70	0.47	40
14	0.46	0.78	0.57	40
15	0.58	0.85	0.69	40
16	0.49	0.50	0.49	40
17	0.88	0.90	0.89	40
18	0.71	0.75	0.73	40
19	0.82	0.90	0.86	40
20	0.50	0.65	0.57	40
21	0.00	0.00	0.00	40
22	0.44	0.75	0.56	40
23	0.00	0.00	0.00	40
24	0.92	0.82	0.87	40
25	0.65	0.78	0.70	40
26	0.36	0.25	0.29	40
27	0.25	0.03	0.05	40
28	0.78	0.62	0.69	40
29	0.24	0.80	0.37	40
30	0.00	0.00	0.00	40
31	0.87	0.85	0.86	40
32	0.88	0.90	0.89	40
33	0.64	0.85	0.73	40
34	0.79	0.78	0.78	40
35	0.47	0.62	0.54	40
36	0.60	0.65	0.63	40
37	0.00	0.00	0.00	40
38	0.00	0.00	0.00	40
39	0.46	0.68	0.55	40
40	0.20	0.03	0.04	40
41	0.43	0.45	0.44	40
42	0.26	0.17	0.21	40
43	0.60	0.23	0.33	40
44	0.00	0.00	0.00	40
45	0.58	0.70	0.64	40
46	0.67	0.40	0.50	40
47	0.45	0.82	0.58	40

48	0.76	0.40	0.52	40
49	0.50	0.03	0.05	40
50	0.67	0.50	0.57	40
51	0.54	0.93	0.69	40
52	0.21	0.07	0.11	40
53	0.63	0.55	0.59	40
54	0.46	0.40	0.43	40
55	0.50	0.03	0.05	40
56	0.60	0.07	0.13	40
57	0.92	0.85	0.88	40
58	0.47	0.40	0.43	40
59	0.49	0.42	0.45	40
60	0.72	0.85	0.78	40
61	0.38	0.15	0.21	40
62	0.67	0.15	0.24	40
63	0.57	0.62	0.60	40
64	0.55	0.75	0.63	40
65	0.81	0.55	0.66	40
66	0.47	0.53	0.49	40
67	0.74	0.42	0.54	40
68	0.00	0.00	0.00	40
69	0.00	0.00	0.00	40
70	0.73	1.00	0.84	40
71	0.00	0.00	0.00	40
72	0.00	0.00	0.00	40
73	1.00	0.82	0.90	40
74	0.10	0.97	0.18	40
75	0.44	0.82	0.57	40
76	0.92	0.88	0.90	40
accuracy				0.50 3080
macro avg				0.50 0.50 0.47 3080
weighted avg				0.50 0.50 0.47 3080

The number of misclassifications: 1525
Proportion of misclassifications: 49.51%
Input Text: locate card
Actual Label: 11
Predicted Label: 41

Input Text: get card
Actual Label: 11
Predicted Label: 43

Input Text: received card
Actual Label: 11
Predicted Label: 43

Input Text: tracking number card mailed
Actual Label: 11
Predicted Label: 22

Input Text: ordered card still havent received two weeks
Actual Label: 11
Predicted Label: 9

Input Text: im starting think card lost still hasnt arrived help
Actual Label: 11
Predicted Label: 41

Input Text: tracking info available
Actual Label: 11
Predicted Label: 42

Input Text: received new card dont see app anywhere
Actual Label: 13
Predicted Label: 12

Input Text: add card account
Actual Label: 13
Predicted Label: 10

Input Text: put old card back system found
Actual Label: 13
Predicted Label: 11

Input Text: hello found card misplaced need reactive
Actual Label: 13
Predicted Label: 41

Input Text: found card add app
Actual Label: 13
Predicted Label: 40

Input Text: link credit card
Actual Label: 13
Predicted Label: 39

Input Text: reactivate lost card found morning jacket pocket
Actual Label: 13
Predicted Label: 12

Input Text: app doesnt show card received
Actual Label: 13
Predicted Label: 11

Input Text: please show find location link card
Actual Label: 13
Predicted Label: 41

Input Text: way make old card usable app
Actual Label: 13
Predicted Label: 29

Input Text: need go app enter card info
Actual Label: 13
Predicted Label: 42

Input Text: found lost stolen card way link card account app
Actual Label: 13
Predicted Label: 41

Input Text: good time exchange
Actual Label: 32
Predicted Label: 33

Input Text: exchange rate like app
Actual Label: 32
Predicted Label: 17

Input Text: currencies exchange rate calculated
Actual Label: 32
Predicted Label: 31

Input Text: kind foreign exchange rate get exchange money
Actual Label: 32
Predicted Label: 31

Input Text: made currency exchange think charged
Actual Label: 17
Predicted Label: 31

Input Text: charged
Actual Label: 17
Predicted Label: 63

Input Text: conversion value card payments incorrect
Actual Label: 17
Predicted Label: 34

Input Text: wrong exchange rate used bought something foriegn currency
Actual Label: 17
Predicted Label: 76

Input Text: would like refund extra pound charged
Actual Label: 34
Predicted Label: 19

Input Text: explain random charge
Actual Label: 34
Predicted Label: 63

Input Text: remember purchasing anything & statement please tell
Actual Label: 34
Predicted Label: 45

```
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344:  
UndefinedMetricWarning: Precision and F-score are ill-defined and being set to  
0.0 in labels with no predicted samples. Use `zero_division` parameter to  
control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344:  
UndefinedMetricWarning: Precision and F-score are ill-defined and being set to  
0.0 in labels with no predicted samples. Use `zero_division` parameter to  
control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344:  
UndefinedMetricWarning: Precision and F-score are ill-defined and being set to  
0.0 in labels with no predicted samples. Use `zero_division` parameter to  
control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

1.5.3 LSTM (with dropout)

```
[ ]: # Define the output dimension for the embedding layer and hidden units
hidden_unit = 30
nlabel = 77

dropout_model = keras.models.Sequential()
e = layers.Embedding(voca_size+1, 300, weights=[embedding_matrix],
    ↪input_length=max_length_train_text, trainable=False)
dropout_model.add(e)
dropout_model.add(layers.LSTM(hidden_unit, dropout=0.2))
dropout_model.add(layers.Dense(nlabel, activation='softmax'))

# Compile the model
dropout_model.compile(optimizer='adam', loss='sparse_categorical_crossentropy',
    ↪metrics=['accuracy'])

# Summary the model
dropout_model.summary()
```

Model: "sequential_7"

Layer (type)	Output Shape	Param #
embedding_7 (Embedding)	(None, 29, 300)	626700
lstm_7 (LSTM)	(None, 30)	39720
dense_7 (Dense)	(None, 77)	2387

=====
 Total params: 668807 (2.55 MB)
 Trainable params: 42107 (164.48 KB)
 Non-trainable params: 626700 (2.39 MB)
 =====

```
[ ]: # Define the folder path to save the model
folder_path = '/content/drive/MyDrive/1. NLP CW/LSTM/'

# Define the file path for the model checkpoint
model_checkpoint_path = folder_path + 'dropout_LSTM_word2vec_model.keras'

# Define the model checkpoint
mc = tf.keras.callbacks.ModelCheckpoint(
    filepath=model_checkpoint_path,
    monitor='val_accuracy',
    mode='max',
    save_best_only=True)
```

```
[ ]: # Import time to measure the elapsed time
import time

# Measure time before training
start_time = time.time()

# Fit the model
history = dropout_model.fit(
    X_train_padded, y_train,
    epochs = 100,
    validation_data = (X_val_padded, y_val),
    callbacks = [mc, es],
    batch_size = 32)

# End the training time
end_time = time.time()

# Measure the training time
training_time = end_time - start_time
print("Training time:", training_time, "seconds")
```

```
Epoch 1/100
230/230 [=====] - 10s 32ms/step - loss: 4.1376 -
accuracy: 0.0193 - val_loss: 3.9000 - val_accuracy: 0.0321
Epoch 2/100
230/230 [=====] - 6s 25ms/step - loss: 3.8377 -
accuracy: 0.0366 - val_loss: 3.7534 - val_accuracy: 0.0359
Epoch 3/100
230/230 [=====] - 10s 45ms/step - loss: 3.7450 -
accuracy: 0.0395 - val_loss: 3.6682 - val_accuracy: 0.0555
Epoch 4/100
230/230 [=====] - 12s 50ms/step - loss: 3.6308 -
accuracy: 0.0502 - val_loss: 3.5310 - val_accuracy: 0.0670
Epoch 5/100
230/230 [=====] - 13s 55ms/step - loss: 3.4824 -
accuracy: 0.0682 - val_loss: 3.3377 - val_accuracy: 0.0773
Epoch 6/100
230/230 [=====] - 20s 87ms/step - loss: 3.3260 -
accuracy: 0.0791 - val_loss: 3.2037 - val_accuracy: 0.0936
Epoch 7/100
230/230 [=====] - 14s 63ms/step - loss: 3.2143 -
accuracy: 0.0913 - val_loss: 3.1157 - val_accuracy: 0.0931
Epoch 8/100
230/230 [=====] - 11s 47ms/step - loss: 3.1142 -
accuracy: 0.0976 - val_loss: 3.0257 - val_accuracy: 0.1181
Epoch 9/100
230/230 [=====] - 10s 45ms/step - loss: 3.0052 -
```

accuracy: 0.1194 - val_loss: 2.9100 - val_accuracy: 0.1399
 Epoch 10/100
 230/230 [=====] - 11s 49ms/step - loss: 2.9193 -
 accuracy: 0.1371 - val_loss: 2.8299 - val_accuracy: 0.1492
 Epoch 11/100
 230/230 [=====] - 11s 49ms/step - loss: 2.8110 -
 accuracy: 0.1563 - val_loss: 2.7365 - val_accuracy: 0.1932
 Epoch 12/100
 230/230 [=====] - 10s 42ms/step - loss: 2.6905 -
 accuracy: 0.2002 - val_loss: 2.6498 - val_accuracy: 0.2150
 Epoch 13/100
 230/230 [=====] - 13s 55ms/step - loss: 2.5966 -
 accuracy: 0.2156 - val_loss: 2.5415 - val_accuracy: 0.2270
 Epoch 14/100
 230/230 [=====] - 11s 49ms/step - loss: 2.4795 -
 accuracy: 0.2476 - val_loss: 2.4292 - val_accuracy: 0.2657
 Epoch 15/100
 230/230 [=====] - 9s 38ms/step - loss: 2.3678 -
 accuracy: 0.2746 - val_loss: 2.3284 - val_accuracy: 0.2831
 Epoch 16/100
 230/230 [=====] - 8s 36ms/step - loss: 2.3066 -
 accuracy: 0.2905 - val_loss: 2.2382 - val_accuracy: 0.3239
 Epoch 17/100
 230/230 [=====] - 6s 26ms/step - loss: 2.2535 -
 accuracy: 0.3101 - val_loss: 2.1902 - val_accuracy: 0.3408
 Epoch 18/100
 230/230 [=====] - 7s 29ms/step - loss: 2.1428 -
 accuracy: 0.3371 - val_loss: 2.1148 - val_accuracy: 0.3353
 Epoch 19/100
 230/230 [=====] - 7s 30ms/step - loss: 2.1001 -
 accuracy: 0.3518 - val_loss: 2.0857 - val_accuracy: 0.3511
 Epoch 20/100
 230/230 [=====] - 6s 25ms/step - loss: 2.0080 -
 accuracy: 0.3764 - val_loss: 1.9976 - val_accuracy: 0.3892
 Epoch 21/100
 230/230 [=====] - 8s 34ms/step - loss: 1.9501 -
 accuracy: 0.4006 - val_loss: 1.9471 - val_accuracy: 0.4192
 Epoch 22/100
 230/230 [=====] - 5s 24ms/step - loss: 1.8516 -
 accuracy: 0.4308 - val_loss: 1.8448 - val_accuracy: 0.4491
 Epoch 23/100
 230/230 [=====] - 8s 33ms/step - loss: 1.8403 -
 accuracy: 0.4406 - val_loss: 1.8490 - val_accuracy: 0.4273
 Epoch 24/100
 230/230 [=====] - 6s 25ms/step - loss: 1.7508 -
 accuracy: 0.4669 - val_loss: 1.7319 - val_accuracy: 0.4932
 Epoch 25/100
 230/230 [=====] - 6s 28ms/step - loss: 1.7503 -

accuracy: 0.4562 - val_loss: 1.6921 - val_accuracy: 0.5046
 Epoch 26/100
 230/230 [=====] - 7s 32ms/step - loss: 1.6505 -
 accuracy: 0.4906 - val_loss: 1.6588 - val_accuracy: 0.4986
 Epoch 27/100
 230/230 [=====] - 6s 24ms/step - loss: 1.5878 -
 accuracy: 0.5074 - val_loss: 1.6278 - val_accuracy: 0.5150
 Epoch 28/100
 230/230 [=====] - 8s 33ms/step - loss: 1.5702 -
 accuracy: 0.5087 - val_loss: 1.6215 - val_accuracy: 0.5253
 Epoch 29/100
 230/230 [=====] - 6s 24ms/step - loss: 1.5098 -
 accuracy: 0.5272 - val_loss: 1.5272 - val_accuracy: 0.5384
 Epoch 30/100
 230/230 [=====] - 8s 34ms/step - loss: 1.4530 -
 accuracy: 0.5452 - val_loss: 1.4939 - val_accuracy: 0.5449
 Epoch 31/100
 230/230 [=====] - 6s 26ms/step - loss: 1.4235 -
 accuracy: 0.5536 - val_loss: 1.4532 - val_accuracy: 0.5612
 Epoch 32/100
 230/230 [=====] - 9s 38ms/step - loss: 1.3776 -
 accuracy: 0.5689 - val_loss: 1.4207 - val_accuracy: 0.5683
 Epoch 33/100
 230/230 [=====] - 9s 38ms/step - loss: 1.3623 -
 accuracy: 0.5715 - val_loss: 1.3973 - val_accuracy: 0.5841
 Epoch 34/100
 230/230 [=====] - 5s 23ms/step - loss: 1.3205 -
 accuracy: 0.5886 - val_loss: 1.3861 - val_accuracy: 0.5955
 Epoch 35/100
 230/230 [=====] - 8s 35ms/step - loss: 1.3162 -
 accuracy: 0.5965 - val_loss: 1.4134 - val_accuracy: 0.5874
 Epoch 36/100
 230/230 [=====] - 6s 24ms/step - loss: 1.2758 -
 accuracy: 0.6103 - val_loss: 1.3703 - val_accuracy: 0.5846
 Epoch 37/100
 230/230 [=====] - 8s 34ms/step - loss: 1.2594 -
 accuracy: 0.6079 - val_loss: 1.3232 - val_accuracy: 0.6173
 Epoch 38/100
 230/230 [=====] - 6s 28ms/step - loss: 1.2538 -
 accuracy: 0.6172 - val_loss: 1.3042 - val_accuracy: 0.6189
 Epoch 39/100
 230/230 [=====] - 6s 26ms/step - loss: 1.2036 -
 accuracy: 0.6342 - val_loss: 1.2918 - val_accuracy: 0.6309
 Epoch 40/100
 230/230 [=====] - 7s 31ms/step - loss: 1.1799 -
 accuracy: 0.6395 - val_loss: 1.2506 - val_accuracy: 0.6462
 Epoch 41/100
 230/230 [=====] - 5s 24ms/step - loss: 1.1737 -

```


accuracy: 0.6425 - val_loss: 1.2694 - val_accuracy: 0.6287
Epoch 42/100
230/230 [=====] - 8s 35ms/step - loss: 1.1052 -
accuracy: 0.6644 - val_loss: 1.2296 - val_accuracy: 0.6440
Epoch 43/100
230/230 [=====] - 6s 26ms/step - loss: 1.1176 -
accuracy: 0.6617 - val_loss: 1.2405 - val_accuracy: 0.6429
Epoch 44/100
230/230 [=====] - 8s 33ms/step - loss: 1.1162 -
accuracy: 0.6654 - val_loss: 1.1781 - val_accuracy: 0.6663
Epoch 45/100
230/230 [=====] - 6s 25ms/step - loss: 1.0243 -
accuracy: 0.6992 - val_loss: 1.1823 - val_accuracy: 0.6685
Epoch 46/100
230/230 [=====] - 6s 26ms/step - loss: 1.0189 -
accuracy: 0.7020 - val_loss: 1.2504 - val_accuracy: 0.6456
Epoch 47/100
230/230 [=====] - 7s 32ms/step - loss: 0.9995 -
accuracy: 0.7035 - val_loss: 1.1060 - val_accuracy: 0.6935
Epoch 48/100
230/230 [=====] - 6s 25ms/step - loss: 1.0028 -
accuracy: 0.7035 - val_loss: 1.1554 - val_accuracy: 0.6761
Epoch 49/100
230/230 [=====] - 8s 34ms/step - loss: 0.9643 -
accuracy: 0.7107 - val_loss: 1.0804 - val_accuracy: 0.7011
Epoch 50/100
230/230 [=====] - 6s 25ms/step - loss: 0.9423 -
accuracy: 0.7140 - val_loss: 1.1377 - val_accuracy: 0.6837
Epoch 51/100
230/230 [=====] - 7s 29ms/step - loss: 0.9571 -
accuracy: 0.7085 - val_loss: 1.0972 - val_accuracy: 0.6924
Epoch 52/100
230/230 [=====] - 7s 30ms/step - loss: 0.8994 -
accuracy: 0.7332 - val_loss: 1.1023 - val_accuracy: 0.6854
Training time: 421.2397985458374 seconds

```


```

[ ]: # Plot the loss
plt.title('Loss Plot')
plt.plot(history.history['loss'], label='train')
plt.plot(history.history['val_loss'], label='validation')
plt.legend()
plt.show()

```

```
[ ]: # Plot the accuracy
plt.title('Accuracy Plot')
plt.plot(history.history['accuracy'], label='train')
plt.plot(history.history['val_accuracy'], label='validation')
plt.legend()
plt.show()
```



```
[ ]: # Load the saved model
saved_model = tf.keras.models.load_model(model_checkpoint_path)

# Evaluate the model with the test set
loss, accuracy = saved_model.evaluate(X_test_padded, y_test_array)

print("Test Loss:", loss)
print("Test Accuracy:", round((accuracy*100), 2))
```

```
97/97 [=====] - 2s 9ms/step - loss: 1.1089 - accuracy:
0.7000
Test Loss: 1.108872652053833
Test Accuracy: 70.0
```

```
[ ]: # Check predictions with the test set
y_test_prob = saved_model.predict(X_test_padded)

# Convert probabilities to class labels
y_test_pred = np.argmax(y_test_prob, axis=1)

# Calculate precision, recall, and f1 score
```

```

precision = precision_score(y_test_array, y_test_pred, average='weighted')
recall = recall_score(y_test_array, y_test_pred, average='weighted')
f1 = f1_score(y_test_array, y_test_pred, average='weighted')

print("Precision:", round((precision*100), 2))
print("Recall:", round((recall*100), 2))
print("F1 Score:", round((f1*100), 2))

```

97/97 [=====] - 1s 9ms/step

Precision: 70.48

Recall: 70.0

F1 Score: 68.02

/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344:
 UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels
 with no predicted samples. Use `zero_division` parameter to control this
 behavior.

```
_warn_prf(average, modifier, msg_start, len(result))
```

```

[ ]: # Error analysis

# Print classification report
print(classification_report(y_test_array, y_test_pred))

# Check misclassified data
misclassified_data = np.where(y_test_pred != y_test_array)[0]
print(f"The number of misclassifications: {len(misclassified_data)}")

# Check the ratio of misclassifications
misclassification_ratio = (len(misclassified_data) / len(y_test_array)) * 100
# Round the number
rounded_ratio = round(misclassification_ratio, 2)
print(f"Proportion of misclassifications: {rounded_ratio}%")

# Iterate over misclassified data for error analysis
for idx in misclassified_data[:30]:
    input_text = X_test[idx]
    true_label = y_test[idx]
    predicted_label = y_test_pred[idx]

    # Print information about the misclassified data
    print("Input Text:", input_text)
    print("Actual Label:", true_label)
    print("Predicted Label:", predicted_label)
    print()

```

```
precision    recall  f1-score   support
```

0	0.92	0.90	0.91	40
1	0.90	0.95	0.93	40
2	1.00	0.97	0.99	40
3	0.72	0.33	0.45	40
4	0.96	0.68	0.79	40
5	0.47	0.68	0.55	40
6	0.81	0.85	0.83	40
7	0.50	0.35	0.41	40
8	0.80	0.80	0.80	40
9	0.95	0.93	0.94	40
10	0.78	0.53	0.63	40
11	0.52	0.65	0.58	40
12	0.90	0.47	0.62	40
13	0.62	0.88	0.73	40
14	0.43	0.70	0.53	40
15	0.71	0.88	0.79	40
16	0.56	0.70	0.62	40
17	0.84	0.80	0.82	40
18	0.67	0.72	0.70	40
19	0.80	0.90	0.85	40
20	0.72	0.72	0.73	40
21	0.72	0.78	0.75	40
22	0.64	0.80	0.71	40
23	0.00	0.00	0.00	40
24	0.81	0.95	0.87	40
25	0.58	0.78	0.67	40
26	0.76	0.88	0.81	40
27	0.75	0.75	0.75	40
28	0.73	0.68	0.70	40
29	0.41	0.85	0.55	40
30	0.79	0.85	0.82	40
31	0.57	0.80	0.67	40
32	0.93	0.95	0.94	40
33	0.77	0.75	0.76	40
34	0.88	0.70	0.78	40
35	0.56	0.72	0.63	40
36	0.85	0.85	0.85	40
37	0.00	0.00	0.00	40
38	0.61	0.78	0.68	40
39	0.47	0.60	0.53	40
40	0.64	0.95	0.77	40
41	0.49	0.70	0.58	40
42	0.71	0.93	0.80	40
43	0.40	0.53	0.45	40
44	0.94	0.82	0.88	40
45	0.77	0.85	0.81	40
46	0.76	0.85	0.80	40
47	0.71	0.60	0.65	40

48	0.68	0.65	0.67	40
49	0.82	0.45	0.58	40
50	0.82	0.57	0.68	40
51	0.70	0.82	0.76	40
52	0.74	0.57	0.65	40
53	0.68	0.53	0.59	40
54	0.73	0.82	0.78	40
55	0.95	0.97	0.96	40
56	0.82	0.82	0.82	40
57	0.97	0.85	0.91	40
58	0.69	0.60	0.64	40
59	0.52	0.78	0.62	40
60	0.76	0.88	0.81	40
61	0.82	0.57	0.68	40
62	0.63	0.55	0.59	40
63	0.69	0.78	0.73	40
64	0.91	0.78	0.84	40
65	0.61	0.70	0.65	40
66	0.83	0.25	0.38	40
67	0.75	0.68	0.71	40
68	0.50	0.03	0.05	40
69	0.00	0.00	0.00	40
70	0.83	1.00	0.91	40
71	0.97	0.80	0.88	40
72	1.00	0.05	0.10	40
73	0.95	0.90	0.92	40
74	0.35	0.97	0.51	40
75	0.84	0.80	0.82	40
76	0.84	0.68	0.75	40
accuracy				0.70 3080
macro avg				0.70 0.70 0.68 3080
weighted avg				0.70 0.70 0.68 3080

The number of misclassifications: 924
Proportion of misclassifications: 30.0%
Input Text: locate card
Actual Label: 11
Predicted Label: 43

Input Text: still received new card ordered week ago
Actual Label: 11
Predicted Label: 13

Input Text: get card
Actual Label: 11
Predicted Label: 39

Input Text: know tracking number new card sent
Actual Label: 11
Predicted Label: 13

Input Text: received card
Actual Label: 11
Predicted Label: 43

Input Text: still waiting card
Actual Label: 11
Predicted Label: 43

Input Text: track card
Actual Label: 11
Predicted Label: 43

Input Text: still dont card weeks
Actual Label: 11
Predicted Label: 43

Input Text: ive waiting longer expected card
Actual Label: 11
Predicted Label: 43

Input Text: hasnt card delivered
Actual Label: 11
Predicted Label: 43

Input Text: card still hasnt arrived weeks lost
Actual Label: 11
Predicted Label: 13

Input Text: get card yet lost
Actual Label: 11
Predicted Label: 13

Input Text: status card ordered
Actual Label: 11
Predicted Label: 43

Input Text: im starting think card lost still hasnt arrived help
Actual Label: 11
Predicted Label: 13

Input Text: would like reactivate card
Actual Label: 13
Predicted Label: 0

Input Text: add card account
Actual Label: 13
Predicted Label: 24

Input Text: link credit card
Actual Label: 13
Predicted Label: 11

Input Text: link replacement card
Actual Label: 13
Predicted Label: 11

Input Text: link another card account
Actual Label: 13
Predicted Label: 43

Input Text: good time exchange
Actual Label: 32
Predicted Label: 31

Input Text: kind foreign exchange rate get exchange money
Actual Label: 32
Predicted Label: 50

Input Text: made currency exchange think charged
Actual Label: 17
Predicted Label: 31

Input Text: rate exchange card payment incorrect
Actual Label: 17
Predicted Label: 76

Input Text: bought something overseas wrong exchange rate statement
Actual Label: 17
Predicted Label: 76

Input Text: exchange rate card payment wrong
Actual Label: 17
Predicted Label: 76

Input Text: charged
Actual Label: 17
Predicted Label: 15

Input Text: believe card payment exchange rate incorrect
Actual Label: 17
Predicted Label: 76

Input Text: conversion value card payments incorrect
Actual Label: 17
Predicted Label: 25

Input Text: exchange rate totally wrong card payment
Actual Label: 17
Predicted Label: 76

Input Text: extra dollar charged account
Actual Label: 34
Predicted Label: 19

```
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344:  
UndefinedMetricWarning: Precision and F-score are ill-defined and being set to  
0.0 in labels with no predicted samples. Use `zero_division` parameter to  
control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344:  
UndefinedMetricWarning: Precision and F-score are ill-defined and being set to  
0.0 in labels with no predicted samples. Use `zero_division` parameter to  
control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344:  
UndefinedMetricWarning: Precision and F-score are ill-defined and being set to  
0.0 in labels with no predicted samples. Use `zero_division` parameter to  
control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

1.5.4 Hyperparameter tuning

```
[ ]: # The code for hyperparameter tuning is derived from the Tensorflow website.  
# (https://www.tensorflow.org/tutorials/keras/keras_tuner)  
  
# Define the model for hyperparameter tuning  
def model_builder(hp):  
    model = keras.models.Sequential()  
    e = layers.Embedding(voca_size+1, 300, weights=[embedding_matrix],  
↳input_length=max_length_train_text, trainable=False)  
    model.add(e)  
    hp_units = hp.Int('units', min_value = 20, max_value = 50, step = 10) # Set  
↳up the hyperparameters  
    model.add(layers.LSTM(units = hp_units)) # We will check the optimal hidden  
↳unit for the LSTM layer  
    model.add(layers.Dense(nlabel, activation='softmax'))  
  
    hp_learning_rate = hp.Choice('learning_rate', values = [0.01, 0.001, 0.0001])  
↳# Set up the hyperparameters
```



```

model.compile(optimizer = keras.optimizers.Adam(learning_rate = hp_learning_rate), # We will check the optimal learning rate
              loss = 'sparse_categorical_crossentropy',
              metrics = ['accuracy'])
return model

```

```

[ ]: # The code for hyperparameter tuning is derived from the Tensorflow website.
      # (https://www.tensorflow.org/tutorials/keras/keras_tuner)

```

```

# Specify the tuner
tuner = kt.Hyperband(model_builder,
                    objective = 'val_accuracy',
                    max_epochs = 100)

```

```

[ ]: # The code for hyperparameter tuning is derived from the Tensorflow website.
      # (https://www.tensorflow.org/tutorials/keras/keras_tuner)

```

```

# Set up a callback for early stopping
stop_early = tf.keras.callbacks.EarlyStopping(monitor='val_loss', patience=5)

```

```

[ ]: # The code for hyperparameter tuning is derived from the Tensorflow website.
      # (https://www.tensorflow.org/tutorials/keras/keras_tuner)

```

```

# Run the tuner
tuner.search(X_train_padded, y_train, epochs = 100, validation_data = (X_val_padded, y_val), callbacks = [stop_early])

# Get the optimal hyperparameters
best_hps = tuner.get_best_hyperparameters(num_trials = 1)[0]

print(f"The optimal number of units: {best_hps.get('units')}. The optimal learning rate: {best_hps.get('learning_rate')}")

```

Trial 12 Complete [00h 00m 13s]
val_accuracy: 0.03647251054644585

Best val_accuracy So Far: 0.17038649320602417
Total elapsed time: 00h 24m 36s
The optimal number of units: 50. The optimal learning rate: 0.01.

1.5.5 Tuned LSTM

```

[ ]: # Define the output dimension for the embedding layer and hidden units
nlabel = 77

tuned_model = keras.models.Sequential()

```

```

e = layers.Embedding(voca_size+1, 300, weights=[embedding_matrix],
    ↪input_length=max_length_train_text, trainable=False)
tuned_model.add(e)
tuned_model.add(layers.LSTM(50))
tuned_model.add(layers.Dense(nlabel, activation='softmax'))

# Compile the model
tuned_model.compile(optimizer=keras.optimizers.Adam(learning_rate = 0.01),
    ↪loss='sparse_categorical_crossentropy', metrics=['accuracy'])

# Summary the model
tuned_model.summary()

```

Model: "sequential_8"

Layer (type)	Output Shape	Param #
embedding_8 (Embedding)	(None, 29, 300)	626700
lstm_8 (LSTM)	(None, 50)	70200
dense_8 (Dense)	(None, 77)	3927

=====
 Total params: 700827 (2.67 MB)
 Trainable params: 74127 (289.56 KB)
 Non-trainable params: 626700 (2.39 MB)
 =====

```

[ ]: # Define the folder path to save the model
folder_path = '/content/drive/MyDrive/1. NLP CW/LSTM/'

# Define the file path for the model checkpoint
model_checkpoint_path = folder_path + 'tuned_LSTM_word2vec_model.keras'

# Define the model checkpoint
mc = tf.keras.callbacks.ModelCheckpoint(
    filepath=model_checkpoint_path,
    monitor='val_accuracy',
    mode='max',
    save_best_only=True)

```

```

[ ]: # Import time to measure the elapsed time
import time

# Measure time before training
start_time = time.time()

```

```

# Fit the model
history = tuned_model.fit(
    X_train_padded, y_train,
    epochs = 100,
    validation_data = (X_val_padded, y_val),
    callbacks = [mc, es],
    batch_size = 32)

# End the training time
end_time = time.time()

# Measure the training time
training_time = end_time - start_time
print("Training time:", training_time, "seconds")

```


```

Epoch 1/100
230/230 [=====] - 14s 49ms/step - loss: 4.1499 -
accuracy: 0.0271 - val_loss: 4.1483 - val_accuracy: 0.0299
Epoch 2/100
230/230 [=====] - 11s 48ms/step - loss: 4.2695 -
accuracy: 0.0267 - val_loss: 4.2910 - val_accuracy: 0.0240
Epoch 3/100
230/230 [=====] - 7s 29ms/step - loss: 4.1253 -
accuracy: 0.0253 - val_loss: 4.0790 - val_accuracy: 0.0289
Epoch 4/100
230/230 [=====] - 10s 45ms/step - loss: 3.7872 -
accuracy: 0.0501 - val_loss: 3.4356 - val_accuracy: 0.0806
Epoch 5/100
230/230 [=====] - 8s 33ms/step - loss: 3.2299 -
accuracy: 0.1010 - val_loss: 2.9360 - val_accuracy: 0.1595
Epoch 6/100
230/230 [=====] - 10s 44ms/step - loss: 2.6225 -
accuracy: 0.2332 - val_loss: 2.2807 - val_accuracy: 0.3005
Epoch 7/100
230/230 [=====] - 9s 38ms/step - loss: 1.9497 -
accuracy: 0.3935 - val_loss: 1.7657 - val_accuracy: 0.4605
Epoch 8/100
230/230 [=====] - 12s 52ms/step - loss: 1.4739 -
accuracy: 0.5501 - val_loss: 1.4363 - val_accuracy: 0.5781
Epoch 9/100
230/230 [=====] - 11s 47ms/step - loss: 1.1382 -
accuracy: 0.6572 - val_loss: 1.2136 - val_accuracy: 0.6587
Epoch 10/100
230/230 [=====] - 8s 35ms/step - loss: 0.8966 -
accuracy: 0.7374 - val_loss: 1.0990 - val_accuracy: 0.6973
Epoch 11/100

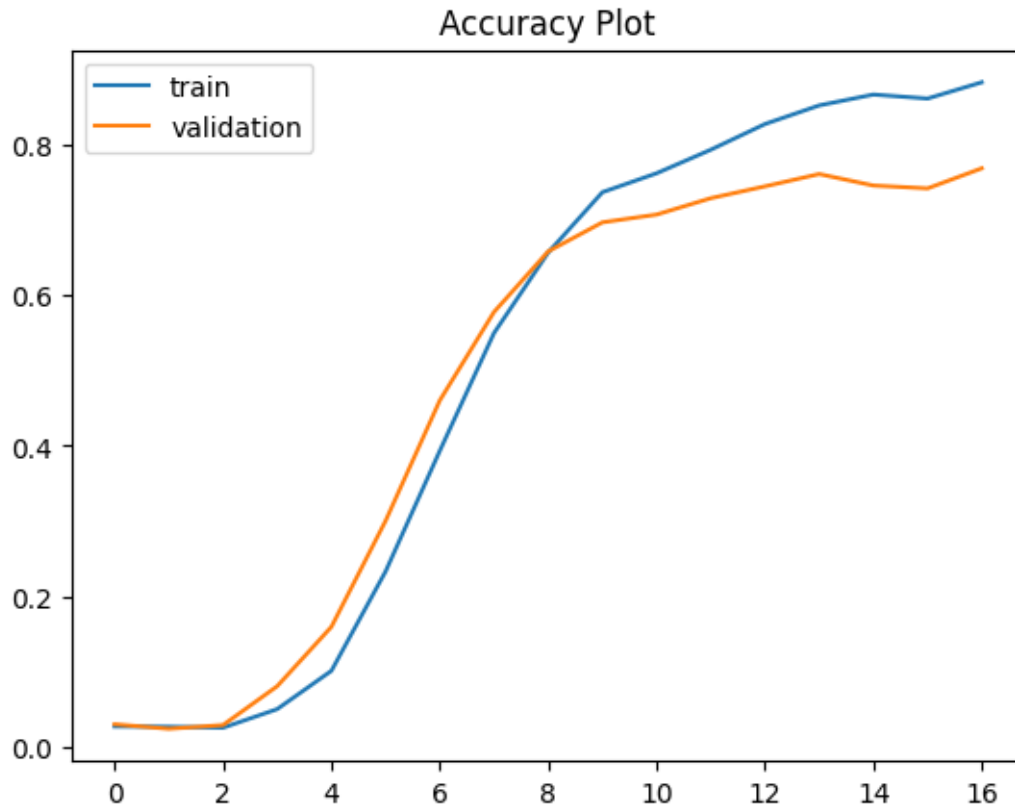
```

```
230/230 [=====] - 10s 42ms/step - loss: 0.8136 -  
accuracy: 0.7626 - val_loss: 1.0372 - val_accuracy: 0.7077  
Epoch 12/100  
230/230 [=====] - 9s 40ms/step - loss: 0.6888 -  
accuracy: 0.7939 - val_loss: 0.9918 - val_accuracy: 0.7295  
Epoch 13/100  
230/230 [=====] - 8s 36ms/step - loss: 0.5707 -  
accuracy: 0.8281 - val_loss: 0.9511 - val_accuracy: 0.7452  
Epoch 14/100  
230/230 [=====] - 10s 44ms/step - loss: 0.4869 -  
accuracy: 0.8528 - val_loss: 0.9241 - val_accuracy: 0.7616  
Epoch 15/100  
230/230 [=====] - 7s 30ms/step - loss: 0.4483 -  
accuracy: 0.8673 - val_loss: 0.9850 - val_accuracy: 0.7463  
Epoch 16/100  
230/230 [=====] - 11s 47ms/step - loss: 0.4489 -  
accuracy: 0.8620 - val_loss: 1.0156 - val_accuracy: 0.7425  
Epoch 17/100  
230/230 [=====] - 7s 32ms/step - loss: 0.4037 -  
accuracy: 0.8836 - val_loss: 0.9544 - val_accuracy: 0.7692  
Training time: 161.79024744033813 seconds
```

```
[ ]: # Plot the loss  
plt.title('Loss Plot')  
plt.plot(history.history['loss'], label='train')  
plt.plot(history.history['val_loss'], label='validation')  
plt.legend()  
plt.show()
```



```
[ ]: # Plot the accuracy
plt.title('Accuracy Plot')
plt.plot(history.history['accuracy'], label='train')
plt.plot(history.history['val_accuracy'], label='validation')
plt.legend()
plt.show()
```



```
[ ]: # Load the saved model
saved_model = tf.keras.models.load_model(model_checkpoint_path)

# Evaluate the model with the test set
loss, accuracy = saved_model.evaluate(X_test_padded, y_test_array)

print("Test Loss:", loss)
print("Test Accuracy:", round((accuracy*100), 2))
```

```
97/97 [=====] - 3s 22ms/step - loss: 0.8925 - accuracy:
0.7731
Test Loss: 0.892530083656311
Test Accuracy: 77.31
```

```
[ ]: # Check predictions with the test set
y_test_prob = saved_model.predict(X_test_padded)

# Convert probabilities to class labels
y_test_pred = np.argmax(y_test_prob, axis=1)

# Calculate precision, recall, and f1 score
```

```

precision = precision_score(y_test_array, y_test_pred, average='weighted')
recall = recall_score(y_test_array, y_test_pred, average='weighted')
f1 = f1_score(y_test_array, y_test_pred, average='weighted')

print("Precision:", round((precision*100), 2))
print("Recall:", round((recall*100), 2))
print("F1 Score:", round((f1*100), 2))

```

```

97/97 [=====] - 2s 12ms/step
Precision: 78.87
Recall: 77.31
F1 Score: 77.25

```

```

[ ]: # Error analysis

# Print classification report
print(classification_report(y_test_array, y_test_pred))

# Check misclassified data
misclassified_data = np.where(y_test_pred != y_test_array)[0]
print(f"The number of misclassifications: {len(misclassified_data)}")

# Check the ratio of misclassifications
misclassification_ratio = (len(misclassified_data) / len(y_test_array)) * 100
# Round the number
rounded_ratio = round(misclassification_ratio, 2)
print(f"Proportion of misclassifications: {rounded_ratio}%")

# Iterate over misclassified data for error analysis
for idx in misclassified_data[:30]:
    input_text = X_test[idx]
    true_label = y_test[idx]
    predicted_label = y_test_pred[idx]

    # Print information about the misclassified data
    print("Input Text:", input_text)
    print("Actual Label:", true_label)
    print("Predicted Label:", predicted_label)
    print()

```

	precision	recall	f1-score	support
0	0.93	0.95	0.94	40
1	0.85	0.97	0.91	40
2	0.97	0.97	0.97	40
3	0.59	0.68	0.63	40
4	0.85	0.70	0.77	40
5	0.66	0.68	0.67	40

6	0.74	0.78	0.76	40
7	0.71	0.72	0.72	40
8	0.86	0.80	0.83	40
9	0.87	1.00	0.93	40
10	0.69	0.50	0.58	40
11	0.67	0.78	0.72	40
12	0.68	0.75	0.71	40
13	0.92	0.90	0.91	40
14	0.55	0.82	0.66	40
15	0.74	0.85	0.79	40
16	0.56	0.55	0.56	40
17	0.78	0.90	0.84	40
18	0.91	0.72	0.81	40
19	0.77	0.93	0.84	40
20	0.50	0.72	0.59	40
21	0.97	0.78	0.86	40
22	0.67	0.65	0.66	40
23	0.97	0.85	0.91	40
24	0.85	0.97	0.91	40
25	0.75	0.82	0.79	40
26	0.73	0.93	0.81	40
27	0.88	0.72	0.79	40
28	0.82	0.70	0.76	40
29	0.86	0.78	0.82	40
30	0.88	0.93	0.90	40
31	0.92	0.82	0.87	40
32	0.85	0.88	0.86	40
33	0.77	0.85	0.81	40
34	0.79	0.78	0.78	40
35	0.68	0.68	0.68	40
36	0.79	0.78	0.78	40
37	0.54	0.80	0.65	40
38	0.87	0.97	0.92	40
39	0.93	0.70	0.80	40
40	0.81	0.75	0.78	40
41	0.73	0.55	0.63	40
42	0.97	0.88	0.92	40
43	0.73	0.68	0.70	40
44	0.98	1.00	0.99	40
45	0.78	0.72	0.75	40
46	0.76	0.78	0.77	40
47	0.76	0.78	0.77	40
48	0.56	0.72	0.63	40
49	0.78	0.72	0.75	40
50	1.00	0.62	0.77	40
51	0.84	0.90	0.87	40
52	0.81	0.88	0.84	40
53	0.64	0.57	0.61	40

54	0.84	0.78	0.81	40
55	0.97	0.88	0.92	40
56	0.88	0.70	0.78	40
57	1.00	0.90	0.95	40
58	0.63	0.68	0.65	40
59	0.83	0.60	0.70	40
60	0.85	0.85	0.85	40
61	0.60	0.68	0.64	40
62	0.76	0.55	0.64	40
63	0.75	0.82	0.79	40
64	0.75	0.82	0.79	40
65	0.88	0.55	0.68	40
66	0.58	0.72	0.64	40
67	0.85	0.57	0.69	40
68	0.81	0.75	0.78	40
69	0.62	0.25	0.36	40
70	0.84	0.95	0.89	40
71	0.93	0.97	0.95	40
72	0.86	0.47	0.61	40
73	0.97	0.88	0.92	40
74	0.44	0.88	0.58	40
75	0.74	0.88	0.80	40
76	0.86	0.80	0.83	40
accuracy				0.77
macro avg				0.79
weighted avg				0.77

The number of misclassifications: 699
Proportion of misclassifications: 22.69%
Input Text: ordered card arrived help please
Actual Label: 11
Predicted Label: 12

Input Text: get card
Actual Label: 11
Predicted Label: 43

Input Text: waiting longer expected bank card could provide information arrive
Actual Label: 11
Predicted Label: 27

Input Text: hasnt card delivered
Actual Label: 11
Predicted Label: 12

Input Text: status card ordered
Actual Label: 11

Predicted Label: 14

Input Text: expecting new card wondering havent received yet

Actual Label: 11

Predicted Label: 40

Input Text: know card arrive

Actual Label: 11

Predicted Label: 12

Input Text: im starting think card lost still hasnt arrived help

Actual Label: 11

Predicted Label: 41

Input Text: tracking info available

Actual Label: 11

Predicted Label: 13

Input Text: add card account

Actual Label: 13

Predicted Label: 18

Input Text: put old card back system found

Actual Label: 13

Predicted Label: 41

Input Text: hello found card misplaced need reactive

Actual Label: 13

Predicted Label: 49

Input Text: already one cards link

Actual Label: 13

Predicted Label: 40

Input Text: good time exchange

Actual Label: 32

Predicted Label: 33

Input Text: currencies exchange rate calculated

Actual Label: 32

Predicted Label: 31

Input Text: im trying figure current exchange rate

Actual Label: 32

Predicted Label: 76

Input Text: kind foreign exchange rate get exchange money

Actual Label: 32

Predicted Label: 31

Input Text: rate get determined

Actual Label: 32

Predicted Label: 17

Input Text: made currency exchange think charged

Actual Label: 17

Predicted Label: 31

Input Text: rate low sure using right exchange rate

Actual Label: 17

Predicted Label: 32

Input Text: charged

Actual Label: 17

Predicted Label: 15

Input Text: conversion value card payments incorrect

Actual Label: 17

Predicted Label: 63

Input Text: explain random charge

Actual Label: 34

Predicted Label: 63

Input Text: transaction credited

Actual Label: 34

Predicted Label: 8

Input Text: fee come

Actual Label: 34

Predicted Label: 19

Input Text: extra charge

Actual Label: 34

Predicted Label: 15

Input Text: extra pound charge card

Actual Label: 34

Predicted Label: 15

Input Text: euro fee come

Actual Label: 34

Predicted Label: 17

Input Text: euro fee statement

Actual Label: 34

Predicted Label: 15

Input Text: reason accounts charged extra dollar

Actual Label: 34

Predicted Label: 15

1.5.6 Tuned LSTM (with dropout)

```
[ ]: # Define the output dimension for the embedding layer and hidden units
nlabel = 77

dropout_tuned_model = keras.models.Sequential()
e = layers.Embedding(voca_size+1, 300, weights=[embedding_matrix],
    ↪input_length=max_length_train_text, trainable=False)
dropout_tuned_model.add(e)
dropout_tuned_model.add(layers.LSTM(50, dropout=0.2))
dropout_tuned_model.add(layers.Dense(nlabel, activation='softmax'))

# Compile the model
dropout_tuned_model.compile(optimizer=keras.optimizers.Adam(learning_rate = 0.
    ↪01), loss='sparse_categorical_crossentropy', metrics=['accuracy'])

# Summary the model
dropout_tuned_model.summary()
```

Model: "sequential_9"

Layer (type)	Output Shape	Param #
embedding_9 (Embedding)	(None, 29, 300)	626700
lstm_9 (LSTM)	(None, 50)	70200
dense_9 (Dense)	(None, 77)	3927

=====
Total params: 700827 (2.67 MB)
Trainable params: 74127 (289.56 KB)
Non-trainable params: 626700 (2.39 MB)
=====

```
[ ]: # Define the folder path to save the model
folder_path = '/content/drive/MyDrive/1. NLP CW/LSTM/'

# Define the file path for the model checkpoint
model_checkpoint_path = folder_path + 'dropout_tuned_LSTM_word2vec_model.keras'
```

```
# Define the model checkpoint
mc = tf.keras.callbacks.ModelCheckpoint(
    filepath=model_checkpoint_path,
    monitor='val_accuracy',
    mode='max',
    save_best_only=True)
```

```
[ ]: # Get the weights of the embedding layer
embedding_weights = dropout_tuned_model.layers[0].get_weights()[0]

# Define the file path to save the weights
embedding_weights_file = '/content/drive/MyDrive/1. NLP CW/embedding_weights.
    ↪ npy'

# Save the weights as a file
np.save(embedding_weights_file, embedding_weights)
```

```
[ ]: # Import time to measure the elapsed time
import time

# Measure time before training
start_time = time.time()

# Fit the model
history = dropout_tuned_model.fit(
    X_train_padded, y_train,
    epochs = 100,
    validation_data = (X_val_padded, y_val),
    callbacks = [mc, es],
    batch_size = 32)

# End the training time
end_time = time.time()

# Measure the training time
training_time = end_time - start_time
print("Training time:", training_time, "seconds")
```

Epoch 1/100

230/230 [=====] - 16s 53ms/step - loss: 3.9733 - accuracy: 0.0348 - val_loss: 3.7733 - val_accuracy: 0.0474

Epoch 2/100

230/230 [=====] - 10s 43ms/step - loss: 3.8154 - accuracy: 0.0566 - val_loss: 3.8087 - val_accuracy: 0.0637


Epoch 3/100

230/230 [=====] - 10s 41ms/step - loss: 3.7495 -


accuracy: 0.0709 - val_loss: 3.7288 - val_accuracy: 0.0680
 Epoch 4/100
 230/230 [=====] - 11s 48ms/step - loss: 3.5694 -
 accuracy: 0.0957 - val_loss: 3.2373 - val_accuracy: 0.1214
 Epoch 5/100
 230/230 [=====] - 8s 37ms/step - loss: 2.9182 -
 accuracy: 0.1797 - val_loss: 2.5324 - val_accuracy: 0.2548
 Epoch 6/100
 230/230 [=====] - 11s 48ms/step - loss: 2.1847 -
 accuracy: 0.3357 - val_loss: 1.8590 - val_accuracy: 0.4192
 Epoch 7/100
 230/230 [=====] - 12s 54ms/step - loss: 1.7277 -
 accuracy: 0.4684 - val_loss: 1.4882 - val_accuracy: 0.5362
 Epoch 8/100
 230/230 [=====] - 9s 41ms/step - loss: 1.3967 -
 accuracy: 0.5702 - val_loss: 1.3007 - val_accuracy: 0.6113
 Epoch 9/100
 230/230 [=====] - 11s 49ms/step - loss: 1.1765 -
 accuracy: 0.6394 - val_loss: 1.1318 - val_accuracy: 0.6717
 Epoch 10/100
 230/230 [=====] - 11s 48ms/step - loss: 0.9851 -
 accuracy: 0.6911 - val_loss: 1.0336 - val_accuracy: 0.6979
 Epoch 11/100
 230/230 [=====] - 9s 38ms/step - loss: 0.8667 -
 accuracy: 0.7352 - val_loss: 0.9671 - val_accuracy: 0.7333
 Epoch 12/100
 230/230 [=====] - 11s 46ms/step - loss: 0.7273 -
 accuracy: 0.7721 - val_loss: 0.9134 - val_accuracy: 0.7420
 Epoch 13/100
 230/230 [=====] - 10s 43ms/step - loss: 0.6583 -
 accuracy: 0.8008 - val_loss: 0.8731 - val_accuracy: 0.7599
 Epoch 14/100
 230/230 [=====] - 9s 38ms/step - loss: 0.6147 -
 accuracy: 0.8119 - val_loss: 0.8417 - val_accuracy: 0.7686
 Epoch 15/100
 230/230 [=====] - 12s 51ms/step - loss: 0.5607 -
 accuracy: 0.8238 - val_loss: 0.8010 - val_accuracy: 0.7703
 Epoch 16/100
 230/230 [=====] - 9s 39ms/step - loss: 0.5102 -
 accuracy: 0.8403 - val_loss: 0.7857 - val_accuracy: 0.7817
 Epoch 17/100
 230/230 [=====] - 11s 48ms/step - loss: 0.4737 -
 accuracy: 0.8486 - val_loss: 0.7466 - val_accuracy: 0.8019
 Epoch 18/100
 230/230 [=====] - 11s 49ms/step - loss: 0.4363 -
 accuracy: 0.8650 - val_loss: 0.7574 - val_accuracy: 0.7980
 Epoch 19/100
 230/230 [=====] - 9s 38ms/step - loss: 0.3991 -

accuracy: 0.8748 - val_loss: 0.7701 - val_accuracy: 0.7948
Epoch 20/100
230/230 [=====] - 11s 48ms/step - loss: 0.3757 -
accuracy: 0.8808 - val_loss: 0.7752 - val_accuracy: 0.7937
Training time: 210.85999464988708 seconds

```
[ ]: # Plot the loss
plt.title('Loss Plot')
plt.plot(history.history['loss'], label='train')
plt.plot(history.history['val_loss'], label='validation')
plt.legend()
plt.show()
```



```
[ ]: # Plot the accuracy
plt.title('Accuracy Plot')
plt.plot(history.history['accuracy'], label='train')
plt.plot(history.history['val_accuracy'], label='validation')
plt.legend()
plt.show()
```



```
[ ]: # Load the saved model
saved_model = tf.keras.models.load_model(model_checkpoint_path)

# Evaluate the model with the test set
loss, accuracy = saved_model.evaluate(X_test_padded, y_test_array)

print("Test Loss:", loss)
print("Test Accuracy:", round((accuracy*100), 2))
```

```
97/97 [=====] - 4s 25ms/step - loss: 0.7372 - accuracy:
0.7984
Test Loss: 0.7371559143066406
Test Accuracy: 79.84
```

```
[ ]: # Check predictions with the test set
y_test_prob = saved_model.predict(X_test_padded)

# Convert probabilities to class labels
y_test_pred = np.argmax(y_test_prob, axis=1)

# Calculate precision, recall, and f1 score
```



```

precision = precision_score(y_test_array, y_test_pred, average='weighted')
recall = recall_score(y_test_array, y_test_pred, average='weighted')
f1 = f1_score(y_test_array, y_test_pred, average='weighted')

print("Precision:", round((precision*100), 2))
print("Recall:", round((recall*100), 2))
print("F1 Score:", round((f1*100), 2))

```

```

97/97 [=====] - 2s 13ms/step
Precision: 81.46
Recall: 79.84
F1 Score: 79.54

```

```

[ ]: # Error analysis

# Print classification report
print(classification_report(y_test_array, y_test_pred))

# Check misclassified data
misclassified_data = np.where(y_test_pred != y_test_array)[0]
print(f"The number of misclassifications: {len(misclassified_data)}")

# Check the ratio of misclassifications
misclassification_ratio = (len(misclassified_data) / len(y_test_array)) * 100
# Round the number
rounded_ratio = round(misclassification_ratio, 2)
print(f"Proportion of misclassifications: {rounded_ratio}%")

# Iterate over misclassified data for error analysis
for idx in misclassified_data[:30]:
    input_text = X_test[idx]
    true_label = y_test[idx]
    predicted_label = y_test_pred[idx]

    # Print information about the misclassified data
    print("Input Text:", input_text)
    print("Actual Label:", true_label)
    print("Predicted Label:", predicted_label)
    print()

```

	precision	recall	f1-score	support
0	0.90	0.95	0.93	40
1	0.90	0.95	0.93	40
2	0.95	1.00	0.98	40
3	0.58	0.70	0.64	40
4	1.00	0.85	0.92	40
5	0.70	0.65	0.68	40

6	1.00	0.75	0.86	40
7	0.79	0.68	0.73	40
8	0.73	0.90	0.81	40
9	0.85	0.97	0.91	40
10	0.79	0.57	0.67	40
11	0.87	0.68	0.76	40
12	0.65	0.90	0.76	40
13	0.93	0.95	0.94	40
14	0.81	0.85	0.83	40
15	0.80	0.80	0.80	40
16	0.50	0.57	0.53	40
17	0.79	0.93	0.85	40
18	0.92	0.82	0.87	40
19	0.88	0.95	0.92	40
20	0.67	0.90	0.77	40
21	0.97	0.95	0.96	40
22	0.67	0.72	0.70	40
23	1.00	0.88	0.93	40
24	0.88	0.93	0.90	40
25	0.83	0.88	0.85	40
26	0.78	0.72	0.75	40
27	0.97	0.70	0.81	40
28	0.88	0.75	0.81	40
29	0.88	0.75	0.81	40
30	0.91	0.97	0.94	40
31	1.00	0.85	0.92	40
32	0.90	0.93	0.91	40
33	0.81	0.85	0.83	40
34	0.79	0.85	0.82	40
35	0.71	0.72	0.72	40
36	0.84	0.78	0.81	40
37	0.63	0.82	0.72	40
38	0.82	0.93	0.87	40
39	0.83	0.88	0.85	40
40	0.74	0.97	0.84	40
41	0.82	0.70	0.76	40
42	0.80	0.88	0.83	40
43	0.82	0.68	0.74	40
44	1.00	1.00	1.00	40
45	0.95	0.88	0.91	40
46	0.71	0.88	0.79	40
47	0.71	0.85	0.77	40
48	0.72	0.72	0.73	40
49	0.88	0.75	0.81	40
50	0.86	0.78	0.82	40
51	0.93	0.68	0.78	40
52	0.61	0.95	0.75	40
53	0.74	0.50	0.60	40

54	0.79	0.78	0.78	40
55	0.93	0.93	0.93	40
56	0.77	0.75	0.76	40
57	0.88	0.88	0.88	40
58	0.56	0.85	0.67	40
59	0.80	0.70	0.75	40
60	0.93	0.95	0.94	40
61	0.77	0.57	0.66	40
62	0.87	0.65	0.74	40
63	0.76	0.88	0.81	40
64	0.79	0.85	0.82	40
65	0.79	0.75	0.77	40
66	0.66	0.68	0.67	40
67	0.68	0.65	0.67	40
68	0.75	0.30	0.43	40
69	0.64	0.23	0.33	40
70	0.95	0.97	0.96	40
71	0.95	1.00	0.98	40
72	0.93	0.33	0.48	40
73	1.00	0.85	0.92	40
74	0.43	0.93	0.58	40
75	0.69	0.82	0.75	40
76	0.97	0.82	0.89	40
accuracy				0.80
macro avg				0.81
weighted avg				0.80

The number of misclassifications: 621
Proportion of misclassifications: 20.16%
Input Text: locate card
Actual Label: 11
Predicted Label: 13

Input Text: way know card arrive
Actual Label: 11
Predicted Label: 12

Input Text: get card
Actual Label: 11
Predicted Label: 12

Input Text: received card
Actual Label: 11
Predicted Label: 12

Input Text: long card delivery take
Actual Label: 11

Predicted Label: 12

Input Text: hasnt card delivered

Actual Label: 11

Predicted Label: 12

Input Text: get card yet lost

Actual Label: 11

Predicted Label: 41

Input Text: status card ordered

Actual Label: 11

Predicted Label: 12

Input Text: im still waiting delivery new card taking long

Actual Label: 11

Predicted Label: 9

Input Text: know card arrive

Actual Label: 11

Predicted Label: 12

Input Text: im still waiting card delivered

Actual Label: 11

Predicted Label: 12

Input Text: im starting think card lost still hasnt arrived help

Actual Label: 11

Predicted Label: 41

Input Text: tracking info available

Actual Label: 11

Predicted Label: 68

Input Text: add card account

Actual Label: 13

Predicted Label: 39

Input Text: link another card account

Actual Label: 13

Predicted Label: 39

Input Text: good time exchange

Actual Label: 32

Predicted Label: 17

Input Text: currencies exchange rate calculated

Actual Label: 32

Predicted Label: 17

Input Text: im trying figure current exchange rate

Actual Label: 32

Predicted Label: 17

Input Text: rate low sure using right exchange rate

Actual Label: 17

Predicted Label: 32

Input Text: charged

Actual Label: 17

Predicted Label: 34

Input Text: conversion value card payments incorrect

Actual Label: 17

Predicted Label: 15

Input Text: explain random charge

Actual Label: 34

Predicted Label: 16

Input Text: transaction credited

Actual Label: 34

Predicted Label: 8

Input Text: fee come

Actual Label: 34

Predicted Label: 15

Input Text: euro fee come

Actual Label: 34

Predicted Label: 29

Input Text: euro fee statement

Actual Label: 34

Predicted Label: 22

Input Text: two weeks transaction reversed

Actual Label: 34

Predicted Label: 63

Input Text: made withdrawal account posted

Actual Label: 46

Predicted Label: 20

Input Text: wheres accounting cash withdrawal

Actual Label: 46

Predicted Label: 20

Input Text: account charged withdraw tried make decline

Actual Label: 46

Predicted Label: 19

1.6 LSTM (with GloVe)

1.6.1 Set up the GloVe model

```
[ ]: # Define the embedding index as a dictionary
embeddings_index = dict()

# Load the GloVe file to the colab notebook
f = open('/content/drive/MyDrive/1. NLP CW/glove.6B.100d.txt', encoding="utf8")
for line in f:
    values = line.split()
    word = values[0]
    coefs = np.asarray(values[1:], dtype='float32')
    embeddings_index[word] = coefs
f.close()

print(f'Total vectors in the GloVe file: {len(embeddings_index)}')
```

Total vectors in the GloVe file: 400001

```
[ ]: # Check the dimension of a word
print(f"The dimension of a word in embedding_index: {len(embeddings_index['card'])}")
```

The dimension of a word in embedding_index: 100

```
[ ]: # Define the embedding matrix
embedding_matrix = np.zeros((voca_size+1, 100)) # For future dimension matching
# with the word_index, add 1 to the vocabulary size, and match 100 from GloVe
print(f'The shape of embedding matrix: {np.shape(embedding_matrix)}')
```

The shape of embedding matrix: (2089, 100)

```
[ ]: # Check the index of the word 'card' in word_index
print(f"The index of card in word_index: {word_index['card']}")
```

The index of card in word_index: 1

```
[ ]: # Check the word 'card' vectors in the embedding index
print(embeddings_index['card'])
```

```
[ 1.6292e-01 -3.1798e-01  4.2328e-01 -8.6767e-01  4.5101e-01  5.7857e-01
 2.6645e-02 -1.2648e-01  3.3465e-01 -4.2047e-02 -4.0596e-02  1.6478e-01
-6.7344e-01 -3.3751e-01  3.5913e-01  5.7383e-01  8.4620e-01  3.6374e-01
 3.0630e-01 -6.8050e-02 -6.7610e-01 -1.9147e-01 -1.4594e-01  3.2621e-03
 6.6949e-01 -3.3588e-01  1.7868e-01 -3.9360e-01  1.7700e-01 -3.3642e-01
 1.9288e-01  1.0030e+00 -2.1794e-01  2.4271e-01  1.0935e+00 -1.0303e-01
-7.9197e-01 -1.3506e-01  1.2156e-01 -9.8377e-01  1.0300e+00 -1.0242e+00
 6.0269e-01 -1.5986e-01 -2.6773e-01 -5.5630e-01  2.5834e-01 -8.5021e-02
-1.5221e-01 -3.3717e-01  2.6358e-02  2.3171e-01 -1.8056e-01  5.7107e-01
 3.8556e-01 -1.5732e+00 -1.4902e-01  3.7826e-02  1.8485e+00  7.0210e-01
-1.1697e-01  7.7822e-02  7.4620e-02  9.9570e-02 -2.1427e-01 -6.0061e-01
 9.4903e-02  8.0589e-01  5.5333e-01 -3.1359e-01 -9.0991e-01  5.3645e-02
-1.4494e-01 -4.8532e-01  1.0335e-01  1.2182e+00 -2.2199e-01 -1.4934e-02
-1.1355e+00  3.2790e-01  1.1733e+00 -5.2838e-01 -6.6953e-01 -6.2109e-01
-1.3660e+00 -4.4052e-01 -2.9538e-01 -7.1655e-01  5.9920e-01 -3.4550e-04
-8.2363e-01  9.3572e-01  6.2134e-01 -2.6649e-01  9.9595e-02 -1.1545e-01
 6.0000e-01  7.2834e-02  6.6487e-01 -6.4510e-01]
```

```
[ ]: # Check the word 'topup' vectors in the embedding index
print(embeddings_index['topup'])
```

```
-----
KeyError                                Traceback (most recent call last)
<ipython-input-153-c4db3aaada12> in <cell line: 2>()
      1 # Check the word 'topup' vectors in the embedding index
----> 2 print(embeddings_index['topup'])

KeyError: 'topup'
```

There is no word 'topup'.

```
[ ]: # Match the index and word to create an embedding matrix
for word, index in word_index.items():
    vector_value = embeddings_index.get(word)
    if vector_value is not None:
        embedding_matrix[index] = vector_value
```

```
[ ]: # Check the word 'card' vectors in the embedding matrix
print(embedding_matrix[1])
```

```
[ 1.62919998e-01 -3.17979991e-01  4.23280001e-01 -8.67670000e-01
 4.51009989e-01  5.78570008e-01  2.66449992e-02 -1.26479998e-01
 3.34650010e-01 -4.20470014e-02 -4.05960009e-02  1.64780006e-01
-6.73439980e-01 -3.37509990e-01  3.59129995e-01  5.73830009e-01
 8.46199989e-01  3.63739997e-01  3.06300014e-01 -6.80499971e-02
-6.76100016e-01 -1.91469997e-01 -1.45940006e-01  3.26210004e-03
 6.69489980e-01 -3.35880011e-01  1.78680003e-01 -3.93599987e-01]
```

```

1.77000001e-01 -3.36420000e-01 1.92880005e-01 1.00300002e+00
-2.17940003e-01 2.42709994e-01 1.09350002e+00 -1.03030004e-01
-7.91970015e-01 -1.35059997e-01 1.21560000e-01 -9.83770013e-01
1.02999997e+00 -1.02419996e+00 6.02689981e-01 -1.59860000e-01
-2.67729998e-01 -5.56299984e-01 2.58340001e-01 -8.50209966e-02
-1.52209997e-01 -3.37170005e-01 2.63579991e-02 2.31710002e-01
-1.80559993e-01 5.71070015e-01 3.85560006e-01 -1.57319999e+00
-1.49020001e-01 3.78260016e-02 1.84850001e+00 7.02099979e-01
-1.16970003e-01 7.78219998e-02 7.46200010e-02 9.95699987e-02
-2.14269996e-01 -6.00610018e-01 9.49029997e-02 8.05890024e-01
5.53330004e-01 -3.13589990e-01 -9.09910023e-01 5.36449999e-02
-1.44940004e-01 -4.85320002e-01 1.03349999e-01 1.21819997e+00
-2.21990004e-01 -1.49339996e-02 -1.13549995e+00 3.27899992e-01
1.17330003e+00 -5.28379977e-01 -6.69529974e-01 -6.21089995e-01
-1.36600006e+00 -4.40519989e-01 -2.95379996e-01 -7.16549993e-01
5.99200010e-01 -3.45500011e-04 -8.23629975e-01 9.35720026e-01
6.21339977e-01 -2.66490012e-01 9.95950028e-02 -1.15450002e-01
6.00000024e-01 7.28340000e-02 6.64870024e-01 -6.45099998e-01]

```

1.6.2 LSTM (baseline)

```

[ ]: # Define the output dimension for the embedding layer and hidden units
hidden_unit = 30
nlabel = 77

model = keras.models.Sequential()
e = layers.Embedding(voca_size+1, 100, weights=[embedding_matrix],
    ↳input_length=max_length_train_text, trainable=False) # Using 100 dimension
    ↳for GloVe
model.add(e)
model.add(layers.LSTM(hidden_unit))
model.add(layers.Dense(nlabel, activation='softmax'))

# Compile the model
model.compile(optimizer='adam', loss='sparse_categorical_crossentropy',
    ↳metrics=['accuracy']) #, run_eagerly=True

# Summary the model
model.summary()

```

Model: "sequential_10"

Layer (type)	Output Shape	Param #
embedding_10 (Embedding)	(None, 29, 100)	208900
lstm_10 (LSTM)	(None, 30)	15720

dense_10 (Dense) (None, 77) 2387

```
=====
Total params: 227007 (886.75 KB)
Trainable params: 18107 (70.73 KB)
Non-trainable params: 208900 (816.02 KB)
-----
```

```
[ ]: # Define the folder path to save the model
folder_path = '/content/drive/MyDrive/1. NLP CW/LSTM/'

# Define the file path for the model checkpoint
model_checkpoint_path = folder_path + 'LSTM_glove_model.keras'

# Define the model checkpoint
mc = tf.keras.callbacks.ModelCheckpoint(
    filepath=model_checkpoint_path,
    monitor='val_accuracy',
    mode='max',
    save_best_only=True)

[ ]: # Define early stopping
es = tf.keras.callbacks.EarlyStopping(monitor='val_loss', patience=3) # Random_
    ↪ number of patience

[ ]: # Import time to measure the elapsed time
import time

# Measure time before training
start_time = time.time()

# Fit the model
history = model.fit(
    X_train_padded, y_train,
    epochs = 100,
    validation_data = (X_val_padded, y_val),
    callbacks = [mc, es],
    batch_size = 32)

# End the training time
end_time = time.time()

# Measure the training time
training_time = end_time - start_time
print("Training time:", training_time, "seconds")
```

Epoch 1/100

230/230 [=====] - 11s 31ms/step - loss: 4.0977 -
accuracy: 0.0253 - val_loss: 3.8938 - val_accuracy: 0.0267
Epoch 2/100
230/230 [=====] - 4s 19ms/step - loss: 3.8196 -
accuracy: 0.0376 - val_loss: 3.7330 - val_accuracy: 0.0435
Epoch 3/100
230/230 [=====] - 4s 19ms/step - loss: 3.6594 -
accuracy: 0.0561 - val_loss: 3.5586 - val_accuracy: 0.0648
Epoch 4/100
230/230 [=====] - 6s 26ms/step - loss: 3.4984 -
accuracy: 0.0724 - val_loss: 3.4428 - val_accuracy: 0.0768
Epoch 5/100
230/230 [=====] - 4s 19ms/step - loss: 3.3714 -
accuracy: 0.0779 - val_loss: 3.3062 - val_accuracy: 0.0936
Epoch 6/100
230/230 [=====] - 4s 19ms/step - loss: 3.2583 -
accuracy: 0.0873 - val_loss: 3.2377 - val_accuracy: 0.0925
Epoch 7/100
230/230 [=====] - 5s 22ms/step - loss: 3.1858 -
accuracy: 0.1032 - val_loss: 3.1524 - val_accuracy: 0.1132
Epoch 8/100
230/230 [=====] - 5s 20ms/step - loss: 3.1273 -
accuracy: 0.1171 - val_loss: 3.1484 - val_accuracy: 0.1083
Epoch 9/100
230/230 [=====] - 4s 18ms/step - loss: 3.0075 -
accuracy: 0.1261 - val_loss: 3.0145 - val_accuracy: 0.1366
Epoch 10/100
230/230 [=====] - 4s 19ms/step - loss: 2.9693 -
accuracy: 0.1262 - val_loss: 2.9630 - val_accuracy: 0.1334
Epoch 11/100
230/230 [=====] - 6s 27ms/step - loss: 2.8887 -
accuracy: 0.1364 - val_loss: 2.9393 - val_accuracy: 0.1366
Epoch 12/100
230/230 [=====] - 5s 20ms/step - loss: 2.8243 -
accuracy: 0.1481 - val_loss: 2.8853 - val_accuracy: 0.1497
Epoch 13/100
230/230 [=====] - 4s 19ms/step - loss: 2.8182 -
accuracy: 0.1501 - val_loss: 2.8631 - val_accuracy: 0.1595
Epoch 14/100
230/230 [=====] - 6s 28ms/step - loss: 2.7213 -
accuracy: 0.1695 - val_loss: 2.7291 - val_accuracy: 0.1796
Epoch 15/100
230/230 [=====] - 4s 19ms/step - loss: 2.6767 -
accuracy: 0.1853 - val_loss: 2.6960 - val_accuracy: 0.1791
Epoch 16/100
230/230 [=====] - 4s 19ms/step - loss: 2.6447 -
accuracy: 0.1854 - val_loss: 2.6630 - val_accuracy: 0.1981
Epoch 17/100

230/230 [=====] - 6s 25ms/step - loss: 2.5574 - accuracy: 0.2058 - val_loss: 2.7203 - val_accuracy: 0.1747
Epoch 18/100
230/230 [=====] - 4s 19ms/step - loss: 2.5111 - accuracy: 0.2171 - val_loss: 2.5729 - val_accuracy: 0.2248
Epoch 19/100
230/230 [=====] - 4s 16ms/step - loss: 2.4451 - accuracy: 0.2340 - val_loss: 2.5427 - val_accuracy: 0.2172
Epoch 20/100
230/230 [=====] - 5s 23ms/step - loss: 2.4228 - accuracy: 0.2322 - val_loss: 2.5926 - val_accuracy: 0.2041
Epoch 21/100
230/230 [=====] - 5s 21ms/step - loss: 2.3821 - accuracy: 0.2520 - val_loss: 2.4346 - val_accuracy: 0.2727
Epoch 22/100
230/230 [=====] - 4s 18ms/step - loss: 2.3258 - accuracy: 0.2723 - val_loss: 2.3848 - val_accuracy: 0.2657
Epoch 23/100
230/230 [=====] - 7s 32ms/step - loss: 2.2492 - accuracy: 0.2883 - val_loss: 2.3145 - val_accuracy: 0.2782
Epoch 24/100
230/230 [=====] - 7s 29ms/step - loss: 2.2229 - accuracy: 0.2875 - val_loss: 2.2886 - val_accuracy: 0.2825
Epoch 25/100
230/230 [=====] - 5s 23ms/step - loss: 2.1747 - accuracy: 0.3068 - val_loss: 2.3042 - val_accuracy: 0.2896
Epoch 26/100
230/230 [=====] - 6s 25ms/step - loss: 2.2218 - accuracy: 0.3072 - val_loss: 2.2399 - val_accuracy: 0.3114
Epoch 27/100
230/230 [=====] - 4s 18ms/step - loss: 2.0938 - accuracy: 0.3252 - val_loss: 2.1766 - val_accuracy: 0.3217
Epoch 28/100
230/230 [=====] - 5s 20ms/step - loss: 2.0396 - accuracy: 0.3357 - val_loss: 2.1520 - val_accuracy: 0.3304
Epoch 29/100
230/230 [=====] - 5s 23ms/step - loss: 1.9676 - accuracy: 0.3560 - val_loss: 2.1121 - val_accuracy: 0.3386
Epoch 30/100
230/230 [=====] - 5s 21ms/step - loss: 1.9361 - accuracy: 0.3601 - val_loss: 2.0977 - val_accuracy: 0.3375
Epoch 31/100
230/230 [=====] - 4s 18ms/step - loss: 1.9150 - accuracy: 0.3695 - val_loss: 2.0477 - val_accuracy: 0.3457
Epoch 32/100
230/230 [=====] - 5s 21ms/step - loss: 1.8701 - accuracy: 0.3891 - val_loss: 2.0085 - val_accuracy: 0.3778
Epoch 33/100

230/230 [=====] - 6s 25ms/step - loss: 1.8323 - accuracy: 0.3998 - val_loss: 2.0175 - val_accuracy: 0.3702
Epoch 34/100
230/230 [=====] - 4s 19ms/step - loss: 1.7985 - accuracy: 0.4123 - val_loss: 1.9501 - val_accuracy: 0.3832
Epoch 35/100
230/230 [=====] - 4s 19ms/step - loss: 1.7482 - accuracy: 0.4300 - val_loss: 1.9111 - val_accuracy: 0.3958
Epoch 36/100
230/230 [=====] - 6s 27ms/step - loss: 1.7472 - accuracy: 0.4306 - val_loss: 1.9328 - val_accuracy: 0.4012
Epoch 37/100
230/230 [=====] - 4s 19ms/step - loss: 1.6903 - accuracy: 0.4510 - val_loss: 1.9369 - val_accuracy: 0.4056
Epoch 38/100
230/230 [=====] - 4s 19ms/step - loss: 1.6844 - accuracy: 0.4468 - val_loss: 1.8676 - val_accuracy: 0.4175
Epoch 39/100
230/230 [=====] - 6s 28ms/step - loss: 1.6393 - accuracy: 0.4631 - val_loss: 2.0641 - val_accuracy: 0.3767
Epoch 40/100
230/230 [=====] - 4s 18ms/step - loss: 1.6475 - accuracy: 0.4624 - val_loss: 1.8607 - val_accuracy: 0.4279
Epoch 41/100
230/230 [=====] - 5s 20ms/step - loss: 1.5949 - accuracy: 0.4741 - val_loss: 1.8212 - val_accuracy: 0.4197
Epoch 42/100
230/230 [=====] - 6s 26ms/step - loss: 1.5768 - accuracy: 0.4857 - val_loss: 1.8403 - val_accuracy: 0.4415
Epoch 43/100
230/230 [=====] - 7s 29ms/step - loss: 1.5028 - accuracy: 0.5140 - val_loss: 1.7729 - val_accuracy: 0.4529
Epoch 44/100
230/230 [=====] - 4s 19ms/step - loss: 1.5005 - accuracy: 0.5244 - val_loss: 1.7320 - val_accuracy: 0.4780
Epoch 45/100
230/230 [=====] - 6s 28ms/step - loss: 1.4529 - accuracy: 0.5357 - val_loss: 1.6892 - val_accuracy: 0.4823
Epoch 46/100
230/230 [=====] - 4s 20ms/step - loss: 1.4240 - accuracy: 0.5520 - val_loss: 1.6691 - val_accuracy: 0.4932
Epoch 47/100
230/230 [=====] - 4s 18ms/step - loss: 1.3856 - accuracy: 0.5652 - val_loss: 1.6269 - val_accuracy: 0.5046
Epoch 48/100
230/230 [=====] - 6s 26ms/step - loss: 1.3377 - accuracy: 0.5784 - val_loss: 1.6021 - val_accuracy: 0.5161
Epoch 49/100

230/230 [=====] - 4s 18ms/step - loss: 1.3058 - accuracy: 0.5849 - val_loss: 1.6068 - val_accuracy: 0.5253
Epoch 50/100
230/230 [=====] - 4s 18ms/step - loss: 1.2801 - accuracy: 0.5923 - val_loss: 1.5313 - val_accuracy: 0.5444
Epoch 51/100
230/230 [=====] - 5s 24ms/step - loss: 1.2841 - accuracy: 0.5953 - val_loss: 1.5498 - val_accuracy: 0.5542
Epoch 52/100
230/230 [=====] - 5s 20ms/step - loss: 1.2555 - accuracy: 0.6066 - val_loss: 1.4948 - val_accuracy: 0.5509
Epoch 53/100
230/230 [=====] - 4s 18ms/step - loss: 1.1757 - accuracy: 0.6262 - val_loss: 1.4880 - val_accuracy: 0.5651
Epoch 54/100
230/230 [=====] - 5s 21ms/step - loss: 1.1375 - accuracy: 0.6418 - val_loss: 1.4567 - val_accuracy: 0.5863
Epoch 55/100
230/230 [=====] - 6s 26ms/step - loss: 1.1200 - accuracy: 0.6525 - val_loss: 1.4595 - val_accuracy: 0.5770
Epoch 56/100
230/230 [=====] - 4s 19ms/step - loss: 1.1394 - accuracy: 0.6474 - val_loss: 1.4413 - val_accuracy: 0.5901
Epoch 57/100
230/230 [=====] - 5s 20ms/step - loss: 1.1099 - accuracy: 0.6604 - val_loss: 1.4607 - val_accuracy: 0.5836
Epoch 58/100
230/230 [=====] - 6s 27ms/step - loss: 1.1738 - accuracy: 0.6447 - val_loss: 1.4562 - val_accuracy: 0.5797
Epoch 59/100
230/230 [=====] - 4s 19ms/step - loss: 1.0876 - accuracy: 0.6695 - val_loss: 1.4259 - val_accuracy: 0.6179
Epoch 60/100
230/230 [=====] - 4s 18ms/step - loss: 1.0330 - accuracy: 0.6914 - val_loss: 1.3771 - val_accuracy: 0.6200
Epoch 61/100
230/230 [=====] - 6s 26ms/step - loss: 1.0152 - accuracy: 0.6910 - val_loss: 1.4332 - val_accuracy: 0.6162
Epoch 62/100
230/230 [=====] - 4s 18ms/step - loss: 0.9939 - accuracy: 0.7056 - val_loss: 1.4023 - val_accuracy: 0.6184
Epoch 63/100
230/230 [=====] - 5s 20ms/step - loss: 0.9923 - accuracy: 0.6983 - val_loss: 1.3735 - val_accuracy: 0.6320
Epoch 64/100
230/230 [=====] - 6s 25ms/step - loss: 0.9576 - accuracy: 0.7137 - val_loss: 1.4352 - val_accuracy: 0.6037
Epoch 65/100

```

230/230 [=====] - 5s 23ms/step - loss: 0.9878 -
accuracy: 0.7065 - val_loss: 1.3652 - val_accuracy: 0.6364
Epoch 66/100
230/230 [=====] - 4s 19ms/step - loss: 0.8919 -
accuracy: 0.7356 - val_loss: 1.3525 - val_accuracy: 0.6380
Epoch 67/100
230/230 [=====] - 5s 22ms/step - loss: 0.8803 -
accuracy: 0.7411 - val_loss: 1.4911 - val_accuracy: 0.6015
Epoch 68/100
230/230 [=====] - 5s 23ms/step - loss: 1.1530 -
accuracy: 0.6639 - val_loss: 1.3941 - val_accuracy: 0.6233
Epoch 69/100
230/230 [=====] - 4s 17ms/step - loss: 0.9495 -
accuracy: 0.7213 - val_loss: 1.3490 - val_accuracy: 0.6473
Epoch 70/100
230/230 [=====] - 5s 20ms/step - loss: 0.8977 -
accuracy: 0.7389 - val_loss: 1.3274 - val_accuracy: 0.6570
Epoch 71/100
230/230 [=====] - 6s 24ms/step - loss: 0.8398 -
accuracy: 0.7610 - val_loss: 1.3034 - val_accuracy: 0.6658
Epoch 72/100
230/230 [=====] - 4s 17ms/step - loss: 0.8164 -
accuracy: 0.7607 - val_loss: 1.3009 - val_accuracy: 0.6625
Epoch 73/100
230/230 [=====] - 4s 18ms/step - loss: 0.8154 -
accuracy: 0.7616 - val_loss: 1.3166 - val_accuracy: 0.6614
Epoch 74/100
230/230 [=====] - 6s 27ms/step - loss: 0.7957 -
accuracy: 0.7712 - val_loss: 1.2847 - val_accuracy: 0.6821
Epoch 75/100
230/230 [=====] - 4s 19ms/step - loss: 0.8666 -
accuracy: 0.7495 - val_loss: 1.4881 - val_accuracy: 0.6200
Epoch 76/100
230/230 [=====] - 4s 18ms/step - loss: 0.8429 -
accuracy: 0.7552 - val_loss: 1.3415 - val_accuracy: 0.6652
Epoch 77/100
230/230 [=====] - 5s 21ms/step - loss: 0.8279 -
accuracy: 0.7642 - val_loss: 1.2907 - val_accuracy: 0.6701
Training time: 385.54614877700806 seconds

```


```

[ ]: # Plot the loss
plt.title('Loss Plot')
plt.plot(history.history['loss'], label='train')
plt.plot(history.history['val_loss'], label='validation')
plt.legend()
plt.show()

```



```
[ ]: # Plot the accuracy
plt.title('Accuracy Plot')
plt.plot(history.history['accuracy'], label='train')
plt.plot(history.history['val_accuracy'], label='validation')
plt.legend()
plt.show()
```



```
[ ]: # Load the saved model
saved_model = tf.keras.models.load_model(model_checkpoint_path)

# Evaluate the model with the test set
loss, accuracy = saved_model.evaluate(X_test_padded, y_test_array)

print("Test Loss:", loss)
print("Test Accuracy:", round((accuracy*100), 2))
```

```
97/97 [=====] - 2s 7ms/step - loss: 1.2908 - accuracy:
0.6769
Test Loss: 1.2908092737197876
Test Accuracy: 67.69
```

```
[ ]: # Check predictions with the test set
y_test_prob = saved_model.predict(X_test_padded)

# Convert probabilities to class labels
y_test_pred = np.argmax(y_test_prob, axis=1)

# Calculate precision, recall, and f1 score
```



```
precision = precision_score(y_test_array, y_test_pred, average='weighted')
recall = recall_score(y_test_array, y_test_pred, average='weighted')
f1 = f1_score(y_test_array, y_test_pred, average='weighted')

print("Precision:", round((precision*100), 2))
print("Recall:", round((recall*100), 2))
print("F1 Score:", round((f1*100), 2))
```

97/97 [=====] - 2s 9ms/step

Precision: 67.88

Recall: 67.69

F1 Score: 66.03

/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344:
 UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels
 with no predicted samples. Use `zero_division` parameter to control this
 behavior.

```
_warn_prf(average, modifier, msg_start, len(result))
```

[]: *# Error analysis*

```
# Print classification report
print(classification_report(y_test_array, y_test_pred))

# Check misclassified data
misclassified_data = np.where(y_test_pred != y_test_array)[0]
print(f"The number of misclassifications: {len(misclassified_data)}")

# Check the ratio of misclassifications
misclassification_ratio = (len(misclassified_data) / len(y_test_array)) * 100
# Round the number
rounded_ratio = round(misclassification_ratio, 2)
print(f"Proportion of misclassifications: {rounded_ratio}%")

# Iterate over misclassified data for error analysis
for idx in misclassified_data[:30]:
    input_text = X_test[idx]
    true_label = y_test[idx]
    predicted_label = y_test_pred[idx]

    # Print information about the misclassified data
    print("Input Text:", input_text)
    print("Actual Label:", true_label)
    print("Predicted Label:", predicted_label)
    print()
```

```
precision    recall  f1-score   support
```

0	0.87	0.85	0.86	40
1	0.90	0.95	0.93	40
2	0.97	0.97	0.97	40
3	0.79	0.65	0.71	40
4	0.77	0.68	0.72	40
5	0.46	0.68	0.55	40
6	0.80	0.90	0.85	40
7	0.62	0.45	0.52	40
8	0.71	0.85	0.77	40
9	0.90	0.93	0.91	40
10	0.50	0.05	0.09	40
11	0.45	0.75	0.57	40
12	0.64	0.17	0.27	40
13	0.74	0.80	0.77	40
14	0.37	0.55	0.44	40
15	0.74	0.80	0.77	40
16	0.64	0.62	0.63	40
17	0.85	0.82	0.84	40
18	0.66	0.62	0.64	40
19	0.80	0.90	0.85	40
20	0.74	0.57	0.65	40
21	0.70	0.75	0.72	40
22	0.40	0.62	0.49	40
23	0.96	0.68	0.79	40
24	0.66	0.88	0.75	40
25	0.53	0.68	0.59	40
26	0.64	0.75	0.69	40
27	0.75	0.68	0.71	40
28	0.70	0.75	0.72	40
29	0.60	0.80	0.69	40
30	0.82	0.93	0.87	40
31	0.89	0.80	0.84	40
32	0.88	0.88	0.88	40
33	0.71	0.90	0.79	40
34	0.78	0.72	0.75	40
35	0.61	0.50	0.55	40
36	0.78	0.62	0.69	40
37	0.75	0.23	0.35	40
38	0.58	0.75	0.65	40
39	0.57	0.78	0.66	40
40	0.44	0.97	0.61	40
41	0.51	0.65	0.57	40
42	0.91	0.75	0.82	40
43	0.46	0.45	0.46	40
44	0.92	0.88	0.90	40
45	0.72	0.72	0.73	40
46	0.64	0.75	0.69	40
47	0.61	0.75	0.67	40

48	0.73	0.47	0.58	40
49	0.73	0.28	0.40	40
50	0.63	0.60	0.62	40
51	0.74	0.93	0.82	40
52	0.73	0.60	0.66	40
53	0.85	0.70	0.77	40
54	0.61	0.62	0.62	40
55	0.85	0.82	0.84	40
56	0.76	0.47	0.58	40
57	0.97	0.85	0.91	40
58	0.82	0.68	0.74	40
59	0.56	0.57	0.57	40
60	0.91	0.75	0.82	40
61	0.65	0.65	0.65	40
62	0.58	0.53	0.55	40
63	0.84	0.93	0.88	40
64	0.63	0.78	0.70	40
65	0.63	0.65	0.64	40
66	0.50	0.68	0.57	40
67	0.44	0.40	0.42	40
68	0.33	0.05	0.09	40
69	0.00	0.00	0.00	40
70	0.92	0.82	0.87	40
71	0.80	0.93	0.86	40
72	0.00	0.00	0.00	40
73	0.82	0.82	0.82	40
74	0.37	0.97	0.53	40
75	0.64	0.75	0.69	40
76	0.74	0.65	0.69	40

accuracy			0.68	3080
macro avg	0.68	0.68	0.66	3080
weighted avg	0.68	0.68	0.66	3080

The number of misclassifications: 995
Proportion of misclassifications: 32.31%
Input Text: locate card
Actual Label: 11
Predicted Label: 43

Input Text: card arrived yet
Actual Label: 11
Predicted Label: 43

Input Text: get card
Actual Label: 11
Predicted Label: 43

Input Text: know tracking number new card sent
Actual Label: 11
Predicted Label: 9

Input Text: received card
Actual Label: 11
Predicted Label: 43

Input Text: still waiting card
Actual Label: 11
Predicted Label: 43

Input Text: normal wait week new card
Actual Label: 11
Predicted Label: 9

Input Text: get card yet lost
Actual Label: 11
Predicted Label: 39

Input Text: card arrived yet
Actual Label: 11
Predicted Label: 43

Input Text: tracking info available
Actual Label: 11
Predicted Label: 71

Input Text: add card account
Actual Label: 13
Predicted Label: 39

Input Text: put old card back system found
Actual Label: 13
Predicted Label: 39

Input Text: hello found card misplaced need reactive
Actual Label: 13
Predicted Label: 0

Input Text: view card received app
Actual Label: 13
Predicted Label: 41

Input Text: found card add app
Actual Label: 13
Predicted Label: 39

Input Text: link credit card

Actual Label: 13

Predicted Label: 41

Input Text: reactivate lost card found morning jacket pocket

Actual Label: 13

Predicted Label: 42

Input Text: app doesnt show card received

Actual Label: 13

Predicted Label: 16

Input Text: exchange rates offer

Actual Label: 32

Predicted Label: 31

Input Text: international exchange rates

Actual Label: 32

Predicted Label: 31

Input Text: good time exchange

Actual Label: 32

Predicted Label: 31

Input Text: currencies exchange rate calculated

Actual Label: 32

Predicted Label: 31

Input Text: rate get determined

Actual Label: 32

Predicted Label: 76

Input Text: made currency exchange think charged

Actual Label: 17

Predicted Label: 32

Input Text: rate exchange card payment incorrect

Actual Label: 17

Predicted Label: 76

Input Text: exchange rate card payment wrong

Actual Label: 17

Predicted Label: 76

Input Text: charged

Actual Label: 17

Predicted Label: 34

Input Text: conversion value card payments incorrect
Actual Label: 17
Predicted Label: 45

Input Text: exchange rate totally wrong card payment
Actual Label: 17
Predicted Label: 76

Input Text: wrong rate applied item bought currency different mine changed
Actual Label: 17
Predicted Label: 76

```
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344:  
UndefinedMetricWarning: Precision and F-score are ill-defined and being set to  
0.0 in labels with no predicted samples. Use `zero_division` parameter to  
control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))  
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344:  
UndefinedMetricWarning: Precision and F-score are ill-defined and being set to  
0.0 in labels with no predicted samples. Use `zero_division` parameter to  
control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))  
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344:  
UndefinedMetricWarning: Precision and F-score are ill-defined and being set to  
0.0 in labels with no predicted samples. Use `zero_division` parameter to  
control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

1.6.3 LSTM (with dropout)

```
[ ]: # Define the output dimension for the embedding layer and hidden units  
hidden_unit = 30  
nlabel = 77  
  
dropout_model = keras.models.Sequential()  
e = layers.Embedding(voca_size+1, 100, weights=[embedding_matrix],  
    ↪input_length=max_length_train_text, trainable=False) # Using 100 dimension  
    ↪for GloVe  
dropout_model.add(e)  
dropout_model.add(layers.LSTM(hidden_unit, dropout=0.2))  
dropout_model.add(layers.Dense(nlabel, activation='softmax'))  
  
# Compile the model  
dropout_model.compile(optimizer='adam', loss='sparse_categorical_crossentropy',  
    ↪metrics=['accuracy']) #, run_eagerly=True  
  
# Summary the model
```

```
dropout_model.summary()
```

Model: "sequential_11"

Layer (type)	Output Shape	Param #
embedding_11 (Embedding)	(None, 29, 100)	208900
lstm_11 (LSTM)	(None, 30)	15720
dense_11 (Dense)	(None, 77)	2387

=====
Total params: 227007 (886.75 KB)
Trainable params: 18107 (70.73 KB)
Non-trainable params: 208900 (816.02 KB)
=====

```
[ ]: # Define the folder path to save the model
folder_path = '/content/drive/MyDrive/1. NLP CW/LSTM/'

# Define the file path for the model checkpoint
model_checkpoint_path = folder_path + 'dropout_LSTM_glove_model.keras'

# Define the model checkpoint
mc = tf.keras.callbacks.ModelCheckpoint(
    filepath=model_checkpoint_path,
    monitor='val_accuracy',
    mode='max',
    save_best_only=True)
```

```
[ ]: # Import time to measure the elapsed time
import time

# Measure time before training
start_time = time.time()

# Fit the model
history = dropout_model.fit(
    X_train_padded, y_train,
    epochs = 100,
    validation_data = (X_val_padded, y_val),
    callbacks = [mc, es],
    batch_size = 32)

# End the training time
end_time = time.time()
```

```
# Measure the training time
training_time = end_time - start_time
print("Training time:", training_time, "seconds")
```

```
Epoch 1/100
230/230 [=====] - 9s 25ms/step - loss: 4.1219 -
accuracy: 0.0373 - val_loss: 3.7914 - val_accuracy: 0.0572
Epoch 2/100
230/230 [=====] - 4s 20ms/step - loss: 3.6903 -
accuracy: 0.0694 - val_loss: 3.5403 - val_accuracy: 0.0925
Epoch 3/100
230/230 [=====] - 6s 28ms/step - loss: 3.4460 -
accuracy: 0.1055 - val_loss: 3.2752 - val_accuracy: 0.1285
Epoch 4/100
230/230 [=====] - 5s 20ms/step - loss: 3.2353 -
accuracy: 0.1286 - val_loss: 3.0684 - val_accuracy: 0.1671
Epoch 5/100
230/230 [=====] - 5s 20ms/step - loss: 3.0482 -
accuracy: 0.1516 - val_loss: 2.8333 - val_accuracy: 0.2047
Epoch 6/100
230/230 [=====] - 7s 29ms/step - loss: 2.8758 -
accuracy: 0.1838 - val_loss: 2.6965 - val_accuracy: 0.2384
Epoch 7/100
230/230 [=====] - 4s 19ms/step - loss: 2.7410 -
accuracy: 0.2053 - val_loss: 2.5757 - val_accuracy: 0.2515
Epoch 8/100
230/230 [=====] - 5s 20ms/step - loss: 2.6136 -
accuracy: 0.2291 - val_loss: 2.4723 - val_accuracy: 0.2711
Epoch 9/100
230/230 [=====] - 6s 27ms/step - loss: 2.5150 -
accuracy: 0.2584 - val_loss: 2.3961 - val_accuracy: 0.2874
Epoch 10/100
230/230 [=====] - 5s 20ms/step - loss: 2.4349 -
accuracy: 0.2784 - val_loss: 2.2997 - val_accuracy: 0.3261
Epoch 11/100
230/230 [=====] - 4s 19ms/step - loss: 2.3619 -
accuracy: 0.2991 - val_loss: 2.2400 - val_accuracy: 0.3381
Epoch 12/100
230/230 [=====] - 6s 27ms/step - loss: 2.3036 -
accuracy: 0.3166 - val_loss: 2.1751 - val_accuracy: 0.3756
Epoch 13/100
230/230 [=====] - 5s 21ms/step - loss: 2.2289 -
accuracy: 0.3357 - val_loss: 2.1284 - val_accuracy: 0.3827
Epoch 14/100
230/230 [=====] - 5s 22ms/step - loss: 2.1670 -
accuracy: 0.3493 - val_loss: 2.0505 - val_accuracy: 0.4017
```


Epoch 15/100
230/230 [=====] - 6s 27ms/step - loss: 2.1111 - accuracy: 0.3641 - val_loss: 1.9942 - val_accuracy: 0.4023
Epoch 16/100
230/230 [=====] - 5s 20ms/step - loss: 2.0460 - accuracy: 0.3814 - val_loss: 1.9737 - val_accuracy: 0.4088
Epoch 17/100
230/230 [=====] - 5s 21ms/step - loss: 1.9883 - accuracy: 0.4048 - val_loss: 1.9501 - val_accuracy: 0.4241
Epoch 18/100
230/230 [=====] - 6s 27ms/step - loss: 1.9483 - accuracy: 0.4097 - val_loss: 1.8844 - val_accuracy: 0.4230
Epoch 19/100
230/230 [=====] - 4s 19ms/step - loss: 1.8871 - accuracy: 0.4268 - val_loss: 1.7997 - val_accuracy: 0.4600
Epoch 20/100
230/230 [=====] - 5s 20ms/step - loss: 1.8451 - accuracy: 0.4345 - val_loss: 1.7730 - val_accuracy: 0.4823
Epoch 21/100
230/230 [=====] - 7s 28ms/step - loss: 1.7948 - accuracy: 0.4594 - val_loss: 1.7160 - val_accuracy: 0.4997
Epoch 22/100
230/230 [=====] - 5s 20ms/step - loss: 1.7520 - accuracy: 0.4713 - val_loss: 1.7165 - val_accuracy: 0.4829
Epoch 23/100
230/230 [=====] - 5s 21ms/step - loss: 1.7203 - accuracy: 0.4790 - val_loss: 1.6868 - val_accuracy: 0.5101
Epoch 24/100
230/230 [=====] - 6s 27ms/step - loss: 1.6815 - accuracy: 0.4814 - val_loss: 1.6795 - val_accuracy: 0.5095
Epoch 25/100
230/230 [=====] - 5s 21ms/step - loss: 1.6439 - accuracy: 0.5052 - val_loss: 1.5830 - val_accuracy: 0.5427
Epoch 26/100
230/230 [=====] - 5s 22ms/step - loss: 1.5916 - accuracy: 0.5203 - val_loss: 1.5895 - val_accuracy: 0.5406
Epoch 27/100
230/230 [=====] - 6s 27ms/step - loss: 1.5704 - accuracy: 0.5313 - val_loss: 1.5709 - val_accuracy: 0.5455
Epoch 28/100
230/230 [=====] - 4s 19ms/step - loss: 1.5318 - accuracy: 0.5433 - val_loss: 1.5212 - val_accuracy: 0.5623
Epoch 29/100
230/230 [=====] - 5s 22ms/step - loss: 1.5055 - accuracy: 0.5506 - val_loss: 1.5391 - val_accuracy: 0.5520
Epoch 30/100
230/230 [=====] - 6s 26ms/step - loss: 1.4706 - accuracy: 0.5658 - val_loss: 1.4595 - val_accuracy: 0.5890

Epoch 31/100
230/230 [=====] - 4s 20ms/step - loss: 1.4534 - accuracy: 0.5672 - val_loss: 1.4532 - val_accuracy: 0.5841

Epoch 32/100
230/230 [=====] - 5s 21ms/step - loss: 1.4122 - accuracy: 0.5824 - val_loss: 1.4450 - val_accuracy: 0.5934

Epoch 33/100
230/230 [=====] - 6s 27ms/step - loss: 1.3942 - accuracy: 0.5847 - val_loss: 1.3758 - val_accuracy: 0.6070

Epoch 34/100
230/230 [=====] - 5s 20ms/step - loss: 1.3746 - accuracy: 0.5922 - val_loss: 1.3792 - val_accuracy: 0.6026

Epoch 35/100
230/230 [=====] - 5s 22ms/step - loss: 1.3337 - accuracy: 0.6116 - val_loss: 1.3454 - val_accuracy: 0.6179

Epoch 36/100
230/230 [=====] - 6s 27ms/step - loss: 1.3192 - accuracy: 0.6063 - val_loss: 1.3635 - val_accuracy: 0.6189

Epoch 37/100
230/230 [=====] - 7s 29ms/step - loss: 1.3119 - accuracy: 0.6096 - val_loss: 1.3324 - val_accuracy: 0.6086

Epoch 38/100
230/230 [=====] - 6s 25ms/step - loss: 1.2998 - accuracy: 0.6112 - val_loss: 1.3223 - val_accuracy: 0.6157

Epoch 39/100
230/230 [=====] - 5s 23ms/step - loss: 1.2733 - accuracy: 0.6231 - val_loss: 1.2850 - val_accuracy: 0.6375

Epoch 40/100
230/230 [=====] - 5s 21ms/step - loss: 1.2412 - accuracy: 0.6349 - val_loss: 1.2724 - val_accuracy: 0.6375

Epoch 41/100
230/230 [=====] - 6s 26ms/step - loss: 1.2065 - accuracy: 0.6446 - val_loss: 1.3188 - val_accuracy: 0.6364

Epoch 42/100
230/230 [=====] - 5s 24ms/step - loss: 1.2054 - accuracy: 0.6418 - val_loss: 1.2413 - val_accuracy: 0.6505

Epoch 43/100
230/230 [=====] - 4s 19ms/step - loss: 1.1724 - accuracy: 0.6518 - val_loss: 1.2153 - val_accuracy: 0.6576

Epoch 44/100
230/230 [=====] - 7s 31ms/step - loss: 1.1606 - accuracy: 0.6593 - val_loss: 1.2107 - val_accuracy: 0.6576

Epoch 45/100
230/230 [=====] - 5s 20ms/step - loss: 1.1813 - accuracy: 0.6476 - val_loss: 1.1999 - val_accuracy: 0.6609

Epoch 46/100
230/230 [=====] - 5s 21ms/step - loss: 1.1356 - accuracy: 0.6665 - val_loss: 1.1637 - val_accuracy: 0.6761

```


Epoch 47/100
230/230 [=====] - 6s 27ms/step - loss: 1.1122 -
accuracy: 0.6711 - val_loss: 1.1611 - val_accuracy: 0.6745
Epoch 48/100
230/230 [=====] - 5s 21ms/step - loss: 1.1101 -
accuracy: 0.6725 - val_loss: 1.1394 - val_accuracy: 0.6897
Epoch 49/100
230/230 [=====] - 5s 21ms/step - loss: 1.0707 -
accuracy: 0.6817 - val_loss: 1.1129 - val_accuracy: 0.6821
Epoch 50/100
230/230 [=====] - 6s 27ms/step - loss: 1.0784 -
accuracy: 0.6796 - val_loss: 1.1557 - val_accuracy: 0.6734
Epoch 51/100
230/230 [=====] - 5s 21ms/step - loss: 1.0550 -
accuracy: 0.6904 - val_loss: 1.1224 - val_accuracy: 0.6886
Epoch 52/100
230/230 [=====] - 4s 19ms/step - loss: 1.0537 -
accuracy: 0.6918 - val_loss: 1.0932 - val_accuracy: 0.6930
Epoch 53/100
230/230 [=====] - 7s 30ms/step - loss: 1.0294 -
accuracy: 0.7005 - val_loss: 1.0722 - val_accuracy: 0.7001
Epoch 54/100
230/230 [=====] - 5s 20ms/step - loss: 1.0248 -
accuracy: 0.6956 - val_loss: 1.0894 - val_accuracy: 0.6935
Epoch 55/100
230/230 [=====] - 5s 21ms/step - loss: 1.0138 -
accuracy: 0.7062 - val_loss: 1.0988 - val_accuracy: 0.6875
Epoch 56/100
230/230 [=====] - 7s 29ms/step - loss: 0.9829 -
accuracy: 0.7115 - val_loss: 1.0786 - val_accuracy: 0.6941
Training time: 301.18429708480835 seconds

```


```

[ ]: # Plot the loss
plt.title('Loss Plot')
plt.plot(history.history['loss'], label='train')
plt.plot(history.history['val_loss'], label='validation')
plt.legend()
plt.show()

```



```
[ ]: # Plot the accuracy
plt.title('Accuracy Plot')
plt.plot(history.history['accuracy'], label='train')
plt.plot(history.history['val_accuracy'], label='validation')
plt.legend()
plt.show()
```



```
[ ]: # Load the saved model
saved_model = tf.keras.models.load_model(model_checkpoint_path)

# Evaluate the model with the test set
loss, accuracy = saved_model.evaluate(X_test_padded, y_test_array)

print("Test Loss:", loss)
print("Test Accuracy:", round((accuracy*100), 2))
```

```
97/97 [=====] - 2s 6ms/step - loss: 1.0427 - accuracy:
0.7166
Test Loss: 1.042669415473938
Test Accuracy: 71.66
```

```
[ ]: # Check predictions with the test set
y_test_prob = saved_model.predict(X_test_padded)

# Convert probabilities to class labels
y_test_pred = np.argmax(y_test_prob, axis=1)

# Calculate precision, recall, and f1 score
```

```

precision = precision_score(y_test_array, y_test_pred, average='weighted')
recall = recall_score(y_test_array, y_test_pred, average='weighted')
f1 = f1_score(y_test_array, y_test_pred, average='weighted')

print("Precision:", round((precision*100), 2))
print("Recall:", round((recall*100), 2))
print("F1 Score:", round((f1*100), 2))

```

```

97/97 [=====] - 1s 7ms/step
Precision: 72.08
Recall: 71.66
F1 Score: 70.69

```

```

[ ]: # Error analysis

# Print classification report
print(classification_report(y_test_array, y_test_pred))

# Check misclassified data
misclassified_data = np.where(y_test_pred != y_test_array)[0]
print(f"The number of misclassifications: {len(misclassified_data)}")

# Check the ratio of misclassifications
misclassification_ratio = (len(misclassified_data) / len(y_test_array)) * 100
# Round the number
rounded_ratio = round(misclassification_ratio, 2)
print(f"Proportion of misclassifications: {rounded_ratio}%")

# Iterate over misclassified data for error analysis
for idx in misclassified_data[:30]:
    input_text = X_test[idx]
    true_label = y_test[idx]
    predicted_label = y_test_pred[idx]

    # Print information about the misclassified data
    print("Input Text:", input_text)
    print("Actual Label:", true_label)
    print("Predicted Label:", predicted_label)
    print()

```

	precision	recall	f1-score	support
0	0.92	0.90	0.91	40
1	0.77	0.85	0.81	40
2	0.86	0.95	0.90	40
3	0.82	0.70	0.76	40
4	0.66	0.68	0.67	40
5	0.59	0.75	0.66	40

6	0.85	0.85	0.85	40
7	0.79	0.47	0.59	40
8	0.77	0.82	0.80	40
9	1.00	0.93	0.96	40
10	0.00	0.00	0.00	40
11	0.41	0.82	0.55	40
12	0.58	0.35	0.44	40
13	0.80	0.88	0.83	40
14	0.66	0.68	0.67	40
15	0.74	0.78	0.76	40
16	0.52	0.62	0.57	40
17	0.94	0.82	0.88	40
18	0.86	0.78	0.82	40
19	0.88	0.90	0.89	40
20	0.61	0.75	0.67	40
21	0.80	0.80	0.80	40
22	0.61	0.50	0.55	40
23	1.00	0.88	0.93	40
24	0.79	0.95	0.86	40
25	0.54	0.68	0.60	40
26	0.65	0.85	0.74	40
27	0.93	0.68	0.78	40
28	0.83	0.72	0.77	40
29	0.76	0.70	0.73	40
30	0.83	0.97	0.90	40
31	0.81	0.85	0.83	40
32	0.85	0.97	0.91	40
33	0.84	0.80	0.82	40
34	0.70	0.82	0.76	40
35	0.59	0.68	0.63	40
36	0.77	0.82	0.80	40
37	0.49	0.68	0.57	40
38	0.86	0.90	0.88	40
39	0.47	0.60	0.53	40
40	0.65	0.93	0.76	40
41	0.63	0.68	0.65	40
42	0.91	0.80	0.85	40
43	0.33	0.38	0.35	40
44	0.84	0.90	0.87	40
45	0.94	0.72	0.82	40
46	0.88	0.75	0.81	40
47	0.70	0.65	0.68	40
48	0.64	0.53	0.58	40
49	0.84	0.65	0.73	40
50	0.74	0.70	0.72	40
51	0.69	0.88	0.77	40
52	0.70	0.47	0.57	40
53	0.54	0.72	0.62	40

	54	0.60	0.60	0.60	40
	55	0.81	0.88	0.84	40
	56	0.94	0.40	0.56	40
	57	0.85	0.88	0.86	40
	58	0.77	0.75	0.76	40
	59	0.57	0.65	0.60	40
	60	0.71	0.80	0.75	40
	61	0.57	0.72	0.64	40
	62	0.75	0.68	0.71	40
	63	0.97	0.78	0.86	40
	64	0.73	0.88	0.80	40
	65	0.61	0.70	0.65	40
	66	0.56	0.35	0.43	40
	67	0.61	0.82	0.70	40
	68	0.79	0.55	0.65	40
	69	0.36	0.23	0.28	40
	70	0.72	0.97	0.83	40
	71	0.83	0.95	0.88	40
	72	0.50	0.03	0.05	40
	73	0.97	0.85	0.91	40
	74	0.43	0.57	0.49	40
	75	0.78	0.53	0.63	40
	76	0.91	0.75	0.82	40
	accuracy			0.72	3080
	macro avg	0.72	0.72	0.71	3080
	weighted avg	0.72	0.72	0.71	3080

The number of misclassifications: 873

Proportion of misclassifications: 28.34%

Input Text: locate card

Actual Label: 11

Predicted Label: 0

Input Text: waiting longer expected bank card could provide information arrive

Actual Label: 11

Predicted Label: 12

Input Text: ive waiting longer expected card

Actual Label: 11

Predicted Label: 12

Input Text: get card yet lost

Actual Label: 11

Predicted Label: 41

Input Text: status card ordered

Actual Label: 11

Predicted Label: 14

Input Text: long new card take arrive

Actual Label: 11

Predicted Label: 12

Input Text: tracking info available

Actual Label: 11

Predicted Label: 30

Input Text: add card account

Actual Label: 13

Predicted Label: 39

Input Text: put old card back system found

Actual Label: 13

Predicted Label: 41

Input Text: hello found card misplaced need reactive

Actual Label: 13

Predicted Label: 42

Input Text: view card received app

Actual Label: 13

Predicted Label: 41

Input Text: app doesnt show card received

Actual Label: 13

Predicted Label: 41

Input Text: good time exchange

Actual Label: 32

Predicted Label: 31

Input Text: made currency exchange think charged

Actual Label: 17

Predicted Label: 31

Input Text: rate low sure using right exchange rate

Actual Label: 17

Predicted Label: 32

Input Text: charged

Actual Label: 17

Predicted Label: 34

Input Text: charged wrong currency exchange purchase

Actual Label: 17

Predicted Label: 31

Input Text: exchange rate seems transaction

Actual Label: 17

Predicted Label: 32

Input Text: conversion value card payments incorrect

Actual Label: 17

Predicted Label: 16

Input Text: check exchange rate applied transaction

Actual Label: 17

Predicted Label: 32

Input Text: would like refund extra pound charged

Actual Label: 34

Predicted Label: 19

Input Text: transaction credited

Actual Label: 34

Predicted Label: 66

Input Text: fee come

Actual Label: 34

Predicted Label: 15

Input Text: extra pound charge card

Actual Label: 34

Predicted Label: 57

Input Text: euro fee come

Actual Label: 34

Predicted Label: 50

Input Text: euro fee statement

Actual Label: 34

Predicted Label: 31

Input Text: two weeks transaction reversed

Actual Label: 34

Predicted Label: 8

Input Text: withdrawal still pending

Actual Label: 46

Predicted Label: 66

Input Text: hi wondering help used city centre atm get cash machine declined
card account shows transaction still pending didnt receive money please cancel

transaction
Actual Label: 46
Predicted Label: 76

Input Text: long til cash goes
Actual Label: 46
Predicted Label: 26

1.6.4 Hyperparameter tuning

```
[ ]: # The code for hyperparameter tuning is derived from the Tensorflow website.  
# (https://www.tensorflow.org/tutorials/keras/keras_tuner)  
  
# Define the model for hyperparameter tuning  
def model_builder(hp):  
    model = keras.models.Sequential()  
    e = layers.Embedding(voca_size+1, 100, weights=[embedding_matrix],  
↳input_length=max_length_train_text, trainable=False)  
    model.add(e)  
    hp_units = hp.Int('units', min_value = 20, max_value = 50, step = 10) # Set  
↳up the hyperparameters  
    model.add(layers.LSTM(units = hp_units)) # We will check the optimal hidden  
↳unit for the LSTM layer  
    model.add(layers.Dense(nlabel, activation='softmax'))  
  
    hp_learning_rate = hp.Choice('learning_rate', values = [0.01, 0.001, 0.0001])  
↳# Set up the hyperparameters  
    model.compile(optimizer = keras.optimizers.Adam(learning_rate =  
↳hp_learning_rate), # We will check the optimal learning rate  
                  loss = 'sparse_categorical_crossentropy',  
                  metrics = ['accuracy'])  
    return model
```

```
[ ]: # The code for hyperparameter tuning is derived from the Tensorflow website.  
# (https://www.tensorflow.org/tutorials/keras/keras_tuner)  
  
# Specify the tuner  
tuner = kt.Hyperband(model_builder,  
                     objective = 'val_accuracy',  
                     max_epochs = 100)
```

```
[ ]: # The code for hyperparameter tuning is derived from the Tensorflow website.  
# (https://www.tensorflow.org/tutorials/keras/keras_tuner)  
  
# Set up a callback for early stopping  
stop_early = tf.keras.callbacks.EarlyStopping(monitor='val_loss', patience=5)
```

```
[ ]: # The code for hyperparameter tuning is derived from the Tensorflow website.
# (https://www.tensorflow.org/tutorials/keras/keras_tuner)

# Run the tuner
tuner.search(X_train_padded, y_train, epochs = 100, validation_data = (X_val_padded, y_val), callbacks = [stop_early])

# Get the optimal hyperparameters
best_hps = tuner.get_best_hyperparameters(num_trials = 1)[0]

print(f"The optimal number of units: {best_hps.get('units')}. The optimal learning rate: {best_hps.get('learning_rate')}")
```

Trial 12 Complete [00h 00m 13s]
val_accuracy: 0.02994011901319027

Best val_accuracy So Far: 0.2449646145105362
Total elapsed time: 00h 09m 37s
The optimal number of units: 50. The optimal learning rate: 0.01.

1.6.5 Tuned LSTM

```
[ ]: # Define the output dimension for the embedding layer and hidden units
nlabel = 77

tuned_model = keras.models.Sequential()
e = layers.Embedding(voca_size+1, 100, weights=[embedding_matrix], input_length=max_length_train_text, trainable=False) # Using 100 dimension for GloVe
tuned_model.add(e)
tuned_model.add(layers.LSTM(50))
tuned_model.add(layers.Dense(nlabel, activation='softmax'))

# Compile the model
tuned_model.compile(optimizer=keras.optimizers.Adam(learning_rate = 0.01), loss='sparse_categorical_crossentropy', metrics=['accuracy']) # run_eagerly=True

# Summary the model
tuned_model.summary()
```

Model: "sequential_12"

Layer (type)	Output Shape	Param #
embedding_12 (Embedding)	(None, 29, 100)	208900

lstm_12 (LSTM)	(None, 50)	30200
dense_12 (Dense)	(None, 77)	3927

```
=====
Total params: 243027 (949.32 KB)
Trainable params: 34127 (133.31 KB)
Non-trainable params: 208900 (816.02 KB)
-----
```

```
[ ]: # Define the folder path to save the model
folder_path = '/content/drive/MyDrive/1. NLP CW/LSTM/'

# Define the file path for the model checkpoint
model_checkpoint_path = folder_path + 'tuned_LSTM_glove_model.keras'

# Define the model checkpoint
mc = tf.keras.callbacks.ModelCheckpoint(
    filepath=model_checkpoint_path,
    monitor='val_accuracy',
    mode='max',
    save_best_only=True)
```

```
[ ]: # Import time to measure the elapsed time
import time

# Measure time before training
start_time = time.time()

# Fit the model
history = tuned_model.fit(
    X_train_padded, y_train,
    epochs = 100,
    validation_data = (X_val_padded, y_val),
    callbacks = [mc, es],
    batch_size = 32)

# End the training time
end_time = time.time()

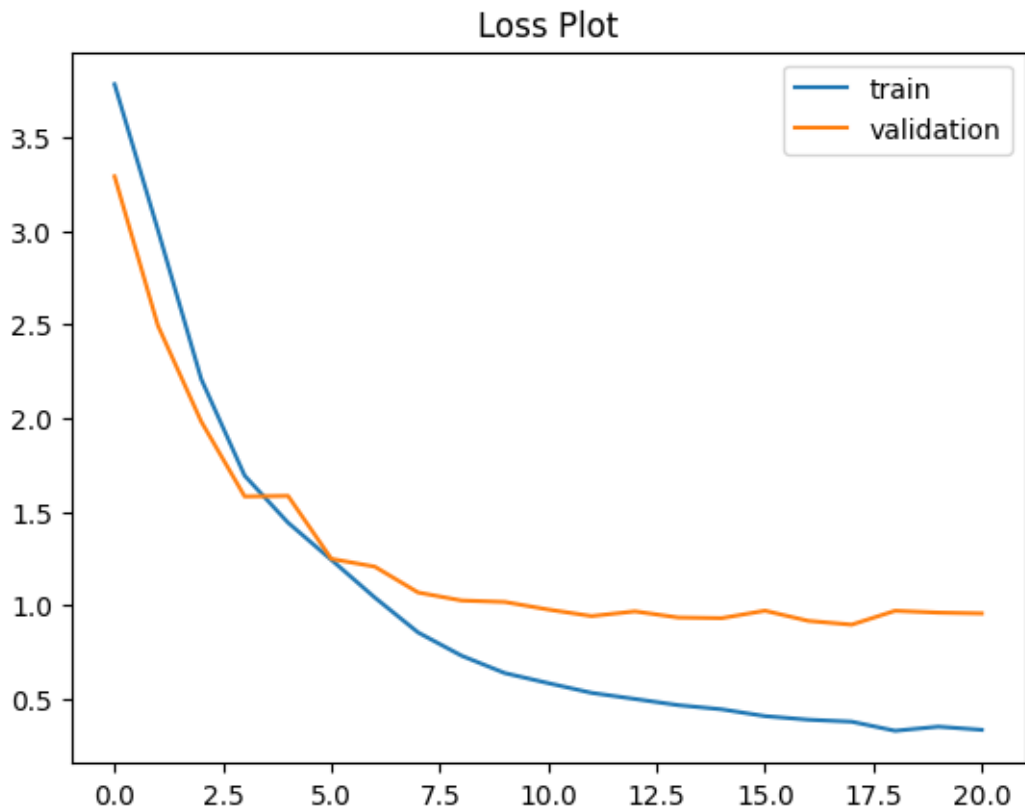
# Measure the training time
training_time = end_time - start_time
print("Training time:", training_time, "seconds")
```

```
Epoch 1/100
230/230 [=====] - 10s 34ms/step - loss: 3.7835 -
accuracy: 0.0475 - val_loss: 3.2911 - val_accuracy: 0.1083
Epoch 2/100
```


230/230 [=====] - 5s 22ms/step - loss: 3.0044 - accuracy: 0.1673 - val_loss: 2.4936 - val_accuracy: 0.2450
Epoch 3/100
230/230 [=====] - 5s 21ms/step - loss: 2.2075 - accuracy: 0.3361 - val_loss: 1.9810 - val_accuracy: 0.3865
Epoch 4/100
230/230 [=====] - 8s 33ms/step - loss: 1.6925 - accuracy: 0.4758 - val_loss: 1.5815 - val_accuracy: 0.5248
Epoch 5/100
230/230 [=====] - 5s 20ms/step - loss: 1.4416 - accuracy: 0.5687 - val_loss: 1.5846 - val_accuracy: 0.5416
Epoch 6/100
230/230 [=====] - 5s 21ms/step - loss: 1.2428 - accuracy: 0.6352 - val_loss: 1.2479 - val_accuracy: 0.6375
Epoch 7/100
230/230 [=====] - 7s 30ms/step - loss: 1.0417 - accuracy: 0.6911 - val_loss: 1.2065 - val_accuracy: 0.6489
Epoch 8/100
230/230 [=====] - 5s 21ms/step - loss: 0.8548 - accuracy: 0.7442 - val_loss: 1.0684 - val_accuracy: 0.7028
Epoch 9/100
230/230 [=====] - 5s 24ms/step - loss: 0.7309 - accuracy: 0.7818 - val_loss: 1.0257 - val_accuracy: 0.7158
Epoch 10/100
230/230 [=====] - 6s 27ms/step - loss: 0.6378 - accuracy: 0.8091 - val_loss: 1.0178 - val_accuracy: 0.7153
Epoch 11/100
230/230 [=====] - 5s 20ms/step - loss: 0.5842 - accuracy: 0.8267 - val_loss: 0.9773 - val_accuracy: 0.7311
Epoch 12/100
230/230 [=====] - 6s 27ms/step - loss: 0.5322 - accuracy: 0.8407 - val_loss: 0.9424 - val_accuracy: 0.7365
Epoch 13/100
230/230 [=====] - 6s 25ms/step - loss: 0.4995 - accuracy: 0.8467 - val_loss: 0.9674 - val_accuracy: 0.7436
Epoch 14/100
230/230 [=====] - 6s 28ms/step - loss: 0.4665 - accuracy: 0.8565 - val_loss: 0.9347 - val_accuracy: 0.7539
Epoch 15/100
230/230 [=====] - 6s 27ms/step - loss: 0.4443 - accuracy: 0.8665 - val_loss: 0.9314 - val_accuracy: 0.7507
Epoch 16/100
230/230 [=====] - 5s 21ms/step - loss: 0.4079 - accuracy: 0.8756 - val_loss: 0.9710 - val_accuracy: 0.7523
Epoch 17/100
230/230 [=====] - 6s 26ms/step - loss: 0.3888 - accuracy: 0.8801 - val_loss: 0.9167 - val_accuracy: 0.7588
Epoch 18/100

```
230/230 [=====] - 6s 24ms/step - loss: 0.3782 -  
accuracy: 0.8866 - val_loss: 0.8966 - val_accuracy: 0.7637  
Epoch 19/100  
230/230 [=====] - 5s 22ms/step - loss: 0.3296 -  
accuracy: 0.8980 - val_loss: 0.9706 - val_accuracy: 0.7539  
Epoch 20/100  
230/230 [=====] - 6s 28ms/step - loss: 0.3516 -  
accuracy: 0.8885 - val_loss: 0.9608 - val_accuracy: 0.7561  
Epoch 21/100  
230/230 [=====] - 5s 23ms/step - loss: 0.3347 -  
accuracy: 0.8945 - val_loss: 0.9567 - val_accuracy: 0.7512  
Training time: 123.43249177932739 seconds
```

```
[ ]: # Plot the loss  
plt.title('Loss Plot')  
plt.plot(history.history['loss'], label='train')  
plt.plot(history.history['val_loss'], label='validation')  
plt.legend()  
plt.show()
```



```
[ ]: # Plot the accuracy
plt.title('Accuracy Plot')
plt.plot(history.history['accuracy'], label='train')
plt.plot(history.history['val_accuracy'], label='validation')
plt.legend()
plt.show()
```



```
[ ]: # Load the saved model
saved_model = tf.keras.models.load_model(model_checkpoint_path)

# Evaluate the model with the test set
loss, accuracy = saved_model.evaluate(X_test_padded, y_test_array)

print("Test Loss:", loss)
print("Test Accuracy:", round((accuracy*100), 2))
```

```
97/97 [=====] - 2s 7ms/step - loss: 0.8815 - accuracy:
0.7769
Test Loss: 0.8814808130264282
Test Accuracy: 77.69
```



```
[ ]: # Check predictions with the test set
y_test_prob = saved_model.predict(X_test_padded)

# Convert probabilities to class labels
y_test_pred = np.argmax(y_test_prob, axis=1)

# Calculate precision, recall, and f1 score
precision = precision_score(y_test_array, y_test_pred, average='weighted')
recall = recall_score(y_test_array, y_test_pred, average='weighted')
f1 = f1_score(y_test_array, y_test_pred, average='weighted')

print("Precision:", round((precision*100), 2))
print("Recall:", round((recall*100), 2))
print("F1 Score:", round((f1*100), 2))
```

```
97/97 [=====] - 1s 7ms/step
Precision: 79.27
Recall: 77.69
F1 Score: 77.84
```

```
[ ]: # Error analysis

# Print classification report
print(classification_report(y_test_array, y_test_pred))

# Check misclassified data
misclassified_data = np.where(y_test_pred != y_test_array)[0]
print(f"The number of misclassifications: {len(misclassified_data)}")

# Check the ratio of misclassifications
misclassification_ratio = (len(misclassified_data) / len(y_test_array)) * 100
# Round the number
rounded_ratio = round(misclassification_ratio, 2)
print(f"Proportion of misclassifications: {rounded_ratio}%")

# Iterate over misclassified data for error analysis
for idx in misclassified_data[:30]:
    input_text = X_test[idx]
    true_label = y_test[idx]
    predicted_label = y_test_pred[idx]

    # Print information about the misclassified data
    print("Input Text:", input_text)
    print("Actual Label:", true_label)
    print("Predicted Label:", predicted_label)
    print()
```

```
precision    recall  f1-score   support
```

0	1.00	0.93	0.96	40
1	0.92	0.88	0.90	40
2	0.95	0.90	0.92	40
3	0.95	0.88	0.91	40
4	0.91	0.78	0.84	40
5	0.56	0.68	0.61	40
6	0.84	0.80	0.82	40
7	0.63	0.55	0.59	40
8	0.85	0.88	0.86	40
9	0.95	0.97	0.96	40
10	0.76	0.70	0.73	40
11	0.54	0.88	0.67	40
12	0.69	0.50	0.58	40
13	0.81	0.95	0.87	40
14	0.63	0.78	0.70	40
15	0.79	0.82	0.80	40
16	0.55	0.68	0.61	40
17	0.84	0.93	0.88	40
18	1.00	0.65	0.79	40
19	0.83	0.85	0.84	40
20	0.71	0.72	0.72	40
21	0.97	0.80	0.88	40
22	0.86	0.60	0.71	40
23	0.95	0.88	0.91	40
24	0.90	0.90	0.90	40
25	0.66	0.82	0.73	40
26	0.68	0.70	0.69	40
27	0.93	0.68	0.78	40
28	0.71	0.75	0.73	40
29	0.93	0.65	0.76	40
30	1.00	0.95	0.97	40
31	0.91	0.80	0.85	40
32	0.93	0.95	0.94	40
33	0.71	0.93	0.80	40
34	0.77	0.82	0.80	40
35	0.57	0.75	0.65	40
36	0.89	0.82	0.86	40
37	0.65	0.75	0.70	40
38	0.81	0.97	0.89	40
39	0.70	0.70	0.70	40
40	0.79	0.95	0.86	40
41	0.87	0.65	0.74	40
42	0.97	0.85	0.91	40
43	0.53	0.65	0.58	40
44	0.93	0.95	0.94	40
45	0.85	0.70	0.77	40
46	0.84	0.78	0.81	40

	47	0.63	0.80	0.70	40
	48	0.73	0.55	0.63	40
	49	0.94	0.80	0.86	40
	50	0.77	0.68	0.72	40
	51	0.74	0.88	0.80	40
	52	0.84	0.78	0.81	40
	53	0.79	0.78	0.78	40
	54	0.71	0.60	0.65	40
	55	0.92	0.88	0.90	40
	56	0.83	0.62	0.71	40
	57	0.89	0.82	0.86	40
	58	0.96	0.65	0.78	40
	59	0.79	0.78	0.78	40
	60	0.79	0.82	0.80	40
	61	0.84	0.68	0.75	40
	62	0.70	0.75	0.72	40
	63	0.81	0.85	0.83	40
	64	0.72	0.85	0.78	40
	65	0.67	0.65	0.66	40
	66	0.55	0.72	0.62	40
	67	0.69	0.62	0.66	40
	68	0.75	0.75	0.75	40
	69	0.56	0.38	0.45	40
	70	0.89	0.97	0.93	40
	71	0.93	0.93	0.93	40
	72	0.85	0.72	0.78	40
	73	0.92	0.90	0.91	40
	74	0.48	0.72	0.57	40
	75	0.55	0.75	0.63	40
	76	0.78	0.72	0.75	40
	accuracy			0.78	3080
	macro avg	0.79	0.78	0.78	3080
	weighted avg	0.79	0.78	0.78	3080

The number of misclassifications: 687

Proportion of misclassifications: 22.31%

Input Text: locate card

Actual Label: 11

Predicted Label: 13

Input Text: way know card arrive

Actual Label: 11

Predicted Label: 12

Input Text: get card

Actual Label: 11

Predicted Label: 43

Input Text: long card delivery take

Actual Label: 11

Predicted Label: 12

Input Text: get card yet lost

Actual Label: 11

Predicted Label: 41

Input Text: add card account

Actual Label: 13

Predicted Label: 39

Input Text: put old card back system found

Actual Label: 13

Predicted Label: 11

Input Text: good time exchange

Actual Label: 32

Predicted Label: 33

Input Text: currencies exchange rate calculated

Actual Label: 32

Predicted Label: 17

Input Text: charged

Actual Label: 17

Predicted Label: 34

Input Text: conversion value card payments incorrect

Actual Label: 17

Predicted Label: 16

Input Text: paid something foreign currency noticed exchange rate incorrect

Actual Label: 17

Predicted Label: 76

Input Text: explain random charge

Actual Label: 34

Predicted Label: 72

Input Text: transaction credited

Actual Label: 34

Predicted Label: 47

Input Text: fee come

Actual Label: 34

Predicted Label: 15

Input Text: extra charge

Actual Label: 34

Predicted Label: 64

Input Text: euro fee come

Actual Label: 34

Predicted Label: 31

Input Text: new customer happened look app charge familiar could tell extra charge

Actual Label: 34

Predicted Label: 28

Input Text: two weeks transaction reversed

Actual Label: 34

Predicted Label: 27

Input Text: hey tried get money earlier machine didnt work saw transaction still seems progress please check whats going seems something broken dont want charged money havent actually received

Actual Label: 46

Predicted Label: 63

Input Text: long til cash goes

Actual Label: 46

Predicted Label: 5

Input Text: hii tried get money machine working transaction still seems progress please check whats going oni dont want charged money received

Actual Label: 46

Predicted Label: 63

Input Text: made withdrawal account posted

Actual Label: 46

Predicted Label: 6

Input Text: tried take money card didnt work later saw transaction still progress whats goign

Actual Label: 46

Predicted Label: 25

Input Text: wheres accounting cash withdrawal

Actual Label: 46

Predicted Label: 76

Input Text: long take post atm drawl

Actual Label: 46

Predicted Label: 4

Input Text: cash withdrawal atm still yet showing confirmed account

Actual Label: 46

Predicted Label: 20

Input Text: whats pending transaction card declined atm account says still
pending cancel payment

Actual Label: 46

Predicted Label: 51

Input Text: incoming payment account deactivated still processed

Actual Label: 36

Predicted Label: 45

Input Text: currencies exchanges

Actual Label: 36

Predicted Label: 32

1.6.6 Tuned LSTM (with dropout)

```
[ ]: # Define the output dimension for the embedding layer and hidden units
nlabel = 77

dropout_tuned_model = keras.models.Sequential()
e = layers.Embedding(voca_size+1, 100, weights=[embedding_matrix],
    ↳input_length=max_length_train_text, trainable=False) # Using 100 dimension
    ↳for GloVe
dropout_tuned_model.add(e)
dropout_tuned_model.add(layers.LSTM(50, dropout=0.2))
dropout_tuned_model.add(layers.Dense(nlabel, activation='softmax'))

# Compile the model
dropout_tuned_model.compile(optimizer=keras.optimizers.Adam(learning_rate = 0.
    ↳01), loss='sparse_categorical_crossentropy', metrics=['accuracy']) #,
    ↳run_eagerly=True

# Summary the model
dropout_tuned_model.summary()
```

Model: "sequential_13"

Layer (type)	Output Shape	Param #
embedding_13 (Embedding)	(None, 29, 100)	208900

lstm_13 (LSTM)	(None, 50)	30200
dense_13 (Dense)	(None, 77)	3927

```
=====
Total params: 243027 (949.32 KB)
Trainable params: 34127 (133.31 KB)
Non-trainable params: 208900 (816.02 KB)
-----
```

```
[ ]: # Define the folder path to save the model
folder_path = '/content/drive/MyDrive/1. NLP CW/LSTM/'

# Define the file path for the model checkpoint
model_checkpoint_path = folder_path + 'dropout_tuned_LSTM_glove_model.keras'

# Define the model checkpoint
mc = tf.keras.callbacks.ModelCheckpoint(
    filepath=model_checkpoint_path,
    monitor='val_accuracy',
    mode='max',
    save_best_only=True)
```

```
[ ]: # Import time to measure the elapsed time
import time

# Measure time before training
start_time = time.time()

# Fit the model
history = dropout_tuned_model.fit(
    X_train_padded, y_train,
    epochs = 100,
    validation_data = (X_val_padded, y_val),
    callbacks = [mc, es],
    batch_size = 32)

# End the training time
end_time = time.time()


# Measure the training time
training_time = end_time - start_time
print("Training time:", training_time, "seconds")
```

```
Epoch 1/100
230/230 [=====] - 11s 32ms/step - loss: 3.9919 -
accuracy: 0.0328 - val_loss: 3.5665 - val_accuracy: 0.0675
Epoch 2/100
```

230/230 [=====] - 5s 23ms/step - loss: 3.1440 - accuracy: 0.1525 - val_loss: 2.5732 - val_accuracy: 0.2580
Epoch 3/100
230/230 [=====] - 7s 32ms/step - loss: 2.2885 - accuracy: 0.3294 - val_loss: 1.9406 - val_accuracy: 0.4066
Epoch 4/100
230/230 [=====] - 5s 23ms/step - loss: 1.8043 - accuracy: 0.4413 - val_loss: 1.5322 - val_accuracy: 0.5291
Epoch 5/100
230/230 [=====] - 6s 25ms/step - loss: 1.5087 - accuracy: 0.5449 - val_loss: 1.3131 - val_accuracy: 0.6091
Epoch 6/100
230/230 [=====] - 6s 26ms/step - loss: 1.2901 - accuracy: 0.6024 - val_loss: 1.2017 - val_accuracy: 0.6358
Epoch 7/100
230/230 [=====] - 5s 22ms/step - loss: 1.1630 - accuracy: 0.6401 - val_loss: 1.0661 - val_accuracy: 0.6924
Epoch 8/100
230/230 [=====] - 7s 30ms/step - loss: 1.0470 - accuracy: 0.6771 - val_loss: 1.0115 - val_accuracy: 0.6962
Epoch 9/100
230/230 [=====] - 5s 23ms/step - loss: 1.0137 - accuracy: 0.6917 - val_loss: 1.1810 - val_accuracy: 0.6538
Epoch 10/100
230/230 [=====] - 5s 22ms/step - loss: 1.0596 - accuracy: 0.6762 - val_loss: 0.9902 - val_accuracy: 0.7039
Epoch 11/100
230/230 [=====] - 7s 32ms/step - loss: 0.9101 - accuracy: 0.7207 - val_loss: 0.9309 - val_accuracy: 0.7267
Epoch 12/100
230/230 [=====] - 5s 23ms/step - loss: 0.8195 - accuracy: 0.7437 - val_loss: 0.8921 - val_accuracy: 0.7414
Epoch 13/100
230/230 [=====] - 6s 28ms/step - loss: 0.7892 - accuracy: 0.7522 - val_loss: 0.8647 - val_accuracy: 0.7420
Epoch 14/100
230/230 [=====] - 6s 27ms/step - loss: 0.7638 - accuracy: 0.7601 - val_loss: 0.8502 - val_accuracy: 0.7583
Epoch 15/100
230/230 [=====] - 5s 21ms/step - loss: 0.7349 - accuracy: 0.7682 - val_loss: 0.8390 - val_accuracy: 0.7605
Epoch 16/100
230/230 [=====] - 7s 32ms/step - loss: 0.6979 - accuracy: 0.7812 - val_loss: 0.8738 - val_accuracy: 0.7572
Epoch 17/100
230/230 [=====] - 5s 23ms/step - loss: 0.6964 - accuracy: 0.7830 - val_loss: 0.8131 - val_accuracy: 0.7681
Epoch 18/100



```
230/230 [=====] - 6s 25ms/step - loss: 0.6616 -  
accuracy: 0.7919 - val_loss: 0.8395 - val_accuracy: 0.7588  
Epoch 19/100  
230/230 [=====] - 7s 31ms/step - loss: 0.6132 -  
accuracy: 0.8117 - val_loss: 0.8524 - val_accuracy: 0.7686  
Epoch 20/100  
230/230 [=====] - 5s 22ms/step - loss: 0.6627 -  
accuracy: 0.7859 - val_loss: 0.8202 - val_accuracy: 0.7567  
Training time: 123.8074312210083 seconds
```

```
[ ]: # Plot the loss  
plt.title('Loss Plot')  
plt.plot(history.history['loss'], label='train')  
plt.plot(history.history['val_loss'], label='validation')  
plt.legend()  
plt.show()
```



```
[ ]: # Plot the accuracy  
plt.title('Accuracy Plot')  
plt.plot(history.history['accuracy'], label='train')  
plt.plot(history.history['val_accuracy'], label='validation')
```

```
plt.legend()
plt.show()
```



```
[ ]: # Load the saved model
saved_model = tf.keras.models.load_model(model_checkpoint_path)

# Evaluate the model with the test set
loss, accuracy = saved_model.evaluate(X_test_padded, y_test_array)

print("Test Loss:", loss)
print("Test Accuracy:", round((accuracy*100), 2))
```

```
97/97 [=====] - 2s 8ms/step - loss: 0.7693 - accuracy:
0.7786
Test Loss: 0.7692508697509766
Test Accuracy: 77.86
```

```
[ ]: # Check predictions with the test set
y_test_prob = saved_model.predict(X_test_padded)

# Convert probabilities to class labels
```

```

y_test_pred = np.argmax(y_test_prob, axis=1)

# Calculate precision, recall, and f1 score
precision = precision_score(y_test_array, y_test_pred, average='weighted')
recall = recall_score(y_test_array, y_test_pred, average='weighted')
f1 = f1_score(y_test_array, y_test_pred, average='weighted')

print("Precision:", round((precision*100), 2))
print("Recall:", round((recall*100), 2))
print("F1 Score:", round((f1*100), 2))

```

```

97/97 [=====] - 2s 11ms/step
Precision: 79.88
Recall: 77.86
F1 Score: 77.85

```

```

[ ]: # Error analysis

# Print classification report
print(classification_report(y_test_array, y_test_pred))

# Check misclassified data
misclassified_data = np.where(y_test_pred != y_test_array)[0]
print(f"The number of misclassifications: {len(misclassified_data)}")

# Check the ratio of misclassifications
misclassification_ratio = (len(misclassified_data) / len(y_test_array)) * 100
# Round the number
rounded_ratio = round(misclassification_ratio, 2)
print(f"Proportion of misclassifications: {rounded_ratio}%")

# Iterate over misclassified data for error analysis
for idx in misclassified_data[:30]:
    input_text = X_test[idx]
    true_label = y_test[idx]
    predicted_label = y_test_pred[idx]

    # Print information about the misclassified data
    print("Input Text:", input_text)
    print("Actual Label:", true_label)
    print("Predicted Label:", predicted_label)
    print()

```

	precision	recall	f1-score	support
0	0.97	0.95	0.96	40
1	0.89	0.97	0.93	40
2	0.93	0.97	0.95	40

3	0.76	0.95	0.84	40
4	0.91	0.78	0.84	40
5	0.48	0.75	0.59	40
6	0.83	0.85	0.84	40
7	0.76	0.70	0.73	40
8	0.88	0.88	0.88	40
9	1.00	0.90	0.95	40
10	0.77	0.82	0.80	40
11	0.79	0.68	0.73	40
12	0.62	0.88	0.73	40
13	0.90	0.88	0.89	40
14	0.76	0.85	0.80	40
15	0.88	0.88	0.88	40
16	0.72	0.65	0.68	40
17	0.77	0.93	0.84	40
18	0.97	0.78	0.86	40
19	0.88	0.88	0.88	40
20	0.88	0.75	0.81	40
21	0.82	0.35	0.49	40
22	0.59	0.75	0.66	40
23	0.89	0.85	0.87	40
24	0.89	0.97	0.93	40
25	0.70	0.80	0.74	40
26	0.69	0.85	0.76	40
27	0.88	0.75	0.81	40
28	0.80	0.70	0.75	40
29	0.89	0.62	0.74	40
30	0.97	0.85	0.91	40
31	0.92	0.88	0.90	40
32	0.92	0.88	0.90	40
33	0.65	0.85	0.74	40
34	0.72	0.78	0.75	40
35	0.76	0.65	0.70	40
36	0.66	0.78	0.71	40
37	0.66	0.68	0.67	40
38	0.40	0.95	0.56	40
39	0.76	0.85	0.80	40
40	0.88	0.93	0.90	40
41	0.68	0.68	0.68	40
42	0.93	0.93	0.93	40
43	0.69	0.68	0.68	40
44	0.81	0.65	0.72	40
45	0.79	0.78	0.78	40
46	0.71	0.80	0.75	40
47	0.71	0.60	0.65	40
48	0.65	0.60	0.62	40
49	0.94	0.72	0.82	40
50	0.82	0.68	0.74	40

	51	0.79	0.82	0.80	40
	52	0.81	0.75	0.78	40
	53	0.72	0.78	0.75	40
	54	0.71	0.60	0.65	40
	55	0.90	0.93	0.91	40
	56	0.95	0.53	0.68	40
	57	0.93	0.93	0.93	40
	58	0.93	0.70	0.80	40
	59	0.83	0.72	0.77	40
	60	0.95	0.90	0.92	40
	61	0.60	0.70	0.64	40
	62	0.72	0.72	0.73	40
	63	0.70	0.88	0.78	40
	64	0.82	0.82	0.82	40
	65	0.68	0.75	0.71	40
	66	0.80	0.70	0.75	40
	67	0.70	0.70	0.70	40
	68	0.93	0.35	0.51	40
	69	0.49	0.75	0.59	40
	70	0.89	1.00	0.94	40
	71	0.91	1.00	0.95	40
	72	0.79	0.82	0.80	40
	73	0.97	0.88	0.92	40
	74	0.65	0.38	0.48	40
	75	0.89	0.78	0.83	40
	76	0.93	0.62	0.75	40
	accuracy			0.78	3080
	macro avg	0.80	0.78	0.78	3080
	weighted avg	0.80	0.78	0.78	3080

The number of misclassifications: 682

Proportion of misclassifications: 22.14%

Input Text: locate card

Actual Label: 11

Predicted Label: 41

Input Text: ordered card arrived help please

Actual Label: 11

Predicted Label: 12

Input Text: way know card arrive

Actual Label: 11

Predicted Label: 12

Input Text: get card

Actual Label: 11

Predicted Label: 43

Input Text: received card

Actual Label: 11

Predicted Label: 43

Input Text: track card

Actual Label: 11

Predicted Label: 12

Input Text: long card delivery take

Actual Label: 11

Predicted Label: 12

Input Text: waiting longer expected bank card could provide information arrive

Actual Label: 11

Predicted Label: 12

Input Text: get card yet lost

Actual Label: 11

Predicted Label: 41

Input Text: status card ordered

Actual Label: 11

Predicted Label: 22

Input Text: long new card take arrive

Actual Label: 11

Predicted Label: 12

Input Text: know card arrive

Actual Label: 11

Predicted Label: 12

Input Text: tracking info available

Actual Label: 11

Predicted Label: 74

Input Text: received new card dont see app anywhere

Actual Label: 13

Predicted Label: 12

Input Text: add card account

Actual Label: 13

Predicted Label: 39

Input Text: put old card back system found

Actual Label: 13

Predicted Label: 41

Input Text: way make old card usable app

Actual Label: 13

Predicted Label: 11

Input Text: found lost stolen card way link card account app

Actual Label: 13

Predicted Label: 42

Input Text: good time exchange

Actual Label: 32

Predicted Label: 33

Input Text: much get exchange rate

Actual Label: 32

Predicted Label: 17

Input Text: im trying figure current exchange rate

Actual Label: 32

Predicted Label: 17

Input Text: kind foreign exchange rate get exchange money

Actual Label: 32

Predicted Label: 17

Input Text: rate get determined

Actual Label: 32

Predicted Label: 17

Input Text: made currency exchange think charged

Actual Label: 17

Predicted Label: 31

Input Text: charged

Actual Label: 17

Predicted Label: 63

Input Text: conversion value card payments incorrect

Actual Label: 17

Predicted Label: 15

Input Text: explain random charge

Actual Label: 34

Predicted Label: 63

Input Text: remember purchasing anything & statement please tell

Actual Label: 34

Predicted Label: 53

Input Text: transaction credited
Actual Label: 34
Predicted Label: 62

Input Text: many fees statement
Actual Label: 34
Predicted Label: 57

1.7 Tokenize the text in the dataset

1.7.1 For DistilBERT Model

ver. 1) Use the preprocessed dataset

```
[ ]: # Change the format as dataframe
banking77_preprocessed.reset_format()

[ ]: # Perform train-test split to make a validation set
# Export only data set to split (training 80%, validation 20% from the training_
↪set)
dataset_dict = banking77_preprocessed['train'].train_test_split(test_size=0.2)

[ ]: # Change the name 'test' to 'validation'
dataset_dict['validation'] = dataset_dict.pop('test')

[ ]: # Check the dataset dictionary
dataset_dict

[ ]: DatasetDict({
  train: Dataset({
    features: ['text', 'label'],
    num_rows: 8002
  })
  validation: Dataset({
    features: ['text', 'label'],
    num_rows: 2001
  })
})

[ ]: # Define training, validation and test sets
trainset = dataset_dict['train']
valset = dataset_dict['validation']
testset = banking77_preprocessed['test']

[ ]: # Change the format as dataframe to save the test set
banking77_preprocessed.set_format(type='pandas')
```



```
[ ]: # Define a dataframe for the test set
test_df = banking77_preprocessed['test'][:]
```

```
[ ]: # Specify the test set file path
csv_file_path = "/content/drive/MyDrive/1. NLP CW/DistilBERT/test set/testset.
↳CSV"

# Save the DataFrame as CSV
test_df.to_csv(csv_file_path, index=False)
```

```
[ ]: # Check the number of data points
pprint(len(trainset))
pprint(len(valset))
pprint(len(testset))
```

```
8002
2001
3080
```

```
[ ]: # This code is derived from lab tutorial 8
# Import libraries
from transformers import DistilBertTokenizer

# Tokenization
tokenizer = DistilBertTokenizer.from_pretrained("distilbert-base-uncased")
```

```
/usr/local/lib/python3.10/dist-packages/huggingface_hub/file_download.py:1132:
FutureWarning: `resume_download` is deprecated and will be removed in version
1.0.0. Downloads always resume when possible. If you want to force a new
download, use `force_download=True`.
warnings.warn(
```

```
[ ]: # This code is derived from lab tutorial 8
# Tokenize the data
def tokenize(batch):
    return tokenizer(batch['text'], padding='max_length', truncation=True,
↳max_length=29) # Define the maximum length as 29

train_set = dataset_dict['train'].map(tokenize, batched=True)
val_set = dataset_dict['validation'].map(tokenize, batched=True)
test_set = banking77_preprocessed['test'].map(tokenize, batched=True)
```

```
Map: 0%|          | 0/8002 [00:00<?, ? examples/s]
Map: 0%|          | 0/2001 [00:00<?, ? examples/s]
Map: 0%|          | 0/3080 [00:00<?, ? examples/s]
```

```
[ ]: # Check the inside of test set
pprint(test_set[:1], sort_dicts=False)
```

```
{'text': ['locate card'],
 'label': [11],
 'input_ids': [[101,
                  12453,
                  4003,
                  102,
                  0,
                  0,
                  0,
                  0,
                  0,
                  0,
                  0,
                  0,
                  0,
                  0,
                  0,
                  0,
                  0,
                  0,
                  0,
                  0,
                  0,
                  0,
                  0,
                  0,
                  0,
                  0,
                  0,
                  0,
                  0,
                  0,
                  0],
 'attention_mask': [[1,
                      1,
                      1,
                      0,
                      0,
                      0,
                      0,
                      0,
                      0,
                      0,
                      0,
                      0,
                      0,
                      0,
                      0,
                      0,
                      0,
                      0,
                      0,
                      0,
                      0,
                      0,
                      0,
                      0,
                      0,
                      0,
                      0,
                      0,
                      0,
                      0,
                      0]]}
```

```
[ ]: # This code is derived from lab tutorial 8
      # Set the data format
      train_set.set_format('pt', columns=['input_ids', 'attention_mask', 'label'])
      val_set.set_format('pt', columns=['input_ids', 'attention_mask', 'label'])
      test_set.set_format('pt', columns=['input_ids', 'attention_mask', 'label'])
```

```
{'attention_mask': tensor([[1, 1, 1, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0],
                           [0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]]),
      'input_ids': tensor([[ 101,  5356,  2134,  2102,  2131,  3065, 10439,   102,
                             0,    0,    0,    0,    0,    0,    0,    0,    0],
                          [0,    0,    0,    0,    0,    0,    0,    0,    0]]),
      'label': tensor([20])}
{'attention_mask': tensor([[1, 1, 1, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0],
                           [0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]]),
      'input_ids': tensor([[ 101, 3841, 12879, 24108, 2854, 23439,   102,    0,
                             0,    0,    0,    0,    0,    0,    0,    0,    0],
                          [0,    0,    0,    0,    0,    0,    0,    0,    0]]),
      'label': tensor([7])}
{'attention_mask': tensor([[1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0],
                           [0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]]),
      'input_ids': tensor([[ 101, 12453,  4003,   102,    0,    0,    0,    0,
```

```
0,      0,
      0,      0,      0,      0,      0,      0,      0,      0,      0,      0,
      0,      0,      0,      0,      0,      0,      0,      0,      0]],
'label': tensor([11])}
```

DistilBERT

```
[ ]: # Import libraries
import torch
import torch.nn as nn

# Check the run time
device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
print(device)

cuda:0

[ ]: # This code is derived from lab tutorial 8
from transformers import DistilBertModel

model = DistilBertModel.from_pretrained('distilbert-base-uncased')

[ ]: # This code is derived from lab tutorial 8
# Define the model architecture
class DistilBERT(nn.Module):
    def __init__(self, model):
        super(DistilBERT, self).__init__()
        self.model = model
        self.linear = nn.Linear(768, 77) # 77 classes

    def forward(self, input_ids, attention_mask):
        outputs = self.model(input_ids=input_ids, attention_mask=attention_mask)
        last_hidden_state = outputs.last_hidden_state[:, 0, :]
        logits = self.linear(last_hidden_state)
        return logits

# Define the model
model = DistilBERT(model)
model.to(device)

[ ]: DistilBERT(
  (model): DistilBertModel(
    (embeddings): Embeddings(
      (word_embeddings): Embedding(30522, 768, padding_idx=0)
      (position_embeddings): Embedding(512, 768)
      (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
      (dropout): Dropout(p=0.1, inplace=False)
    )
```

```

(transformer): Transformer(
  (layer): ModuleList(
    (0-5): 6 x TransformerBlock(
      (attention): MultiHeadSelfAttention(
        (dropout): Dropout(p=0.1, inplace=False)
        (q_lin): Linear(in_features=768, out_features=768, bias=True)
        (k_lin): Linear(in_features=768, out_features=768, bias=True)
        (v_lin): Linear(in_features=768, out_features=768, bias=True)
        (out_lin): Linear(in_features=768, out_features=768, bias=True)
      )
      (sa_layer_norm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
      (ffn): FFN(
        (dropout): Dropout(p=0.1, inplace=False)
        (lin1): Linear(in_features=768, out_features=3072, bias=True)
        (lin2): Linear(in_features=3072, out_features=768, bias=True)
        (activation): GELUActivation()
      )
      (output_layer_norm): LayerNorm((768,), eps=1e-12,
elementwise_affine=True)
    )
  )
)
(linear): Linear(in_features=768, out_features=77, bias=True)
)

```

```

[ ]: # This code is derived from lab tutorial 8
# Set up the optimizer and loss function
optimizer = torch.optim.Adam(model.parameters(), lr=2e-5)
loss_fn = nn.CrossEntropyLoss()

```

```

[ ]: # Set up the data loader for each dataset
train_loader = torch.utils.data.DataLoader(train_set, batch_size=32,
↪shuffle=True)
val_loader = torch.utils.data.DataLoader(val_set, batch_size=32, shuffle=True)
test_loader = torch.utils.data.DataLoader(test_set, batch_size=32,
↪shuffle=False)

```

```

[ ]: # Import the library to save the test set for the distilBERT model
import os

# Define the save function for test loader
def save_test_data(test_loader, save_dir):
    for i, batch in enumerate(test_loader):
        input_ids = batch['input_ids']
        attention_mask = batch['attention_mask']
        labels = batch['label']

```

```

        # Save each batch of data
        torch.save((input_ids, attention_mask, labels), os.path.join(save_dir,
↪f"test_batch_{i}.pt"))

# Specify the directory to save test data
save_dir = "/content/drive/MyDrive/1. NLP CW/DistilBERT"
save_test_data(test_loader, save_dir)

```

```

[ ]: # Define the folder path to save the state dictionary
folder_path = '/content/drive/MyDrive/1. NLP CW/DistilBERT/'

# Define the dictionary file path for the model checkpoint
model_save_path = folder_path + 'distilBERT_model.pth'

```

```

[ ]: #Free up GPU memory
torch.cuda.empty_cache()

```

```

[ ]: # Define the train function
# Import time to measure the training time
import time

def train_and_evaluate(model, train_loader, val_loader, optimizer, loss_fn,
↪device, model_save_path):
    train_losses, val_losses = [], [] # Empty lists to store losses
    train_accuracies, val_accuracies = [], [] # Empty lists to store accuracies

    # Measure the total training time
    total_start_time = time.time()

    for epoch in range(5):
        # Training
        start_time = time.time() # Measure each training time
        model.train()

        epoch_train_loss = 0.0
        correct_train, total_train = 0, 0

        for batch in train_loader:
            input_ids = batch['input_ids'].to(device)
            attention_mask = batch['attention_mask'].to(device)
            labels = batch['label'].to(device)

            optimizer.zero_grad()
            outputs = model(input_ids, attention_mask)
            loss = loss_fn(outputs, labels)
            loss.backward()
            optimizer.step()

```

```

        epoch_train_loss += loss.item() * input_ids.size(0) # Check the
↪train loss per epoch

        predictions_train = torch.round(torch.softmax(outputs, dim=1))
        predicted_train = torch.argmax(predictions_train, dim=1)
        total_train += labels.size(0)
        correct_train += (predicted_train == labels).sum().item()

    train_loss = epoch_train_loss / len(train_loader.dataset)
    train_accuracy = correct_train / total_train
    train_losses.append(train_loss) # Total train loss
    train_accuracies.append(train_accuracy)

    # Validation
    model.eval()
    correct_val, total_val = 0, 0
    epoch_val_loss = 0.0

    with torch.no_grad():
        for batch in val_loader:
            input_ids = batch['input_ids'].to(device)
            attention_mask = batch['attention_mask'].to(device)
            labels = batch['label'].to(device)

            outputs = model(input_ids, attention_mask)
            predictions_val = torch.round(torch.softmax(outputs, dim=1))
            predicted_val = torch.argmax(predictions_val, dim=1)

            loss_val = loss_fn(outputs, labels)
            epoch_val_loss += loss_val.item() * input_ids.size(0) # Check
↪the validation loss per epoch

            total_val += labels.size(0)
            correct_val += (predicted_val == labels).sum().item()

    val_loss = epoch_val_loss / len(val_loader.dataset)
    val_accuracy = correct_val / total_val
    val_losses.append(val_loss) # Total validation loss
    val_accuracies.append(val_accuracy)

    end_time = time.time()
    each_train_duration = end_time - start_time

    # Add 1 to epoch as it starts from 0

```

```

        print(f'Epoch {epoch+1} - Training Time: {each_train_duration:.3f}␣
↪seconds, Train Loss: {train_loss:.4f}, Train Accuracy: {train_accuracy:.4f},␣
↪Validation Loss: {val_loss:.4f}, Validation Accuracy: {val_accuracy:.4f}')

    total_end_time = time.time()
    total_train_duration = end_time - start_time
    print(f'Total training time: {total_train_duration:.3f} seconds')

    # Save the state dictionary
    torch.save(model.state_dict(), model_save_path)
    # Define the file name and path to save the model itself
    # This saving model code is derived from the tutorial of Huggingface's␣
↪DistilBERT (A notebook on how to finetune DistilBERT for multiclass␣
↪classification with PyTorch)
    # (https://huggingface.co/docs/transformers/en/model_doc/
↪distilbert#transformers.DistilBertConfig)
    output_model_file = '/content/drive/MyDrive/1. NLP CW/DistilBERT/
↪processed_distilbert.bin'
    model_to_save = model
    # Save the model itself
    torch.save(model_to_save, output_model_file)
    print('Model and state dictionary have been saved')

    return train_losses, val_losses, train_accuracies, val_accuracies

```

```

[ ]: # Train the model
train_losses, val_losses, train_accuracies, val_accuracies =␣
↪train_and_evaluate(model, train_loader, val_loader, optimizer, loss_fn,␣
↪device, model_save_path)

```

```

Epoch 1 - Training Time: 24.662 seconds, Train Loss: 2.8097, Train Accuracy:
0.0452, Validation Loss: 1.3046, Validation Accuracy: 0.3323
Epoch 2 - Training Time: 23.877 seconds, Train Loss: 1.0237, Train Accuracy:
0.4834, Validation Loss: 0.7290, Validation Accuracy: 0.7026
Epoch 3 - Training Time: 24.640 seconds, Train Loss: 0.5813, Train Accuracy:
0.7652, Validation Loss: 0.5454, Validation Accuracy: 0.8086
Epoch 4 - Training Time: 24.036 seconds, Train Loss: 0.3852, Train Accuracy:
0.8585, Validation Loss: 0.4629, Validation Accuracy: 0.8471
Epoch 5 - Training Time: 24.423 seconds, Train Loss: 0.2616, Train Accuracy:
0.9136, Validation Loss: 0.4399, Validation Accuracy: 0.8576
Total training time: 24.423 seconds
Model and state dictionary have been saved

```

```

[ ]: # Plot the loss and accuracy
# Define the plot fuction
def plot_training_curve(train_losses, val_losses, train_accuracies,␣
↪val_accuracies):

```



```
epochs = range(1, len(train_losses) + 1) # Add 1 to the length of the list,
↳ as the epoch starts from 0
```


```
plt.figure(figsize=(12, 8))

# Plot training and validation losses
plt.subplot(1, 2, 1)
plt.plot(epochs, train_losses, label='Training Loss')
plt.plot(epochs, val_losses, label='Validation Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.title('Training and Validation Loss')
plt.legend()

# Plot validation accuracy
plt.subplot(1, 2, 2)
plt.plot(epochs, train_accuracies, label='Training Accuracy')
plt.plot(epochs, val_accuracies, label='Validation Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.title('Validation Accuracy')
plt.legend()

plt.tight_layout()
plt.show()
```

```
[ ]: # Plot the loss and accuracy of train and validation
plot_training_curve(train_losses, val_losses, train_accuracies, val_accuracies)
```



```
[ ]: # Load the model to test
test_model = torch.load('/content/drive/MyDrive/1. NLP CW/DistilBERT/
    processed_distilbert.bin', map_location=torch.device('cpu'))
# Match the state dictionary to the loaded model
state_dict = torch.load(model_save_path, map_location=torch.device('cpu'))
test_model.load_state_dict(state_dict)
```

```
[ ]: <All keys matched successfully>
```

```
[ ]: # Define the test function
def evaluate(model, test_loader):
    model.eval()
    correct = 0
    total = 0
    predictions_list = []
    labels_list = []
    with torch.no_grad():
        for batch in test_loader:
            input_ids = batch['input_ids']
            attention_mask = batch['attention_mask']
            labels = batch['label']

            outputs = model(input_ids, attention_mask)
```

```

        predictions = torch.argmax(outputs, dim=1)
        correct += (predictions == labels).sum().item()
        total += labels.size(0)

        predictions_list.extend(predictions.cpu().numpy()) # Make sure it
        ↳will run in CPU
        labels_list.extend(labels.cpu().numpy()) # Make sure it will run in
        ↳CPU

    accuracy = correct / total
    precision = precision_score(labels_list, predictions_list,
        ↳average='weighted')
    recall = recall_score(labels_list, predictions_list, average='weighted')
    f1 = f1_score(labels_list, predictions_list, average='weighted')

    return accuracy, precision, recall, f1

```

```

[ ]: # Test the model
test_model.eval()

# Get the test accuracy
test_accuracy, test_precision, test_recall, test_f1 = evaluate(test_model,
    ↳test_loader)

print(f'Test Accuracy: {round((test_accuracy*100), 2)}')
print(f'Test Precision: {round((test_precision*100), 2)}')
print(f'Test Recall: {round((test_recall*100), 2)}')
print(f'Test F1 Score: {round((test_f1*100), 2)}')

```

```

Test Accuracy: 89.55
Test Precision: 90.08
Test Recall: 89.55
Test F1 Score: 89.56

```

```

[ ]: # Make predictions on the test dataset
predictions = model.predict(test_set)

# Get the predicted labels
predicted_labels = [np.argmax(pred) for pred in predictions]

# Get the ground truth labels
true_labels = test_set['label']

# Initialize lists to store misclassified instances
misclassified_texts = []
misclassified_predicted_labels = []
misclassified_true_labels = []

```

```

# Compare predictions with ground truth labels
for i in range(len(true_labels)):
    if predicted_labels[i] != true_labels[i]:
        # Add misclassified instance to lists
        misclassified_texts.append(test_set['text'][i])
        misclassified_predicted_labels.append(predicted_labels[i])
        misclassified_true_labels.append(true_labels[i])

# Print some misclassified instances for analysis
for i in range(min(10, len(misclassified_texts))):
    print("Text:", misclassified_texts[i])
    print("Predicted Label:", misclassified_predicted_labels[i])
    print("True Label:", misclassified_true_labels[i])
    print()

```

ver. 2) Use the unprocessed dataset

```

[ ]: # Perform train-test split to make a validation set
# Export only data set to split (training 80%, validation 20% from the training_
    ↪set)
dataset_dict2 = banking77['train'].train_test_split(test_size=0.2)

```

```

[ ]: # Change the name 'test' to 'validation'
dataset_dict2['validation'] = dataset_dict2.pop('test')

```

```

[ ]: # Check the dataset dictionary
dataset_dict2

```

```

[ ]: DatasetDict({
  train: Dataset({
    features: ['text', 'label'],
    num_rows: 8002
  })
  validation: Dataset({
    features: ['text', 'label'],
    num_rows: 2001
  })
})

```

```

[ ]: # Define a dataframe for the test set
test_df2 = banking77['test'][:]

```

```

[ ]: # Check the dataframe of the test set
test_df2

```

```
[ ]:
      text  label
0      How do I locate my card?      11
1  I still have not received my new card, I order...      11
2  I ordered a card but it has not arrived. Help ...      11
3      Is there a way to know when my card will arrive?      11
4      My card has not arrived yet.      11
...
3075  If i'm not in the UK, can I still get a card?      24
3076      How many countries do you support?      24
3077      What countries do you do business in?      24
3078      What are the countries you operate in.      24
3079      Can the card be mailed and used in Europe?      24
```

[3080 rows x 2 columns]

```
[ ]: # Specify the test set file path
csv_file_path = "/content/drive/MyDrive/1. NLP CW/DistilBERT/test set/testset2.
↳CSV"

# Save the DataFrame as CSV
test_df2.to_csv(csv_file_path, index=False)
```

```
[ ]: # Define training, validation and test sets
trainset2 = dataset_dict2['train']
valset2 = dataset_dict2['validation']
testset2 = banking77['test']
```

```
[ ]: # This code is derived from lab tutorial 8
# Import libraries
from transformers import DistilBertTokenizer

# Tokenization
tokenizer = DistilBertTokenizer.from_pretrained("distilbert-base-uncased")
```

```
/usr/local/lib/python3.10/dist-packages/huggingface_hub/file_download.py:1132:
FutureWarning: `resume_download` is deprecated and will be removed in version
1.0.0. Downloads always resume when possible. If you want to force a new
download, use `force_download=True`.
warnings.warn(
```

```
[ ]: # This code is derived from lab tutorial 8
# Tokenize the data
def tokenize(batch):
    return tokenizer(batch['text'], padding='max_length', truncation=True,
↳max_length=29) # Define the maximum length as 29

train_set2 = dataset_dict2['train'].map(tokenize, batched=True)
```

```
val_set2 = dataset_dict2['validation'].map(tokenize, batched=True)
test_set2 = banking77['test'].map(tokenize, batched=True)
```

```
Map:   0%|          | 0/8002 [00:00<?, ? examples/s]
```

```
Map:   0%|          | 0/2001 [00:00<?, ? examples/s]
```

```
[ ]: # This code is derived from lab tutorial 8
      # Set the data format
      train_set2.set_format('pt', columns=['input_ids', 'attention_mask', 'label'])
      val_set2.set_format('pt', columns=['input_ids', 'attention_mask', 'label'])
      test_set2.set_format('pt', columns=['input_ids', 'attention_mask', 'label'])
```

DistilBERT

```
[ ]: # Import libraries
      import torch
      import torch.nn as nn

      # Check the run time
      device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
      print(device)
```

```
cuda:0
```

```
[ ]: # This code is derived from lab tutorial 8
      from transformers import DistilBertModel

      model = DistilBertModel.from_pretrained('distilbert-base-uncased')
```

```
[ ]: # This code is derived from lab tutorial 8
      # Define the model architecture
      class DistilBERT(nn.Module):
          def __init__(self, model):
              super(DistilBERT, self).__init__()
              self.model = model
              self.linear = nn.Linear(768, 77) # 77 classes

          def forward(self, input_ids, attention_mask):
              outputs = self.model(input_ids=input_ids, attention_mask=attention_mask)
              last_hidden_state = outputs.last_hidden_state[:, 0, :]
              logits = self.linear(last_hidden_state)
              return logits

      # Define the model
      model = DistilBERT(model)
      model.to(device)
```

```
[ ]: DistilBERT(
  (model): DistilBertModel(
    (embeddings): Embeddings(
      (word_embeddings): Embedding(30522, 768, padding_idx=0)
      (position_embeddings): Embedding(512, 768)
      (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
      (dropout): Dropout(p=0.1, inplace=False)
    )
    (transformer): Transformer(
      (layer): ModuleList(
        (0-5): 6 x TransformerBlock(
          (attention): MultiHeadSelfAttention(
            (dropout): Dropout(p=0.1, inplace=False)
            (q_lin): Linear(in_features=768, out_features=768, bias=True)
            (k_lin): Linear(in_features=768, out_features=768, bias=True)
            (v_lin): Linear(in_features=768, out_features=768, bias=True)
            (out_lin): Linear(in_features=768, out_features=768, bias=True)
          )
          (sa_layer_norm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
          (ffn): FFN(
            (dropout): Dropout(p=0.1, inplace=False)
            (lin1): Linear(in_features=768, out_features=3072, bias=True)
            (lin2): Linear(in_features=3072, out_features=768, bias=True)
            (activation): GELUActivation()
          )
          (output_layer_norm): LayerNorm((768,), eps=1e-12,
elementwise_affine=True)
        )
      )
    )
  )
  (linear): Linear(in_features=768, out_features=77, bias=True)
)
```

```
[ ]: # This code is derived from lab tutorial 8
# Set up the optimizer and loss function
optimizer = torch.optim.Adam(model.parameters(), lr=2e-5)
loss_fn = nn.CrossEntropyLoss()
```

```
[ ]: # Set up the data loader for each dataset
train_loader2 = torch.utils.data.DataLoader(train_set2, batch_size=32,
↪shuffle=True)
val_loader2 = torch.utils.data.DataLoader(val_set2, batch_size=32, shuffle=True)
test_loader2 = torch.utils.data.DataLoader(test_set2, batch_size=32,
↪shuffle=False)
```

```
[ ]: # Define the folder path to save the state dictionary
folder_path = '/content/drive/MyDrive/1. NLP CW/DistilBERT/'

# Define the dictionary file path for the model checkpoint
model_save_path = folder_path + 'unprocessed_distilBERT_model.pth'

[ ]: #Free up GPU memory
torch.cuda.empty_cache()

[ ]: # Define the train function
# Import time to measure the training time
import time

def train_and_evaluate(model, train_loader, val_loader, optimizer, loss_fn,
↪device, model_save_path):
    train_losses, val_losses = [], [] # Empty lists to store losses
    train_accuracies, val_accuracies = [], [] # Empty lists to store accuracies

    # Measure the total training time
    total_start_time = time.time()

    for epoch in range(5):
        # Training
        start_time = time.time() # Measure each training time
        model.train()

        epoch_train_loss = 0.0
        correct_train, total_train = 0, 0

        for batch in train_loader:
            input_ids = batch['input_ids'].to(device)
            attention_mask = batch['attention_mask'].to(device)
            labels = batch['label'].to(device)

            optimizer.zero_grad()
            outputs = model(input_ids, attention_mask)
            loss = loss_fn(outputs, labels)
            loss.backward()
            optimizer.step()

            epoch_train_loss += loss.item() * input_ids.size(0)

        predictions_train = torch.round(torch.softmax(outputs, dim=1))
        predicted_train = torch.argmax(predictions_train, dim=1)
        total_train += labels.size(0)
        correct_train += (predicted_train == labels).sum().item()
```



```

train_loss = epoch_train_loss / len(train_loader.dataset)
train_accuracy = correct_train / total_train
train_losses.append(train_loss)
train_accuracies.append(train_accuracy)

# Validation
model.eval()
correct_val, total_val = 0, 0
epoch_val_loss = 0.0

with torch.no_grad():
    for batch in val_loader:
        input_ids = batch['input_ids'].to(device)
        attention_mask = batch['attention_mask'].to(device)
        labels = batch['label'].to(device)

        outputs = model(input_ids, attention_mask)
        predictions_val = torch.round(torch.softmax(outputs, dim=1))
        predicted_val = torch.argmax(predictions_val, dim=1)

        loss_val = loss_fn(outputs, labels)
        epoch_val_loss += loss_val.item() * input_ids.size(0)

        total_val += labels.size(0)
        correct_val += (predicted_val == labels).sum().item()

val_loss = epoch_val_loss / len(val_loader.dataset)
val_accuracy = correct_val / total_val
val_losses.append(val_loss)
val_accuracies.append(val_accuracy)

end_time = time.time()
each_train_duration = end_time - start_time

# Add 1 to epoch as it starts from 0
print(f'Epoch {epoch+1} - Training Time: {each_train_duration:.3f}␣
↪seconds, Train Loss: {train_loss:.4f}, Train Accuracy: {train_accuracy:.4f},␣
↪Validation Loss: {val_loss:.4f}, Validation Accuracy: {val_accuracy:.4f}')
```

```

total_end_time = time.time()
total_train_duration = end_time - start_time
print(f'Total training time: {total_train_duration:.3f} seconds')
```

```

# Save the state dictionary
torch.save(model.state_dict(), model_save_path)
# Define the file name and path to save the model itself
```

```

    # This saving model code is derived from the tutorial of Huggingface's
    ↪ DistilBERT (A notebook on how to finetune DistilBERT for multiclass
    ↪ classification with PyTorch)
    # (https://huggingface.co/docs/transformers/en/model_doc/
    ↪ distilbert#transformers.DistilBertConfig)
    output_model_file = '/content/drive/MyDrive/1. NLP CW/DistilBERT/
    ↪ unprocessed_distilbert.bin'
    model_to_save = model
    # Save the model itself
    torch.save(model_to_save, output_model_file)
    print('Model and state dictionary have been saved')

    return train_losses, val_losses, train_accuracies, val_accuracies

```

```

[ ]: # Train the model
train_losses, val_losses, train_accuracies, val_accuracies =
    ↪ train_and_evaluate(model, train_loader2, val_loader2, optimizer, loss_fn,
    ↪ device, model_save_path)

```

```

Epoch 1 - Training Time: 24.963 seconds, Train Loss: 2.7830, Train Accuracy:
0.0427, Validation Loss: 1.2840, Validation Accuracy: 0.3243
Epoch 2 - Training Time: 25.533 seconds, Train Loss: 0.9567, Train Accuracy:
0.5292, Validation Loss: 0.6202, Validation Accuracy: 0.7576
Epoch 3 - Training Time: 23.989 seconds, Train Loss: 0.5104, Train Accuracy:
0.8082, Validation Loss: 0.4246, Validation Accuracy: 0.8546
Epoch 4 - Training Time: 24.265 seconds, Train Loss: 0.3171, Train Accuracy:
0.8967, Validation Loss: 0.3638, Validation Accuracy: 0.8801
Epoch 5 - Training Time: 24.323 seconds, Train Loss: 0.2051, Train Accuracy:
0.9405, Validation Loss: 0.3408, Validation Accuracy: 0.8831
Total training time: 24.323 seconds
Model and state dictionary have been saved

```

```

[ ]: # Plot the loss and accuracy
# Define the plot fuction
def plot_training_curve(train_losses, val_losses, train_accuracies,
    ↪ val_accuracies):
    epochs = range(1, len(train_losses) + 1) # Add 1 to the length of the list
    ↪ as the epoch starts from 0

    plt.figure(figsize=(12, 8))

    # Plot training and validation losses
    plt.subplot(1, 2, 1)
    plt.plot(epochs, train_losses, label='Training Loss')
    plt.plot(epochs, val_losses, label='Validation Loss')
    plt.xlabel('Epoch')
    plt.ylabel('Loss')


```

```
plt.title('Training and Validation Loss')
plt.legend()

# Plot validation accuracy
plt.subplot(1, 2, 2)
plt.plot(epochs, train_accuracies, label='Training Accuracy')
plt.plot(epochs, val_accuracies, label='Validation Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.title('Validation Accuracy')
plt.legend()

plt.tight_layout()
plt.show()
```

```
[ ]: # Plot the loss and accuracy of train and validation
plot_training_curve(train_losses, val_losses, train_accuracies, val_accuracies)
```



```
[ ]: # Load the model to test
test_model = torch.load('/content/drive/MyDrive/1. NLP CW/DistilBERT/
↳unprocessed_distilbert.bin', map_location=torch.device('cpu'))
# Match the state dictionary to the loaded model
state_dict = torch.load(model_save_path, map_location=torch.device('cpu'))
```

```
test_model.load_state_dict(state_dict)
```

```
[ ]: <All keys matched successfully>
```

```
[ ]: # Define the test function
def evaluate(model, test_loader):
    model.eval()
    correct = 0
    total = 0
    predictions_list = []
    labels_list = []
    with torch.no_grad():
        for batch in test_loader:
            input_ids = batch['input_ids']
            attention_mask = batch['attention_mask']
            labels = batch['label']

            outputs = model(input_ids, attention_mask)
            predictions = torch.argmax(outputs, dim=1)
            correct += (predictions == labels).sum().item()
            total += labels.size(0)

            predictions_list.extend(predictions.cpu().numpy()) # Make sure it
            →will run in CPU
            labels_list.extend(labels.cpu().numpy()) # Make sure it will run in
            →CPU

    accuracy = correct / total
    precision = precision_score(labels_list, predictions_list,
    →average='weighted')
    recall = recall_score(labels_list, predictions_list, average='weighted')
    f1 = f1_score(labels_list, predictions_list, average='weighted')

    return accuracy, precision, recall, f1
```

```
[ ]: # Test the model
test_model.eval()

# Get the test accuracy
test_accuracy, test_precision, test_recall, test_f1 = evaluate(test_model,
    →test_loader)

print(f'Test Accuracy: {round((test_accuracy*100), 2)}')
print(f'Test Precision: {round((test_precision*100), 2)}')
print(f'Test Recall: {round((test_recall*100), 2)}')
print(f'Test F1 Score: {round((test_f1*100), 2)}')
```

Test Accuracy: 83.28
Test Precision: 85.21
Test Recall: 83.28
Test F1 Score: 83.44

1.8 Intermediate results

The model below does not learn properly when applying the true maximum length.

```
[ ]: # Define the maximum length as 303
true_max_length_train_text = 303
true_max_length_train_text
```

```
[ ]: 303
```

```
[ ]: # Import the library for padding
from tensorflow.keras.preprocessing.sequence import pad_sequences

# Padding and truncation will be added to post-texts
X_train_true_padded = pad_sequences(X_train_sequences,
    ↳maxlen=true_max_length_train_text, padding="post", truncating="post")
X_val_true_padded = pad_sequences(X_val_sequences,
    ↳maxlen=true_max_length_train_text, padding="post", truncating="post")
X_test_true_padded = pad_sequences(X_test_sequences,
    ↳maxlen=true_max_length_train_text, padding="post", truncating="post")
```

```
[ ]: # Check the dimension of the variables
print(X_train_true_padded.shape)
print(X_val_true_padded.shape)
print(X_test_true_padded.shape)
```

```
(7346, 303)
(1837, 303)
(3080, 303)
```

```
[ ]: # Check the first 3 elements of all X train variables
print(X_train_array[3])
print(X_train_sequences[3])
print(X_train_true_padded[3])
```

```
something wrong account balance didnt change transferred money
[50, 31, 2, 93, 23, 68, 127, 3]
[ 50  31   2  93  23  68 127   3   0   0   0   0   0   0   0   0   0   0
   0   0   0   0   0   0   0   0   0   0   0   0   0   0   0   0   0
   0   0   0   0   0   0   0   0   0   0   0   0   0   0   0   0   0
   0   0   0   0   0   0   0   0   0   0   0   0   0   0   0   0   0
   0   0   0   0   0   0   0   0   0   0   0   0   0   0   0   0   0
   0   0   0   0   0   0   0   0   0   0   0   0   0   0   0   0   0]
```

[illegible]

```
[ ]: # Define the output dimension for the embedding layer and hidden units
embedding_output_dim = 100
hidden_unit = 30
nlabel = 77

model = keras.models.Sequential()
model.add(layers.Embedding(voca_size, embedding_output_dim))
model.add(layers.LSTM(hidden_unit))
model.add(layers.Dense(nlabel, activation='softmax'))

# Compile the model
model.compile(optimizer='adam', loss='sparse_categorical_crossentropy',
              metrics=['accuracy']) #, run_eagerly=True

# Summary the model
model.summary()
```

```
Model: "sequential"
```

Layer (type)	Output Shape	Param #
embedding (Embedding)	(None, None, 100)	206900
lstm (LSTM)	(None, 30)	15720
dense (Dense)	(None, 77)	2387
Total params: 225007 (878.93 KB)		
Trainable params: 225007 (878.93 KB)		
Non-trainable params: 0 (0.00 Byte)		

```
[ ]: # Define the folder path to save the model
folder path = '/content/drive/MyDrive/1. NLP CW/'
```

```

# Define the file path for the model checkpoint
model_checkpoint_path = folder_path + 'LSTM1.keras'

# Define the model checkpoint
mc = tf.keras.callbacks.ModelCheckpoint(
    filepath=model_checkpoint_path,
    monitor='val_accuracy',
    mode='max',
    save_best_only=True)

```

```

[ ]: # Define early stopping
es = tf.keras.callbacks.EarlyStopping(monitor='val_loss', patience=3) # Random_
    ↳ number of patience

```

```

[ ]: # Import time to measure the elapsed time
import time

# Measure time before training
start_time = time.time()

# Fit the model
history = model.fit(
    X_train_true_padded, y_train,
    epochs = 100,
    validation_data = (X_val_true_padded, y_val),
    callbacks = [mc, es],
    batch_size = 32)

# End the training time
end_time = time.time()

# Measure the training time
training_time = end_time - start_time
print("Training time:", training_time, "seconds")

```

```

Epoch 1/100
230/230 [=====] - 54s 199ms/step - loss: 4.3260 -
accuracy: 0.0174 - val_loss: 4.3061 - val_accuracy: 0.0196
Epoch 2/100
230/230 [=====] - 34s 147ms/step - loss: 4.3104 -
accuracy: 0.0166 - val_loss: 4.3029 - val_accuracy: 0.0196
Epoch 3/100
230/230 [=====] - 35s 152ms/step - loss: 4.3080 -
accuracy: 0.0186 - val_loss: 4.3034 - val_accuracy: 0.0196
Epoch 4/100
230/230 [=====] - 33s 145ms/step - loss: 4.3072 -

```

```

accuracy: 0.0180 - val_loss: 4.3038 - val_accuracy: 0.0191
Epoch 5/100
230/230 [=====] - 34s 147ms/step - loss: 4.3070 -
accuracy: 0.0181 - val_loss: 4.3025 - val_accuracy: 0.0196
Epoch 6/100
230/230 [=====] - 32s 141ms/step - loss: 4.3070 -
accuracy: 0.0200 - val_loss: 4.3025 - val_accuracy: 0.0196
Epoch 7/100
230/230 [=====] - 34s 147ms/step - loss: 4.3062 -
accuracy: 0.0189 - val_loss: 4.3024 - val_accuracy: 0.0191
Epoch 8/100
230/230 [=====] - 33s 143ms/step - loss: 4.3063 -
accuracy: 0.0188 - val_loss: 4.3028 - val_accuracy: 0.0196
Epoch 9/100
230/230 [=====] - 36s 154ms/step - loss: 4.3059 -
accuracy: 0.0180 - val_loss: 4.3022 - val_accuracy: 0.0196
Epoch 10/100
230/230 [=====] - 32s 139ms/step - loss: 4.3059 -
accuracy: 0.0197 - val_loss: 4.3021 - val_accuracy: 0.0191
Epoch 11/100
230/230 [=====] - 33s 143ms/step - loss: 4.3056 -
accuracy: 0.0166 - val_loss: 4.3023 - val_accuracy: 0.0196
Epoch 12/100
230/230 [=====] - 34s 147ms/step - loss: 4.3058 -
accuracy: 0.0193 - val_loss: 4.3020 - val_accuracy: 0.0196
Epoch 13/100
230/230 [=====] - 32s 139ms/step - loss: 4.3056 -
accuracy: 0.0199 - val_loss: 4.3023 - val_accuracy: 0.0196
Epoch 14/100
230/230 [=====] - 33s 142ms/step - loss: 4.3055 -
accuracy: 0.0186 - val_loss: 4.3020 - val_accuracy: 0.0196
Epoch 15/100
230/230 [=====] - 33s 145ms/step - loss: 4.3052 -
accuracy: 0.0193 - val_loss: 4.3021 - val_accuracy: 0.0191
Epoch 16/100
230/230 [=====] - 33s 143ms/step - loss: 4.3054 -
accuracy: 0.0180 - val_loss: 4.3021 - val_accuracy: 0.0196
Epoch 17/100
230/230 [=====] - 36s 156ms/step - loss: 4.3054 -
accuracy: 0.0174 - val_loss: 4.3021 - val_accuracy: 0.0196
Training time: 589.6596763134003 seconds

```

[]: